

Esperimenti

deeplabv3plus_resnet101

* Esperimento: **deeplabv3plus_resnet101** - **dlprn1010000**

- **Architettura:** deeplabv3plus_resnet101
- **Backbone Weights:** best_deeplabv3plus_resnet101_cityscapes_os16.pth
- **Input size:** 256x256
- **Batch size:** 2
- **Augmentation:** off
- **Scheduler:** off
- **Optimizer:** Adama con lr=1e-4
- **Freezing:** none
- **Auxiliary Head:** off
- **Loss Function / Weighting:** cross-entropy.
- **Main classifier Kernel size:** 1
- **Valutazione risultati:**
 - Migliore Validation Loss ottenuta all'epoca 10:
 Train Loss: 0.3957 | Val Loss: 0.6023 | Mean IoU: 0.5009
 - ▶ IoU classe 0: 0.5078
 - ▶ IoU classe 1: 0.5474
 - ▶ IoU classe 2: 0.5975
 - ▶ IoU classe 3: 0.5411
 - ▶ IoU classe 4: 0.0020
 - ▶ IoU classe 5: 0.3486
 - ▶ IoU classe 6: 0.2300
 - ▶ IoU classe 7: 0.8328
 - ▶ IoU classe 8: 0.9079
 -
 - Migliore mean IoU ottenuta all'epoca 20:
 Train Loss: 0.1819 | Val Loss: 0.6963 | Mean IoU: 0.5805
 - ▶ IoU classe 0: 0.4769

- ► IoU classe 1: 0.6141
- ► IoU classe 2: 0.5972
- ► IoU classe 3: 0.5399
- ► IoU classe 4: 0.3949
- ► IoU classe 5: 0.4916
- ► IoU classe 6: 0.2641
- ► IoU classe 7: 0.8350
- ► IoU classe 8: 0.9075

- **Rappresentazione grafica del training:**



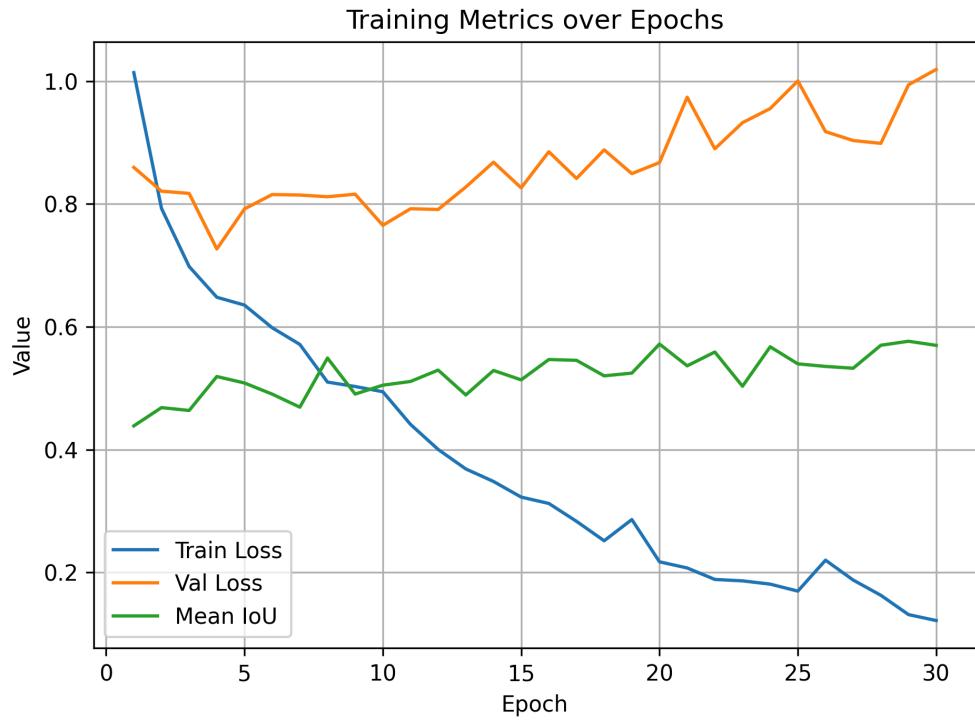
- **Massima memoria RAM occupata durante il training:** 1.28GB
- **Log di training:**
`risultati_training\deeplabv3plus\resnet101\dlprn101000\training_log_dlrn101000.json`
- **Note / Bug Fix / Annotazioni:**

◦

* Esperimento: `deeplabv3plus_resnet101 - dlprn1010001`

- **Architettura:** `deeplabv3plus_resnet101`
- **Backbone Weights:** `best_deeplabv3plus_resnet101_cityscapes_os16.pth`
- **Input size:** 256x256
- **Batch size:** 2

- **Augmentation:** off
- **Scheduler:** off
- **Optimizer:** Adama con lr=1e-4
- **Freezing:** none
- **Auxiliary Head:** off
- **Loss Function / Weighting:** cross-entropy+ class weights euristici [1.0, 1.0, 1.0, 1.0, 5.0, 4.0, 3.5, 0.8, 0.8].
- **Main classifier Kernel size:** 1
- **Valutazione risultati:**
 - Migliore mean IoU ottenuta all'epoca 29:
 - Train Loss: 0.1314 | Val Loss: 0.9941 | Mean IoU: 0.5765
 - ▶ IoU classe 0: 0.5415
 - ▶ IoU classe 1: 0.5790
 - ▶ IoU classe 2: 0.5913
 - ▶ IoU classe 3: 0.5590
 - ▶ IoU classe 4: 0.4902
 - ▶ IoU classe 5: 0.5384
 - ▶ IoU classe 6: 0.1770
 - ▶ IoU classe 7: 0.8164
 - ▶ IoU classe 8: 0.8604
 - Migliore Validation Loss ottenuta all'epoca 4:
 - Train Loss: 0.6481 | Val Loss: 0.7267 | Mean IoU: 0.5191
 - ▶ IoU classe 0: 0.5164
 - ▶ IoU classe 1: 0.5790
 - ▶ IoU classe 2: 0.5396
 - ▶ IoU classe 3: 0.5323
 - ▶ IoU classe 4: 0.3134
 - ▶ IoU classe 5: 0.3297
 - ▶ IoU classe 6: 0.1812
 - ▶ IoU classe 7: 0.7911
 - ▶ IoU classe 8: 0.8862
- **Rappresentazione grafica del training:**

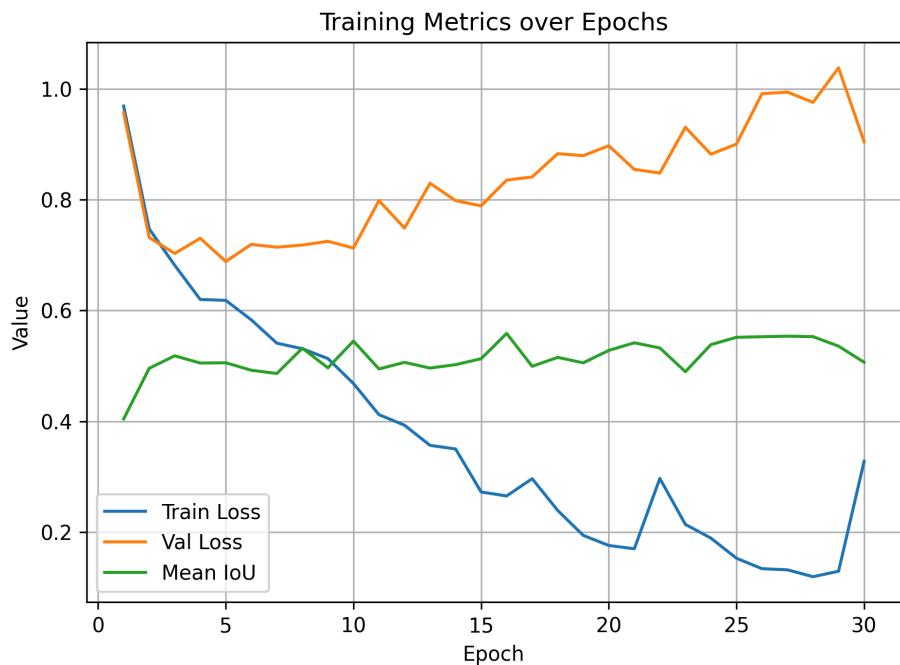


- **Massima memoria RAM occupata durante il training:** 1.28GB
- **Log di training:**
`risultati_training\deeplabv3plus\resnet101\dlprn1010001\training_log_dlrn1010001.json`
- **Note / Bug Fix / Annotazioni:**
 - Bisogna mirare a occupare più memoria GPU, tuttavia raddoppiare la batch size ci porta a 5+GB. Si propone di aumentare la dimensione delle immagini a 512x512

* Esperimento: `deeplabv3plus_resnet101 - dlprn1010002`

- **Architettura:** `deeplabv3plus_resnet101`
- **Backbone Weights:** `best_deeplabv3plus_resnet101_cityscapes_os16.pth`
- **Input size:** 512x512
- **Batch size:** 2
- **Augmentation:** off
- **Scheduler:** off

- **Optimizer:** Adama con lr=1e-4
- **Freezing:** none
- **Auxiliary Head:** off
- **Loss Function / Weighting:** cross-entropy+ class weights euristici [1.0, 1.0, 1.0, 1.0, 5.0, 4.0, 3.5, 0.8, 0.8].
- **Main classifier Kernel size:** 1
- **Valutazione risultati:**
 - Migliore Validation Loss ottenuta all'epoca 5:
 - Train Loss: 0.6183 | Val Loss: 0.6884 | Mean IoU: 0.5057
 - ▶ IoU classe 0: 0.5540
 - ▶ IoU classe 1: 0.5372
 - ▶ IoU classe 2: 0.5092
 - ▶ IoU classe 3: 0.5227
 - ▶ IoU classe 4: 0.1871
 - ▶ IoU classe 5: 0.4341
 - ▶ IoU classe 6: 0.1759
 - ▶ IoU classe 7: 0.7924
 - ▶ IoU classe 8: 0.8871
 - Migliore mean IoU ottenuta all'epoca 16:
 - Train Loss: 0.2655 | Val Loss: 0.8353 | Mean IoU: 0.5586
 - ▶ IoU classe 0: 0.5718
 - ▶ IoU classe 1: 0.5898
 - ▶ IoU classe 2: 0.5835
 - ▶ IoU classe 3: 0.5237
 - ▶ IoU classe 4: 0.3945
 - ▶ IoU classe 5: 0.4986
 - ▶ IoU classe 6: 0.1603
 - ▶ IoU classe 7: 0.8198
 - ▶ IoU classe 8: 0.8990
- **Rappresentazione grafica del training:**



- **Massima memoria RAM occupata durante il training:** 2.88GB
- **Log di training:**
`risultati_training\deeplabv3plus\resnet101\dlprn1010002\training_log_dlprn1010002.json`
- **Note / Bug Fix / Annotazioni:**

○

* Esperimento: deeplabv3plus_resnet101 - dlprn1010003

- **Architettura:** deeplabv3plus_resnet101
- **Backbone Weights:** best_deeplabv3plus_resnet101_cityscapes_os16.pth
- **Input size:** 512x512
- **Batch size:** 2
- **Augmentation:** on
- **Scheduler:** off
- **Optimizer:** Adama con lr=1e-4
- **Freezing:** none
- **Auxiliary Head:** off

- **Loss Function / Weighting:** cross-entropy.
- **Main classifier Kernel size:** 1
- **Valutazione risultati:**
 - Migliore Validation Loss ottenuta all'epoca 4:
 - Train Loss: 0.4977 | Val Loss: 0.5863 | Mean IoU: 0.5376
 - ► IoU classe 0: 0.4975
 - ► IoU classe 1: 0.6046
 - ► IoU classe 2: 0.5578
 - ► IoU classe 3: 0.5504
 - ► IoU classe 4: 0.0461
 - ► IoU classe 5: 0.5470
 - ► IoU classe 6: 0.2577
 - ► IoU classe 7: 0.8245
 - ► IoU classe 8: 0.9130
 - Migliore mean IoU ottenuta all'epoca 26:
 - Train Loss: 0.1132 | Val Loss: 0.7897 | Mean IoU: 0.6194
 - ► IoU classe 0: 0.5176
 - ► IoU classe 1: 0.6358
 - ► IoU classe 2: 0.6137
 - ► IoU classe 3: 0.5713
 - ► IoU classe 4: 0.5040
 - ► IoU classe 5: 0.6013
 - ► IoU classe 6: 0.2633
 - ► IoU classe 7: 0.8464
 - ► IoU classe 8: 0.9194
- **Rappresentazione grafica del training:**



- **Massima memoria RAM occupata durante il training:** 2.42GB
- **Log di training:**
`risultati_training\deeplabv3plus\resnet101\dlprn1010003\training_log_dlrn1010003.json`
- **Note / Bug Fix / Annotazioni:**

○

* Esperimento: `deeplabv3plus_resnet101 - dlprn1010004`

- **Architettura:** `deeplabv3plus_resnet101`
- **Backbone Weights:** `best_deeplabv3plus_resnet101_cityscapes_os16.pth`
- **Input size:** 512x512
- **Batch size:** 2
- **Augmentation:** off
- **Scheduler:** off
- **Optimizer:** Adama con lr=1e-4
- **Freezing:** none

- **Auxiliary Head:** off
- **Loss Function / Weighting:** cross-entropy.
- **Main classifier Kernel size:** 1
- **Valutazione risultati:**
 - Migliore Validation Loss ottenuta all'epoca 12:
 -  Train Loss: 0.3850 | Val Loss: 0.5505 | Mean IoU: 0.5496
 - ▶ IoU classe 0: 0.5188
 - ▶ IoU classe 1: 0.6292
 - ▶ IoU classe 2: 0.5853
 - ▶ IoU classe 3: 0.5909
 - ▶ IoU classe 4: 0.1127
 - ▶ IoU classe 5: 0.5059
 - ▶ IoU classe 6: 0.2224
 - ▶ IoU classe 7: 0.8431
 - ▶ IoU classe 8: 0.9075
 - Migliore mean IoU ottenuta all'epoca 26:
 -  Train Loss: 0.1648 | Val Loss: 0.6548 | Mean IoU: 0.6209
 - ▶ IoU classe 0: 0.5104
 - ▶ IoU classe 1: 0.6407
 - ▶ IoU classe 2: 0.6047
 - ▶ IoU classe 3: 0.6071
 - ▶ IoU classe 4: 0.5819
 - ▶ IoU classe 5: 0.4862
 - ▶ IoU classe 6: 0.2926
 - ▶ IoU classe 7: 0.8411
 - ▶ IoU classe 8: 0.9124
- **Rappresentazione grafica del training:**



- **Massima memoria RAM occupata durante il training:** 2.42GB
- **Log di training:**
`risultati_training\deeplabv3plus\resnet101\dlprn1010004\training_log_dlrn1010004.json`
- **Note / Bug Fix / Annotazioni:**
 - Questo training (img size 512x512) mostra risultati migliori rispetto al corrispettivo con Augmentation e al corrispettivo con pesi euristici per la Loss.

* Esperimento: `deeplabv3plus_resnet101 - dlprn1010005`

- **Architettura:** `deeplabv3plus_resnet101`
- **Backbone Weights:** `best_deeplabv3plus_resnet101_cityscapes_os16.pth`
- **Input size:** 512x512
- **Batch size:** 2
- **Splitting:** on

```

o import numpy as np
o from PIL import Image
o import os
o from iterstrat.ml_stratifiers import
    MultilabelStratifiedShuffleSplit
o import json
o

```

```

o   # sample male etichettati
o   excluded_dirs = {
o       "0007", "0009", "0090", "0095", "0101", "0104",
o       "0105",
o       "0162", "0163", "0284", "0305", "0306", "0307",
o       "0308", "0309", "0310", "0311",
o       "0351", "0372", "0373", "0376",
o       "0498", "0499", "0500", "0501", "0526", "0527",
o       "0530", "0531", "0542",
o       "0564", "0585", "0586", "0587", "0588", "0589",
o       "0590",
o       #seguono nuove immagini da escludere
o       "0000", "0001", "0052"
o   }
o
o
o   # Elenco di cartelle numeriche valide
o   all_dirs = sorted([
o       d for d in os.listdir(DATASET_DIR)
o       if os.path.isdir(os.path.join(DATASET_DIR, d)) and
o       d.isdigit() and d not in excluded_dirs
o   ])
o
o
o   # -----
o   # COSTRUZIONE DELLA MATRICE MULTILABEL
o   # -----
o   def build_presence_matrix(dir_list, n_classes=9):
o       """Return np.array (n_samples, n_classes) with
o       booleans indicating class presence"""
o       presence = np.zeros((len(dir_list), n_classes),
o                           dtype=int)
o
o       for idx, d in enumerate(dir_list):
o           label_path = os.path.join(DATASET_DIR, d,
o                                       "labels.png")
o           lbl = np.array(Image.open(label_path))
o           uniq = np.unique(lbl)
o           presence[idx, uniq] = 1                      # segna le
o           classi presenti
o       return presence
o
o
o   Y = build_presence_matrix(all_dirs, NUM_CLASSES)
o
o
o   # -----
o   # MULTILABEL STRATIFIED SPLIT
o   # -----

```

```

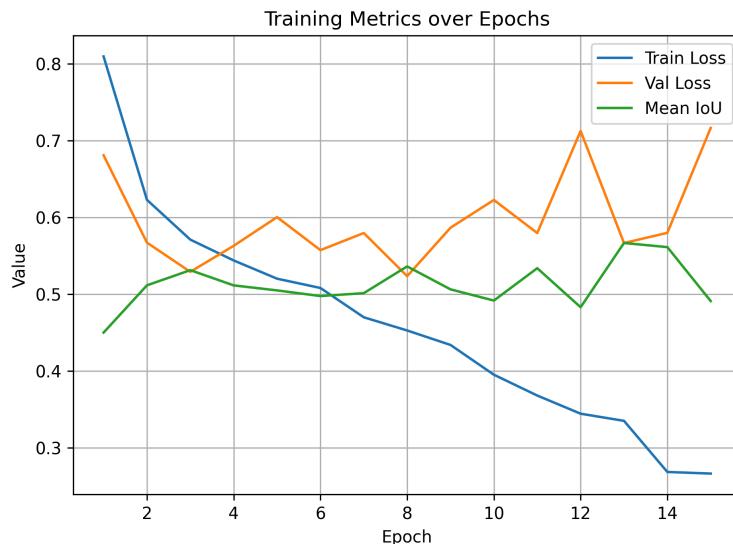
o msss = MultilabelStratifiedShuffleSplit(n_splits=1,
   test_size=0.20, random_state=42)
o train_idx, val_idx =
next(msss.split(np.zeros(len(all_dirs)), Y))
o
o train_ids = [all_dirs[i] for i in train_idx]
o val_ids    = [all_dirs[i] for i in val_idx]
o
o print("Campioni:", len(train_ids), "train |",
  len(val_ids), "val")

```

- **Augmentation:** off
- **Scheduler:** off
- **Optimizer:** Adama con lr=1e-4
- **Freezing:** none
- **Auxiliary Head:** off
- **Loss Function / Weighting:** cross-entropy.
- **Main classifier Kernel size:** 1
- **Valutazione risultati:**
 - Migliore Validation Loss ottenuta all'epoca 8:
 - Train Loss: 0.4527 | Val Loss: 0.5234 | Mean IoU: 0.5361
 - ▶ IoU classe 0: 0.5547
 - ▶ IoU classe 1: 0.6503
 - ▶ IoU classe 2: 0.5166
 - ▶ IoU classe 3: 0.6187
 - ▶ IoU classe 4: 0.0004
 - ▶ IoU classe 5: 0.4538
 - ▶ IoU classe 6: 0.3110
 - ▶ IoU classe 7: 0.8308
 - ▶ IoU classe 8: 0.9075
 - Migliore mean IoU ottenuta all'epoca 13:
 - Train Loss: 0.3349 | Val Loss: 0.5666 | Mean IoU: 0.5666
 - ▶ IoU classe 0: 0.5491
 - ▶ IoU classe 1: 0.6573
 - ▶ IoU classe 2: 0.5543
 - ▶ IoU classe 3: 0.5907
 - ▶ IoU classe 4: 0.2466
 - ▶ IoU classe 5: 0.4339

- ► IoU classe 6: 0.3033
- ► IoU classe 7: 0.8390
- ► IoU classe 8: 0.9078

- **Rappresentazione grafica del training:**



- **Massima memoria RAM occupata durante il training:** 2.42GB
- **Log di training:**
`risultati_training\deeplabv3plus\resnet101\dlprn1010005\training_log_dlrn1010005.json`
- **Note / Bug Fix / Annotazioni:**
 - Questo training (aggiunta di splitting proporzionale alle classi tra training e validation) mostra risultati migliori rispetto alla sua versione originale (1010004).

* Esperimento: `deeplabv3plus_resnet101 - dlprn1010006`

- **Architettura:** `deeplabv3plus_resnet101`
- **Backbone Weights:** `best_deeplabv3plus_resnet101_cityscapes_os16.pth`
- **Input size:** 512x512
- **Batch size:** 2
- **Splitting:** on
 - ```
import numpy as np
from PIL import Image
import os
from iterstrat.ml_stratifiers import
 MultilabelStratifiedShuffleSplit
```

```

o import json
o
o # sample male etichettati
o excluded_dirs = {
o "0007", "0009", "0090", "0095", "0101", "0104",
o "0105",
o "0162", "0163", "0284", "0305", "0306", "0307",
o "0308", "0309", "0310", "0311",
o "0351", "0372", "0373", "0376",
o "0498", "0499", "0500", "0501", "0526", "0527",
o "0530", "0531", "0542",
o "0564", "0585", "0586", "0587", "0588", "0589",
o "0590",
o #seguono nuove immagini da escludere
o "0000", "0001", "0052"
o }
o
o
o # Elenco di cartelle numeriche valide
o all_dirs = sorted([
o d for d in os.listdir(DATASET_DIR)
o if os.path.isdir(os.path.join(DATASET_DIR, d)) and
o d.isdigit() and d not in excluded_dirs
o])
o
o
o # -----
o # COSTRUZIONE DELLA MATRICE MULTILABEL
o # -----
o def build_presence_matrix(dir_list, n_classes=9):
o """Return np.array (n_samples, n_classes) with
o booleans indicating class presence"""
o presence = np.zeros((len(dir_list), n_classes),
o dtype=int)
o
o for idx, d in enumerate(dir_list):
o label_path = os.path.join(DATASET_DIR, d,
o "labels.png")
o lbl = np.array(Image.open(label_path))
o uniq = np.unique(lbl)
o presence[idx, uniq] = 1 # segna le
o classi presenti
o return presence
o
o
o Y = build_presence_matrix(all_dirs, NUM_CLASSES)
o
o # -----

```

```

o # MULTILABEL STRATIFIED SPLIT
o # -----
o msss = MultilabelStratifiedShuffleSplit(n_splits=1,
 test_size=0.20, random_state=42)
o train_idx, val_idx =
 next(msss.split(np.zeros(len(all_dirs)), Y))
o
o train_ids = [all_dirs[i] for i in train_idx]
o val_ids = [all_dirs[i] for i in val_idx]
o
o print("Campioni:", len(train_ids), "train |",
 len(val_ids), "val")

```

- **Augmentation:** off
- **Scheduler:** off
- **Optimizer:** Adama con lr=1e-4
- **Freezing:** none
- **Auxiliary Head:** off
- **Loss Function / Weighting:** cross-entropy.
- **Main classifier Kernel size:** 1

- **Kernel size:**

```

class DeepLabHeadV3Plus(nn.Module):
 def __init__(self, in_channels, low_level_channels,
num_classes, aspp_dilate=[4, 8, 12]): #aspp_dilate=[12, 24,
36]
 super(DeepLabHeadV3Plus, self).__init__()
 self.project = nn.Sequential(
 nn.Conv2d(low_level_channels, 48, 1, bias=False),
 nn.BatchNorm2d(48),
 nn.ReLU(inplace=True),
)

 self.aspp = ASPP(in_channels, aspp_dilate)

 self.classifier = nn.Sequential(
 nn.Conv2d(304, 256, 5, padding=2, bias=False),
#kernel size era 3, padding 1
 nn.BatchNorm2d(256),
 nn.ReLU(inplace=True),
 nn.Conv2d(256, num_classes, 1)

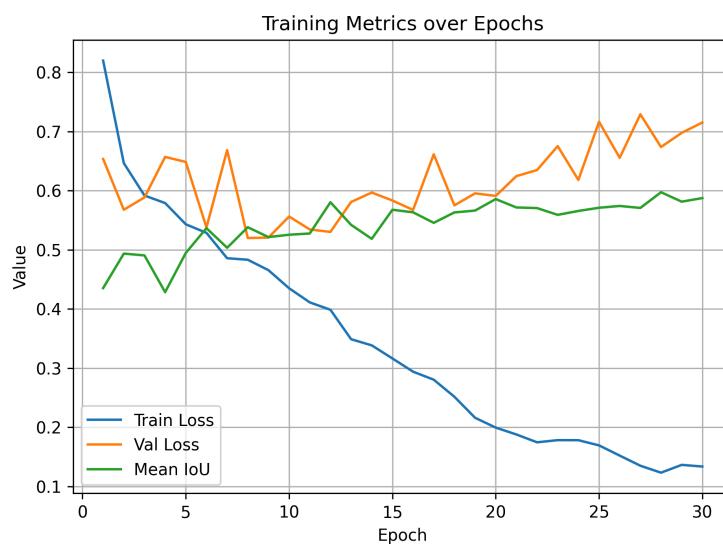
```

```
)
 self._init_weight()
```

- **Valutazione risultati:**

- Migliore Validation Loss ottenuta all'epoca 8:
  - Train Loss: 0.4832 | Val Loss: 0.5199 | Mean IoU: 0.5382
  - ▶ IoU classe 0: 0.5486
  - ▶ IoU classe 1: 0.6654
  - ▶ IoU classe 2: 0.5843
  - ▶ IoU classe 3: 0.5862
  - ▶ IoU classe 4: 0.1720
  - ▶ IoU classe 5: 0.3127
  - ▶ IoU classe 6: 0.2387
  - ▶ IoU classe 7: 0.8411
  - ▶ IoU classe 8: 0.9047
- Migliore mean IoU ottenuta all'epoca 28:
  - Train Loss: 0.1231 | Val Loss: 0.6738 | Mean IoU: 0.5973
  - ▶ IoU classe 0: 0.4991
  - ▶ IoU classe 1: 0.6931
  - ▶ IoU classe 2: 0.5652
  - ▶ IoU classe 3: 0.6404
  - ▶ IoU classe 4: 0.3634
  - ▶ IoU classe 5: 0.4661
  - ▶ IoU classe 6: 0.2992
  - ▶ IoU classe 7: 0.8389
  - ▶ IoU classe 8: 0.9121

- **Rappresentazione grafica del training:**



- **Massima memoria RAM occupata durante il training:** 2.45GB

- **Log di training:**  
`risultati_training\deeplabv3plus\resnet101\dlprn1010006\training_log_dlnr1010006.json`
- **Note / Bug Fix / Annotazioni:**
  - Questo training aggiunge un utilizzo custom del kernel e di ASPP ma peggiora le prestazioni.

## \* Esperimento: `deeplabv3plus_resnet101 - dlprn1010007`

- **Architettura:** `deeplabv3plus_resnet101`
- **Backbone Weights:** `best_deeplabv3plus_resnet101_cityscapes_os16.pth`
- **Input size:** 512x512
- **Batch size:** 2
- **Splitting:** on

```

o import numpy as np
o from PIL import Image
o import os
o from iterstrat.ml_stratifiers import
 MultilabelStratifiedShuffleSplit
o import json
o
o # sample male etichettati
o excluded_dirs = {
o "0007", "0009", "0090", "0095", "0101", "0104",
o "0105",
o "0162", "0163", "0284", "0305", "0306", "0307",
o "0308", "0309", "0310", "0311",
o "0351", "0372", "0373", "0376",
o "0498", "0499", "0500", "0501", "0526", "0527",
o "0530", "0531", "0542",
o "0564", "0585", "0586", "0587", "0588", "0589",
o "0590",
o #seguono nuove immagini da escludere
o "0000", "0001", "0052"
o }
o
o # Elenco di cartelle numeriche valide
o all_dirs = sorted([
o d for d in os.listdir(DATASET_DIR)

```

```

o if os.path.isdir(os.path.join(DATASET_DIR, d)) and
o d.isdigit() and d not in excluded_dirs
o])
o
o # -----
o # COSTRUZIONE DELLA MATRICE MULTILABEL
o # -----
o def build_presence_matrix(dir_list, n_classes=9):
o """Return np.array (n_samples, n_classes) with
o booleans indicating class presence"""
o presence = np.zeros((len(dir_list), n_classes),
o dtype=int)
o
o for idx, d in enumerate(dir_list):
o label_path = os.path.join(DATASET_DIR, d,
o "labels.png")
o lbl = np.array(Image.open(label_path))
o uniq = np.unique(lbl)
o presence[idx, uniq] = 1 # segna le
o classi presenti
o
o return presence
o
o
o Y = build_presence_matrix(all_dirs, NUM_CLASSES)
o
o # -----
o # MULTILABEL STRATIFIED SPLIT
o # -----
o msss = MultilabelStratifiedShuffleSplit(n_splits=1,
o test_size=0.20, random_state=42)
o train_idx, val_idx =
o next(msss.split(np.zeros(len(all_dirs)), Y))
o
o train_ids = [all_dirs[i] for i in train_idx]
o val_ids = [all_dirs[i] for i in val_idx]
o
o print("Campioni:", len(train_ids), "train |",
o len(val_ids), "val")

```

- **Augmentation:** off
- **Scheduler:** off
- **Optimizer:** Adama con lr=1e-4
- **Freezing:** none

- **Auxiliary Head:** off
- **Loss Function / Weighting:** cross-entropy.
- **Main classifier Kernel size:** 1
- **Kernel size:**

```
def __init__(self, in_channels, low_level_channels, num_classes,
aspp_dilate=[3, 7, 3]): #aspp_dilate=[12, 24, 36]
 super(DeepLabHeadV3Plus, self).__init__()
 self.project = nn.Sequential(
 nn.Conv2d(low_level_channels, 48, 1, bias=False),
 nn.BatchNorm2d(48),
 nn.ReLU(inplace=True),
)

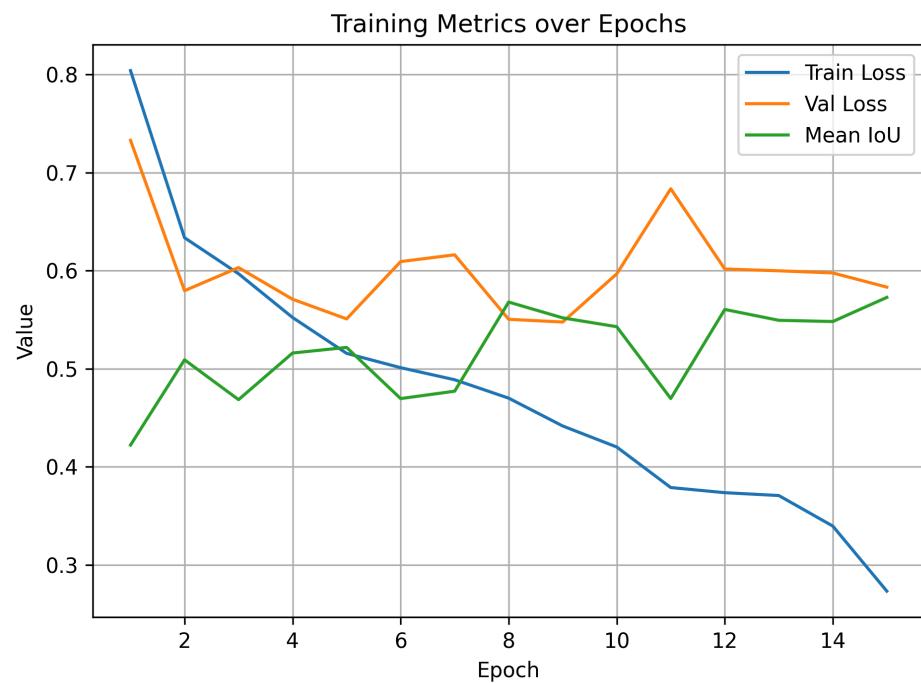
 self.aspp = ASPP(in_channels, aspp_dilate)

 self.classifier = nn.Sequential(
 nn.Conv2d(304, 256, 5, padding=2, bias=False),
#kernel size era 3, padding 1
nn.BatchNorm2d(256),
 nn.ReLU(inplace=True),
 nn.Dropout(0.1),
 nn.Conv2d(256, num_classes, 1)
)
 self.__init_weight()
```

- **Valutazione risultati:**

- Migliore Validation Loss ottenuta all'epoca 9:
  -  Train Loss: 0.4417 | Val Loss: 0.5476 | Mean IoU: 0.5518
  - ▶ IoU classe 0: 0.5367
  - ▶ IoU classe 1: 0.6628
  - ▶ IoU classe 2: 0.5508
  - ▶ IoU classe 3: 0.6003
  - ▶ IoU classe 4: 0.2714
  - ▶ IoU classe 5: 0.3063
  - ▶ IoU classe 6: 0.2910
  - ▶ IoU classe 7: 0.8303
  - ▶ IoU classe 8: 0.9017

- Migliore mean IoU ottenuta all'epoca 15:
  - Train Loss: 0.2733 | Val Loss: 0.5832 | Mean IoU: 0.5727
  - ▶ IoU classe 0: 0.5586
  - ▶ IoU classe 1: 0.7004
  - ▶ IoU classe 2: 0.5607
  - ▶ IoU classe 3: 0.6475
  - ▶ IoU classe 4: 0.3450
  - ▶ IoU classe 5: 0.2963
  - ▶ IoU classe 6: 0.2879
  - ▶ IoU classe 7: 0.8315
  - ▶ IoU classe 8: 0.9122
- Rappresentazione grafica del training:



- Massima memoria RAM occupata durante il training: 2.45GB
- Log di training:  
`risultati_training\deeplabv3plus\resnet101\dlprn1010007\training_log_dlrn1010007.json`
- Note / Bug Fix / Annotazioni:

\* Esperimento: `deeplabv3plus_resnet101 - dlprn1010008`

- Architettura: `deeplabv3plus_resnet101`
- Backbone Weights: `best_deeplabv3plus_resnet101_cityscapes_os16.pth`
- Input size: 512x512

- **Batch size:** 2

- **Splitting:** on

```

o import numpy as np
o from PIL import Image
o import os
o from iterstrat.ml_stratifiers import
 MultilabelStratifiedShuffleSplit
o import json
o
o # sample male etichettati
o excluded_dirs = {
o "0007", "0009", "0090", "0095", "0101", "0104",
o "0105",
o "0162", "0163", "0284", "0305", "0306", "0307",
o "0308", "0309", "0310", "0311",
o "0351", "0372", "0373", "0376",
o "0498", "0499", "0500", "0501", "0526", "0527",
o "0530", "0531", "0542",
o "0564", "0585", "0586", "0587", "0588", "0589",
o "0590",
o #seguono nuovo immagini da escludere
o "0000", "0001", "0052"
o }
o
o # Elenco di cartelle numeriche valide
o all_dirs = sorted([
o d for d in os.listdir(DATASET_DIR)
o if os.path.isdir(os.path.join(DATASET_DIR, d)) and
o d.isdigit() and d not in excluded_dirs
o])
o
o # -----
o # COSTRUZIONE DELLA MATRICE MULTILABEL
o # -----
o def build_presence_matrix(dir_list, n_classes=9):
o """Return np.array (n_samples, n_classes) with
o booleans indicating class presence"""
o presence = np.zeros((len(dir_list), n_classes),
o dtype=int)
o
o for idx, d in enumerate(dir_list):
o label_path = os.path.join(DATASET_DIR, d,
o "labels.png")
o lbl = np.array(Image.open(label_path))

```

```

o uniq = np.unique(lbl)
o presence[idx, uniq] = 1 # segna le
o classi presenti
o return presence
o
o
o Y = build_presence_matrix(all_dirs, NUM_CLASSES)
o
o
o # -----
o # MULTILABEL STRATIFIED SPLIT
o # -----
o msss = MultilabelStratifiedShuffleSplit(n_splits=1,
o test_size=0.20, random_state=42)
o train_idx, val_idx =
o next(msss.split(np.zeros(len(all_dirs)), Y))
o
o train_ids = [all_dirs[i] for i in train_idx]
o val_ids = [all_dirs[i] for i in val_idx]
o
o print("Campioni:", len(train_ids), "train |",
o len(val_ids), "val")

```

- **Augmentation:** off
- **Scheduler:** off
- **Optimizer:** Adama con lr=1e-4
- **Freezing:** none
- **Auxiliary Head:** off
- **Loss Function / Weighting:** cross-entropy.
- **Main classifier Kernel size:** 1
- **Kernel size:**

```

def __init__(self, in_channels, low_level_channels, num_classes,
aspp_dilate=[1, 3, 5]):

 super(DeepLabHeadV3Plus, self).__init__()
 self.project = nn.Sequential(
 nn.Conv2d(low_level_channels, 48, 1, bias=False),
 nn.BatchNorm2d(48),
 nn.ReLU(inplace=True),
)

```

```

 self.aspp = ASPP(in_channels, aspp_dilate)

 self.classifier = nn.Sequential(
 nn.Conv2d(304, 256, 5, padding=2, bias=False),
#kernel size era 3, padding 1
 nn.BatchNorm2d(256),
 nn.ReLU(inplace=True),
 nn.Conv2d(256, 256, 5, padding=2, bias=False),
 nn.ReLU(inplace=True),
 nn.Dropout(0.1),
 nn.Conv2d(256, num_classes, 1)
)
 self._init_weight()

```

- **Valutazione risultati:**

Migliore Validation Loss ottenuta all'epoca 17:

 Train Loss: 0.2485 | Val Loss: 0.5676 | Mean IoU: 0.5315

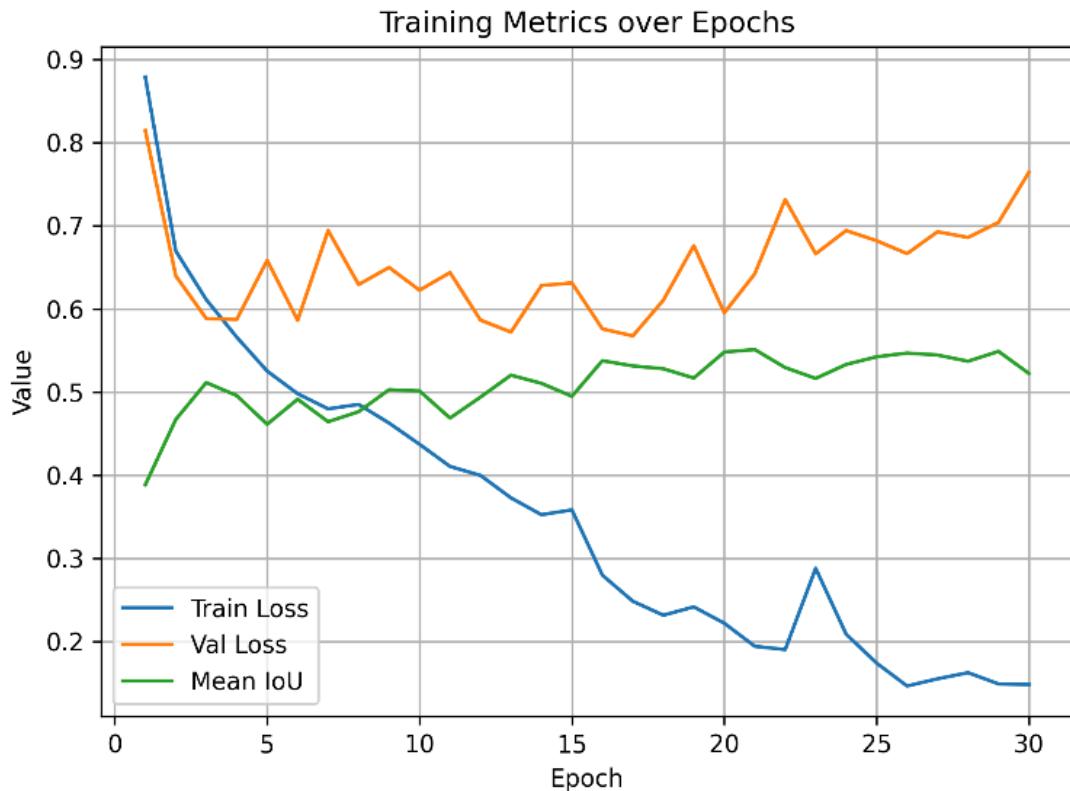
- ▶ IoU classe 0: 0.6081
- ▶ IoU classe 1: 0.6245
- ▶ IoU classe 2: 0.5524
- ▶ IoU classe 3: 0.6435
- ▶ IoU classe 4: 0.0000
- ▶ IoU classe 5: 0.4427
- ▶ IoU classe 6: 0.2647
- ▶ IoU classe 7: 0.8195
- ▶ IoU classe 8: 0.9046

Migliore mean IoU ottenuta all'epoca 21:

 Train Loss: 0.1944 | Val Loss: 0.6422 | Mean IoU: 0.5512

- ▶ IoU classe 0: 0.6414
- ▶ IoU classe 1: 0.6371
- ▶ IoU classe 2: 0.5762
- ▶ IoU classe 3: 0.6136
- ▶ IoU classe 4: 0.0000
- ▶ IoU classe 5: 0.4852
- ▶ IoU classe 6: 0.3657
- ▶ IoU classe 7: 0.8310
- ▶ IoU classe 8: 0.9011

- **Rappresentazione grafica del training:**



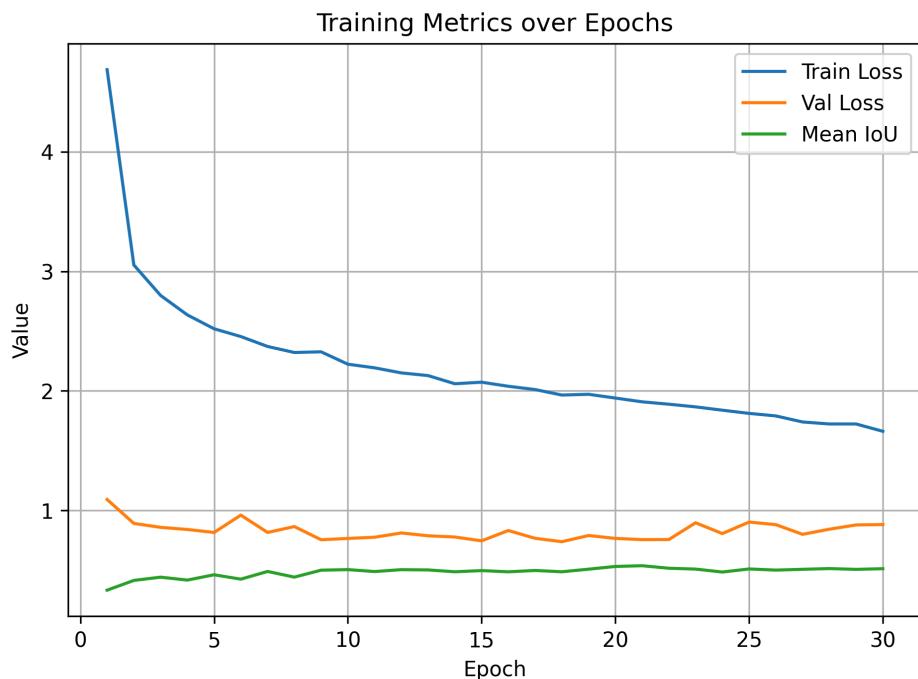
- **Massima memoria RAM occupata durante il training:** 2.48GB
- **Log di training:**  
`risultati_training\deeplabv3plus\resnet101\dlprn1010007\training_log_dlprn1010008.json`
- **Note / Bug Fix / Annotazioni:**

## BiSeNetV2

\* Esperimento: **BiSeNetV2 - bsnv20001**

- **Architettura:** BiSeNetV2
- **Backbone Weights:** model\_final\_rugd.pth
- **Input size:** 256x256
- **Batch size:** 2
- **Augmentation:** off
- **Scheduler:** off

- **Optimizer:** Adama con lr=1e-4
- **Freezing:** none
- **Auxiliary Head:** off
- **Loss Function / Weighting:** cross-entropy+ class weights euristici [1.0, 1.0, 1.0, 1.0, 5.0, 4.0, 3.5, 0.8, 0.8].
- **Main classifier Kernel size:** 1
- **Valutazione risultati:**
  - Migliore Validation Loss ottenuta all'epoca 18:
    - Train Loss: 1.9656 | Val Loss: 0.7396 | Mean IoU: 0.4882
    - ▶ IoU classe 0: 0.5773
    - ▶ IoU classe 1: 0.5773
    - ▶ IoU classe 2: 0.4750
    - ▶ IoU classe 3: 0.4569
    - ▶ IoU classe 4: 0.1991
    - ▶ IoU classe 5: 0.4224
    - ▶ IoU classe 6: 0.1655
    - ▶ IoU classe 7: 0.7791
    - ▶ IoU classe 8: 0.8919
  - Migliore mean IoU ottenuta all'epoca 21:
    - Train Loss: 1.9089 | Val Loss: 0.7570 | Mean IoU: 0.5388
    - ▶ IoU classe 0: 0.5848
    - ▶ IoU classe 1: 0.4562
    - ▶ IoU classe 2: 0.5351
    - ▶ IoU classe 3: 0.5078
    - ▶ IoU classe 4: 0.3807
    - ▶ IoU classe 5: 0.5554
    - ▶ IoU classe 6: 0.1715
    - ▶ IoU classe 7: 0.8067
    - ▶ IoU classe 8: 0.8966
- **Rappresentazione grafica del training:**



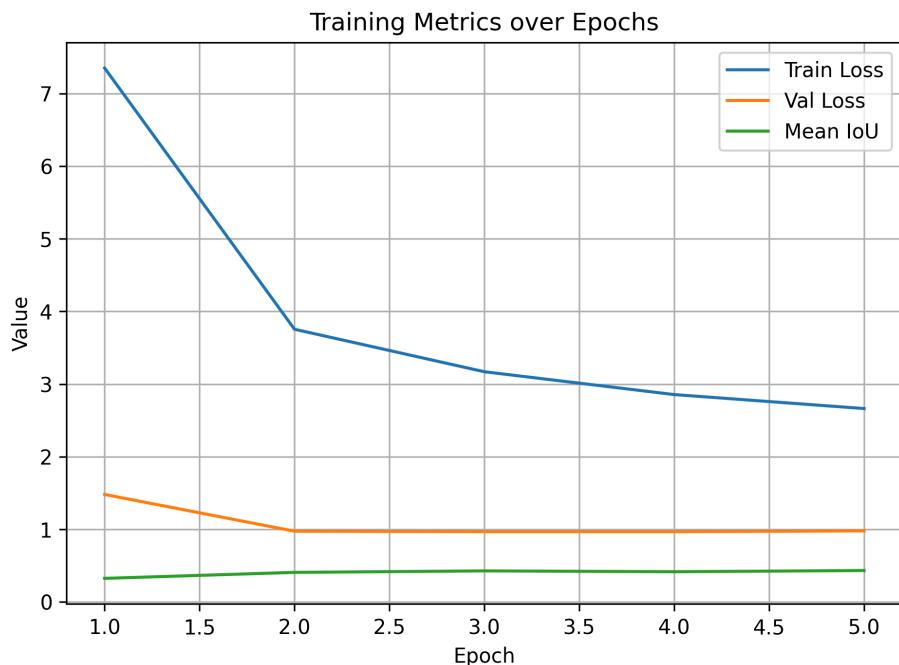
- **Massima memoria RAM occupata durante il training:** 1.96GB
- **Log di training:**  
`risultati_training\deeplabv3plus\resnet101\bsnv20001\training_log_bsnv20001.json`
- **Note / Bug Fix / Annotazioni:**

○ -

## \* Esperimento: BiSeNetV2 - bsnv20002

- **Architettura:** BiSeNetV2
- **Backbone Weights:** model\_final\_rellis3d.pth
- **Input size:** 256x256
- **Batch size:** 2
- **Augmentation:** off
- **Scheduler:** off
- **Optimizer:** Adama con lr=1e-4
- **Freezing:** none
- **Auxiliary Head:** off

- **Loss Function / Weighting:** cross-entropy+ class weights euristici [1.0, 1.0, 1.0, 1.0, 5.0, 4.0, 3.5, 0.8, 0.8].
- **Main classifier Kernel size: 1**
- **Valutazione risultati:**
  - Migliore Validation Loss ottenuta all'epoca 4:
    - Train Loss: 2.8555 | Val Loss: 0.9681 | Mean IoU: 0.4166
      - ▶ IoU classe 0: 0.4735
      - ▶ IoU classe 1: 0.4346
      - ▶ IoU classe 2: 0.3779
      - ▶ IoU classe 3: 0.4268
      - ▶ IoU classe 4: 0.0000
      - ▶ IoU classe 5: 0.3164
      - ▶ IoU classe 6: 0.1450
      - ▶ IoU classe 7: 0.7682
      - ▶ IoU classe 8: 0.8637
  - Migliore mean IoU ottenuta all'epoca 5:
    - Train Loss: 2.6632 | Val Loss: 0.9786 | Mean IoU: 0.4331
      - ▶ IoU classe 0: 0.4779
      - ▶ IoU classe 1: 0.4879
      - ▶ IoU classe 2: 0.4225
      - ▶ IoU classe 3: 0.3771
      - ▶ IoU classe 4: 0.0000
      - ▶ IoU classe 5: 0.3911
      - ▶ IoU classe 6: 0.1437
      - ▶ IoU classe 7: 0.7777
      - ▶ IoU classe 8: 0.8644
- **Rappresentazione grafica del training:**



- **Massima memoria RAM occupata durante il training:** 1.96GB
- **Log di training:**  
`risultati_training\deeplabv3plus\resnet101\bsnv20002\training_log_bsnv20002.json`
- **Note / Bug Fix / Annotazioni:**

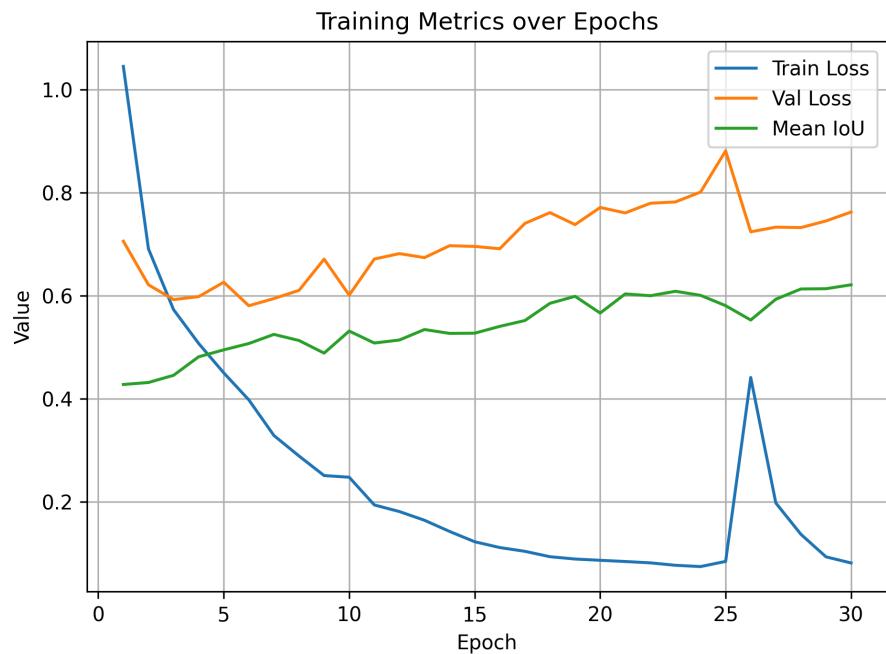
○ -

## deeplabv3\_resnet50

\* Esperimento: `deeplabv3_resnet50 - dlrn500000`

- **Architettura:** `deeplabv3plus_resnet50`
- **Backbone Weights:** DeepLabV3\_ResNet50\_Weights.DEFAULT
- **Input size:** 256x256
- **Batch size:** 4
- **Augmentation:** off
- **Scheduler:** off
- **Optimizer:** Adama con lr=1e-4

- **Freezing:** none
- **Auxiliary Head:** off
- **Loss Function / Weighting:** cross-entropy
- **Main classifier Kernel size:** 1
- **Valutazione risultati:**
  - Migliore Validation Loss ottenuta all'epoca 6:
    - Train Loss: 0.3976 | Val Loss: 0.5804 | Mean IoU: 0.5070
      - ▶ IoU classe 0: 0.4787
      - ▶ IoU classe 1: 0.5573
      - ▶ IoU classe 2: 0.6179
      - ▶ IoU classe 3: 0.5495
      - ▶ IoU classe 4: 0.0000
      - ▶ IoU classe 5: 0.4713
      - ▶ IoU classe 6: 0.1508
      - ▶ IoU classe 7: 0.8290
      - ▶ IoU classe 8: 0.8799
  - Migliore **mean IoU** ottenuta all'epoca 30:
    - Train Loss: 0.0809 | Val Loss: 0.7622 | Mean IoU: 0.6210
      - ▶ IoU classe 0: 0.5350
      - ▶ IoU classe 1: 0.6204
      - ▶ IoU classe 2: 0.6361
      - ▶ IoU classe 3: 0.5720
      - ▶ IoU classe 4: 0.6263
      - ▶ IoU classe 5: 0.6018
      - ▶ IoU classe 6: 0.1782
      - ▶ IoU classe 7: 0.8382
      - ▶ IoU classe 8: 0.8949
- **Rappresentazione grafica del training:**

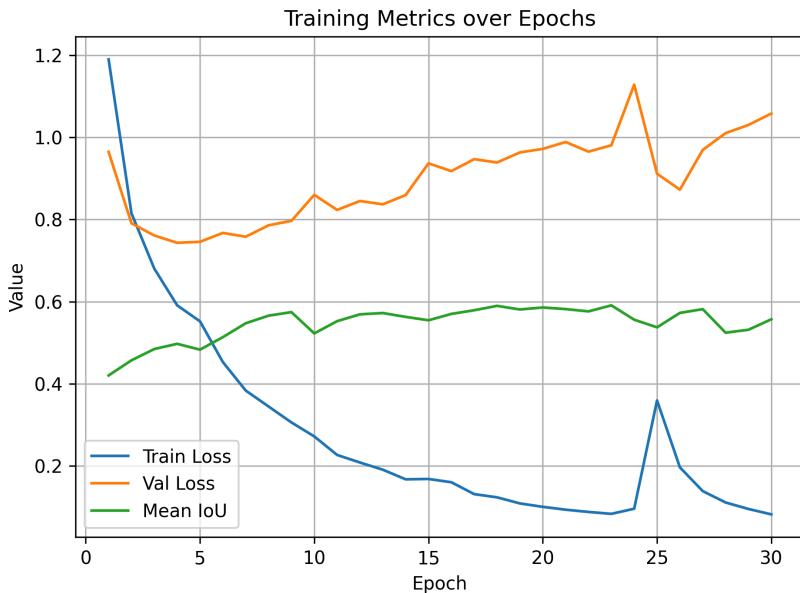


- **Massima memoria RAM occupata durante il training:** 1.54GB
- **Log di training:**  
`risultati_training\deeplabv3\resnet50\dlnrn50000\training_log_dlnrn50000.json`
- **Note / Bug Fix / Annotazioni:**
  - —

### \* Esperimento: `deeplabv3_resnet50 - dlnrn500001`

- **Architettura:** `deeplabv3plus_resnet50`
- **Backbone Weights:** DeepLabV3\_ResNet50\_Weights.DEFAULT
- **Input size:** 256x256
- **Batch size:** 4
- **Augmentation:** off
- **Scheduler:** off
- **Optimizer:** Adama con lr=1e-4
- **Freezing:** none

- **Auxiliary Head:** off
- **Loss Function / Weighting:** cross-entropy + class weights euristici [1.0, 1.0, 1.0, 1.0, 5.0, 4.0, 3.5, 0.8, 0.8].
- **Main classifier Kernel size:** 1
- **Valutazione risultati:**
  - Migliore Validation Loss ottenuta all'epoca 4:
    - Train Loss: 0.5912 | Val Loss: 0.7434 | Mean IoU: 0.4975
    - ▶ IoU classe 0: 0.4978
    - ▶ IoU classe 1: 0.5695
    - ▶ IoU classe 2: 0.5240
    - ▶ IoU classe 3: 0.5543
    - ▶ IoU classe 4: 0.0000
    - ▶ IoU classe 5: 0.4847
    - ▶ IoU classe 6: 0.1588
    - ▶ IoU classe 7: 0.8133
    - ▶ IoU classe 8: 0.8751
  - Migliore **mean IoU** ottenuta all'epoca 23:
    - Train Loss: 0.0836 | Val Loss: 0.9807 | Mean IoU: 0.5911
    - ▶ IoU classe 0: 0.5239
    - ▶ IoU classe 1: 0.6022
    - ▶ IoU classe 2: 0.6362
    - ▶ IoU classe 3: 0.5781
    - ▶ IoU classe 4: 0.4076
    - ▶ IoU classe 5: 0.5827
    - ▶ IoU classe 6: 0.1911
    - ▶ IoU classe 7: 0.8366
    - ▶ IoU classe 8: 0.8943
- **Rappresentazione grafica del training:**



- **Massima memoria RAM occupata durante il training:** 1.54GB
- **Log di training:**  
`risultati_training\deeplabv3\resnet50\dlnrn500001\training_log_dlnrn500001.json`
- **Note / Bug Fix / Annotazioni:**
  - —

## \* Esperimento: deeplabv3\_resnet50 - dlnrn500002

- **Architettura:** deeplabv3plus\_resnet50
- **Backbone Weights:** DeepLabV3\_ResNet50\_Weights.DEFAULT
- **Input size:** 256x256
- **Batch size:** 4
- **Augmentation :** on

```

○ # Augmentazione per immagini con classi rare
○ rare_aug = A.Compose([
○ A.HorizontalFlip(p=0.5),
○ A.RandomBrightnessContrast(p=0.4),
○ A.Resize(256, 256),
○ A.Normalize(mean=(0.485, 0.456, 0.406),
○ std=(0.229, 0.224, 0.225)),
○ ToTensorV2()
○])

```

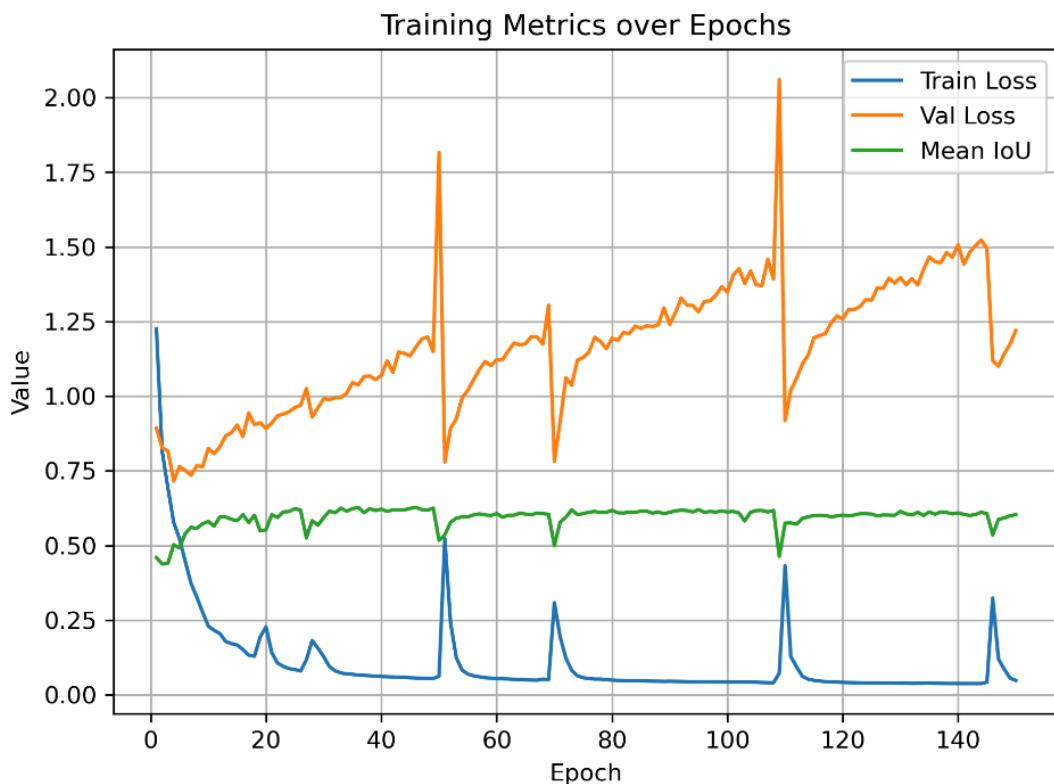
```

o
o # Augmentazione di base
o base_aug = A.Compose([
o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])
o
o
o
o train_loader = DataLoader(train_set,
o batch_size=BATCH_SIZE, shuffle=True, drop_last=True)
o def __init__(self, root_dir, is_val = False,
o sample_ids=None, allow_augmentation = False):
o self.samples = []
o self.is_val = is_val
o self.sample_ids = sample_ids
o self.allow_augmentation = allow_augmentation
o self._load_and_augment(root_dir)
o
o
o def _load_and_augment(self, root_dir):
o sample_dirs = []
o for id in self.sample_ids:
o sample_dirs.append(os.path.join(root_dir, id))
o
o
o for dir_path in sample_dirs:
o img_path = os.path.join(dir_path, 'rgb.jpg')
o label_path = os.path.join(dir_path, 'labels.png')
o
o label_img = Image.open(label_path)
o is_rare =
o SegmentationDataset.contains_rare_classes(label_img,
o RARE_CLASSES)
o
o if not self.allow_augmentation:
o is_rare = False
o
o # Sempre aggiungi l'originale
o self.samples.append((img_path, label_path,
o is_rare))
o
o # Duplicazione per oversampling (solo se contiene
o classi rare)
o if is_rare and not self.is_val:
o for _ in range(DUPLICATE_FACTOR - 1):

```

```
o self.samples.append((img_path,
| label_path, is_rare))
```

- **Scheduler:** off
- **Optimizer:** Adama con lr=1e-4
- **Freezing:** none
- **Auxiliary Head:** off
- **Loss Function / Weighting:** cross-entropy + class weights euristici [1.0, 1.0, 1.0, 1.0, 5.0, 4.0, 3.5, 0.8, 0.8].
- **Main classifier Kernel size:** 1
- **Valutazione risultati:**
  - Migliore Validation Loss ottenuta all'epoca 4:
    - Train Loss: 0.5772 | Val Loss: 0.7160 | Mean IoU: 0.5025
      - ► IoU classe 0: 0.5101
      - ► IoU classe 1: 0.5496
      - ► IoU classe 2: 0.5533
      - ► IoU classe 3: 0.5210
      - ► IoU classe 4: 0.0
      - ► IoU classe 5: 0.5110
      - ► IoU classe 6: 0.1892
      - ► IoU classe 7: 0.8196
      - ► IoU classe 8: 0.8759
  - Migliore **mean IoU** ottenuta all'epoca 46:
    - Train Loss: 0.0560 | Val Loss: 1.1647 | Mean IoU: 0.6277
      - ► IoU classe 0: 0.5581
      - ► IoU classe 1: 0.6294
      - ► IoU classe 2: 0.6459
      - ► IoU classe 3: 0.5936
      - ► IoU classe 4: 0.5776
      - ► IoU classe 5: 0.6460
      - ► IoU classe 6: 0.1898
      - ► IoU classe 7: 0.8455
      - ► IoU classe 8: 0.8937
- **Rappresentazione grafica del training:**



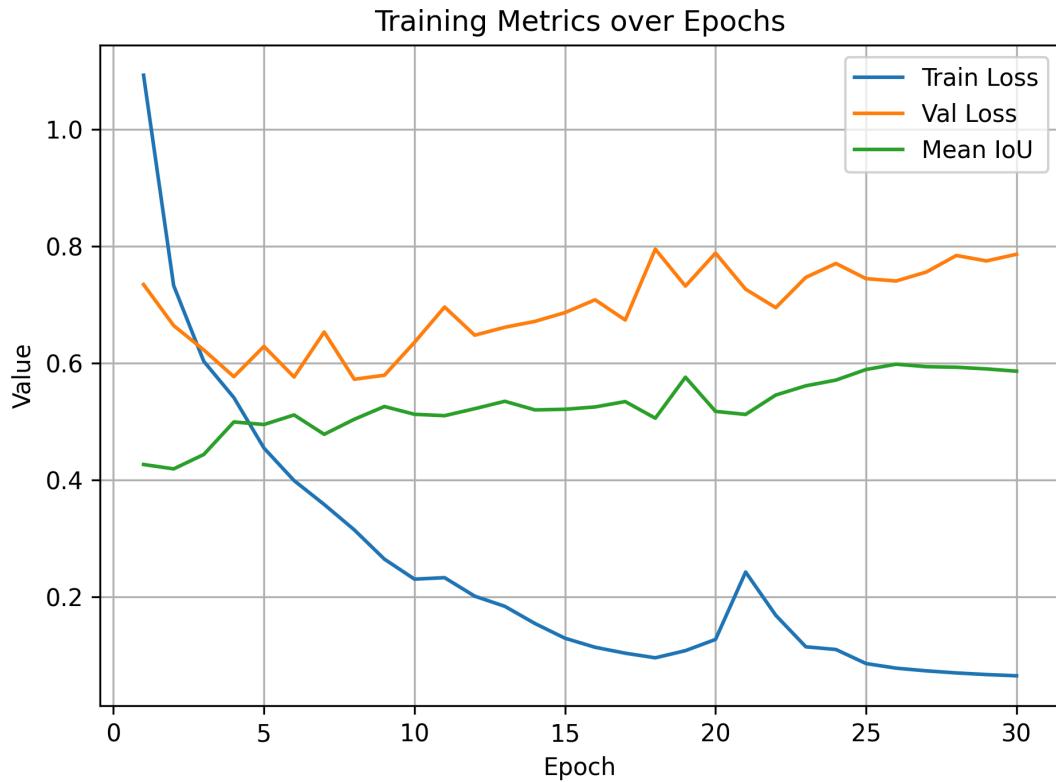
- **Massima memoria RAM occupata durante il training:** —GB
- **Log di training:**  
`risultati_training\deeplabv3\resnet50\dln500001\training_log_dln500002.json`
- **Note / Bug Fix / Annotazioni:**
  - grafico con picchi anomali ogni 30 epochhe

## deeplabv3\_resnet101

\* Esperimento: `deeplabv3_resnet101 - dln101000`

- **Architettura:** `deeplabv3plus_resnet101`
- **Backbone Weights:** DeepLabV3\_ResNet101\_Weights.DEFAULT
- **Input size:** 256x256
- **Batch size:** 4
- **Augmentation:** off

- **Scheduler:** off
- **Optimizer:** Adama con lr=1e-4
- **Freezing:** none
- **Auxiliary Head:** off
- **Loss Function / Weighting:** cross-entropy
- **Main classifier Kernel size:** 1
- **Valutazione risultati:**
  - Migliore Validation Loss ottenuta all'epoca 7:
    - Train Loss: 0.3993 | Val Loss: 0.5766 | Mean IoU: 0.5117
    - ▶ IoU classe 0: 0.5169
    - ▶ IoU classe 1: 0.5172
    - ▶ IoU classe 2: 0.6249
    - ▶ IoU classe 3: 0.5701
    - ▶ IoU classe 4: 0.0000
    - ▶ IoU classe 5: 0.5097
    - ▶ IoU classe 6: 0.1662
    - ▶ IoU classe 7: 0.8303
    - ▶ IoU classe 8: 0.8747
  - Migliore **mean IoU** ottenuta all'epoca 26:
    - Train Loss: 0.0782 | Val Loss: 0.7409 | Mean IoU: 0.5985
    - ▶ IoU classe 0: 0.5401
    - ▶ IoU classe 1: 0.6034
    - ▶ IoU classe 2: 0.6366
    - ▶ IoU classe 3: 0.5879
    - ▶ IoU classe 4: 0.3909
    - ▶ IoU classe 5: 0.6283
    - ▶ IoU classe 6: 0.2004
    - ▶ IoU classe 7: 0.8454
    - ▶ IoU classe 8: 0.8949
- **Rappresentazione grafica del training:**

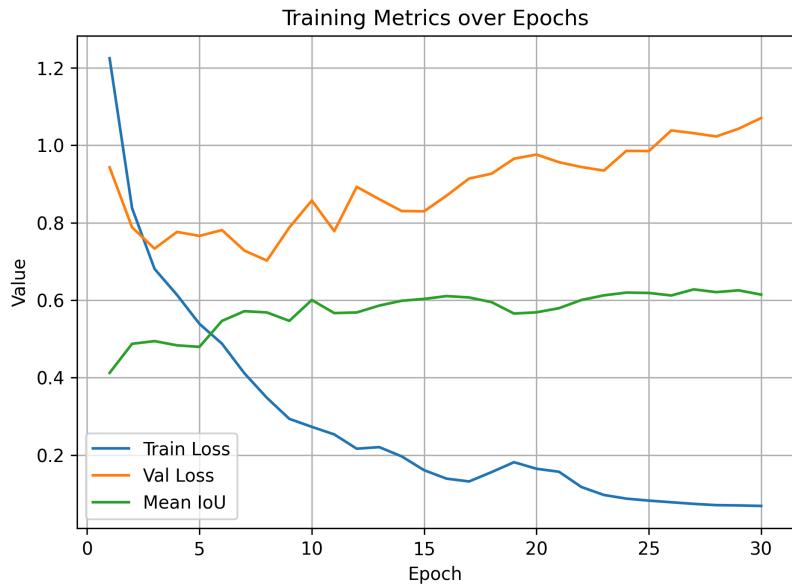


- **Massima memoria RAM occupata durante il training:** 2.62GB
- **Log di training:**  
`risultati_training\deeplabv3\resnet101\dlnr101000\training_log_dlnr101001.json`
- **Note / Bug Fix / Annotazioni:**
  - —

\* Esperimento: `deeplabv3_resnet101 - dlnr101001`

- **Architettura:** `deeplabv3plus_resnet101`
- **Backbone Weights:** DeepLabV3\_ResNet101\_Weights.DEFAULT
- **Input size:** 256x256
- **Batch size:** 4
- **Augmentation:** off
- **Scheduler:** off
- **Optimizer:** Adama con lr=1e-4

- **Freezing:** none
- **Auxiliary Head:** off
- **Loss Function / Weighting:** cross-entropy + class weights euristici [1.0, 1.0, 1.0, 1.0, 5.0, 4.0, 3.5, 0.8, 0.8].
- **Main classifier Kernel size:** 1
- **Valutazione risultati:**
  - Migliore Validation Loss ottenuta all'epoca 8:
    - Train Loss: 0.3480 | Val Loss: 0.7024 | Mean IoU: 0.5685
    - ▶ IoU classe 0: 0.5130
    - ▶ IoU classe 1: 0.6119
    - ▶ IoU classe 2: 0.5904
    - ▶ IoU classe 3: 0.5739
    - ▶ IoU classe 4: 0.3759
    - ▶ IoU classe 5: 0.5182
    - ▶ IoU classe 6: 0.1781
    - ▶ IoU classe 7: 0.8242
    - ▶ IoU classe 8: 0.8751
  - Migliore **mean IoU** ottenuta all'epoca 27:
    - Train Loss: 0.0740 | Val Loss: 1.0314 | Mean IoU: 0.6280
    - ▶ IoU classe 0: 0.5373
    - ▶ IoU classe 1: 0.6072
    - ▶ IoU classe 2: 0.6516
    - ▶ IoU classe 3: 0.5739
    - ▶ IoU classe 4: 0.6328
    - ▶ IoU classe 5: 0.6454
    - ▶ IoU classe 6: 0.1861
    - ▶ IoU classe 7: 0.8358
    - ▶ IoU classe 8: 0.8912
- **Rappresentazione grafica del training:**



- **Massima memoria RAM occupata durante il training:** 2.62GB
- **Log di training:**  
`risultati_training\deeplabv3\resnet101\dlnr1010001\training_log_dlnr1010001.json`
- **Note / Bug Fix / Annotazioni:**
  - —

## \* Esperimento: deeplabv3\_resnet101 - dlnr101002

- **Architettura:** `deeplabv3plus_resnet101`
- **Backbone Weights:** DeepLabV3\_ResNet101\_Weights.DEFAULT
- **Input size:** 256x256
- **Batch size:** 4
- **Augmentation:** on

```

○ # Augmentazione per immagini con classi rare
○ rare_aug = A.Compose([
○ A.HorizontalFlip(p=0.5),
○ A.RandomBrightnessContrast(p=0.4),
○ #A.Rotate(limit=20, p=0.3), potenzialmente inutile
○ A.Resize(256, 256),
○ A.Normalize(mean=(0.485, 0.456, 0.406),
○ std=(0.229, 0.224, 0.225)),
```

```

o ToTensorV2()
o])
o
o # Augmentazione di base
o base_aug = A.Compose([
o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])
o

-
def _load_and_augment(self, root_dir):
 sample_dirs = []
 for id in self.sample_ids:
 sample_dirs.append(os.path.join(root_dir, id))

 for dir_path in sample_dirs:
 img_path = os.path.join(dir_path, 'rgb.jpg')
 label_path = os.path.join(dir_path, 'labels.png')

 label_img = Image.open(label_path)
 is_rare =
SegmentationDataset.contains_rare_classes(label_img,
RARE_CLASSES)

 if not self.allow_augmentation:
 is_rare = False

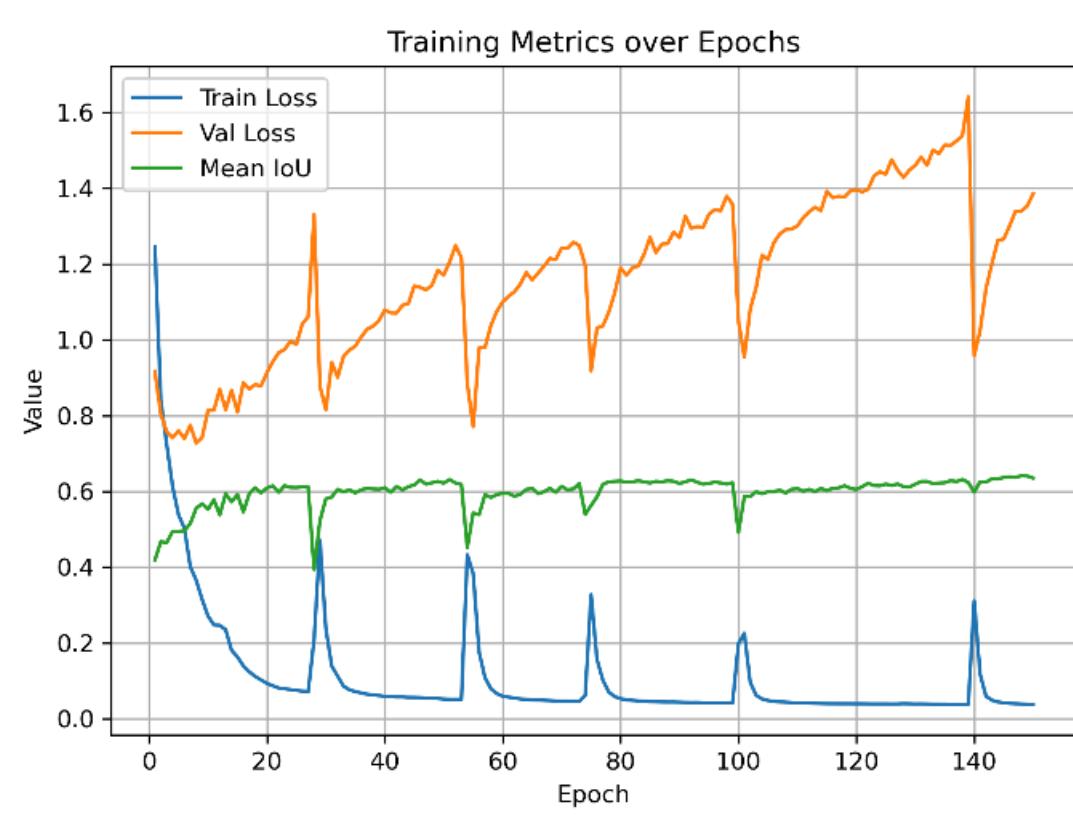
 # Sempre aggiungi l'originale
 self.samples.append((img_path, label_path,
is_rare))

 # Duplicazione per oversampling (solo se contiene
classi rare)
 if is_rare and not self.is_val:
 for _ in range(DUPLICATE_FACTOR - 1):
 self.samples.append((img_path,
label_path, is_rare))

```

- **Scheduler:** off
- **Optimizer:** Adama con lr=1e-4
- **Freezing:** none

- **Auxiliary Head:** off
- **Loss Function / Weighting:** cross-entropy + class weights euristici [1.0, 1.0, 1.0, 1.0, 5.0, 4.0, 3.5, 0.8, 0.8].
- **Main classifier Kernel size:** 1
- **Valutazione risultati:**
  - Migliore Validation Loss ottenuta all'epoca 4:
    - Train Loss: 0.3602 | Val Loss: 0.7229 | Mean IoU: 0.5917
      - IoU classe 0: 0.4873
      - IoU classe 1: 0.5685
      - IoU classe 2: 0.5609
      - IoU classe 3: 0.5867
      - IoU classe 4: 0.5617
      - IoU classe 5: 0.5661
      - IoU classe 6: 0.1902
      - IoU classe 7: 0.8245
      - IoU classe 8: 0.8751
  - Migliore mean IoU ottenuta all'epoca 18:
    - Train Loss: 0.0781 | Val Loss: 0.9812 | Mean IoU: 0.6285
      - IoU classe 0: 0.5387
      - IoU classe 1: 0.6149
      - IoU classe 2: 0.6248
      - IoU classe 3: 0.5763
      - IoU classe 4: 0.6435
      - IoU classe 5: 0.6451
      - IoU classe 6: 0.1964
      - IoU classe 7: 0.8354
      - IoU classe 8: 0.8913
- **Rappresentazione grafica del training:**



- **Massima memoria RAM occupata durante il training:** 2.61GB
- **Log di training:**  
`risultati_training\deeplabv3\resnet101\dlnr101001\training_log_dlnr101002.json`
- **Note / Bug Fix / Annotazioni:**
  - picchi anomali ogni 25 epoche circa

### \* Esperimento: `deeplabv3_resnet101 - dlnr101003`

- **Architettura:** `deeplabv3plus_resnet101`
- **Backbone Weights:** DeepLabV3\_ResNet101\_Weights.DEFAULT
- **Input size:** 256x256
- **Test size:** 512x512

```
Augmentazione di base
val_aug = A.Compose([
 A.Resize(512, 512),
 A.Normalize(mean=(0.485, 0.456, 0.406),
 std=(0.229, 0.224, 0.225)),
```

```
○ ToTensorV2()
○])
○
○
```

- **Batch size:** 4

- **Augmentation:** on

```
○ # Augmentazione per immagini con classi rare
○ rare_aug = A.Compose([
○ A.HorizontalFlip(p=0.5),
○ A.RandomBrightnessContrast(p=0.4),
○ #A.Rotate(limit=20, p=0.3), potenzialmente inutile
○ A.Resize(256, 256),
○ A.Normalize(mean=(0.485, 0.456, 0.406),
○ std=(0.229, 0.224, 0.225)),
○ ToTensorV2()
○])
○
○ # Augmentazione di base
○ base_aug = A.Compose([
○ A.Resize(256, 256),
○ A.Normalize(mean=(0.485, 0.456, 0.406),
○ std=(0.229, 0.224, 0.225)),
○ ToTensorV2()
○])
○])
```

```
- def _load_and_augment(self, root_dir):
- sample_dirs = []
- for id in self.sample_ids:
- sample_dirs.append(os.path.join(root_dir, id))
-
- for dir_path in sample_dirs:
- img_path = os.path.join(dir_path, 'rgb.jpg')
- label_path = os.path.join(dir_path, 'labels.png')
-
- label_img = Image.open(label_path)
- is_rare =
- SegmentationDataset.contains_rare_classes(label_img,
- RARE_CLASSES)
-
- if not self.allow_augmentation:
- is_rare = False
```

```

 # Sempre aggiungi l'originale
 self.samples.append((img_path, label_path,
is_rare))

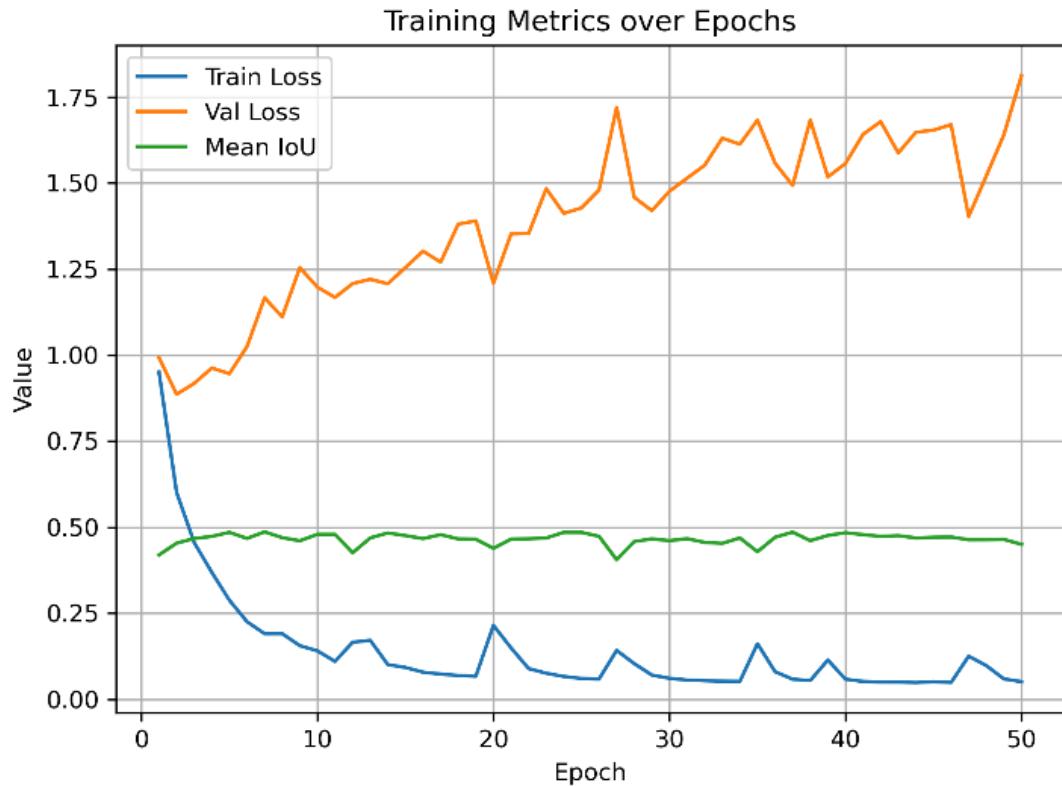
 # Duplicazione per oversampling (solo se contiene
 classi rare)
 if is_rare and not self.is_val:
 for _ in range(DUPLICATE_FACTOR - 1):
 self.samples.append((img_path,
label_path, is_rare))

```

- **Scheduler:** off
- **Optimizer:** Adama con lr=1e-4
- **Freezing:** none
- **Auxiliary Head:** off
- **Loss Function / Weighting:** cross-entropy + class weights euristici [1.0, 1.0, 1.0, 1.0, 5.0, 4.0, 3.5, 0.8, 0.8].
- **Main classifier Kernel size:** 1
- **Valutazione risultati:**
  - Migliore Validation Loss ottenuta all'epoca 2:
    - Train Loss: 0.6010 | Val Loss: 0.8861 | Mean IoU: 0.4538
    - IoU classe 0: 0.4030
    - IoU classe 1: 0.3244
    - IoU classe 2: 0.4519
    - IoU classe 3: 0.4907
    - IoU classe 4: 0.1171
    - IoU classe 5: 0.4992
    - IoU classe 6: 0.1440
    - IoU classe 7: 0.7271
    - IoU classe 8: 0.8758
- Migliore mean IoU ottenuta all'epoca 7:
  - Train Loss: 0.1906 | Val Loss: 1.1668 | Mean IoU: 0.4866
  - IoU classe 0: 0.2706
  - IoU classe 1: 0.3583
  - IoU classe 2: 0.4325
  - IoU classe 3: 0.4450
  - IoU classe 4: 0.2861
  - IoU classe 5: 0.5752

- IoU classe 6: 0.1388
- IoU classe 7: 0.7701
- IoU classe 8: 0.8867

- **Rappresentazione grafica del training:**



- **Massima memoria RAM occupata durante il training:** 2.61GB
- **Log di training:**  
`risultati_training\deeplabv3\resnet101\dlnr1010003\training_log_dlnr1010003.json`
- **Note / Bug Fix / Annotazioni:**

- picchi anomali ogni 25 epoche circa

## \* Esperimento: `deeplabv3_resnet101 - dlnr101004`

- **Architettura:** `deeplabv3plus_resnet101`
- **Backbone Weights:** DeepLabV3\_ResNet101\_Weights.DEFAULT
- **Input size:** 512x512
- **Test size:** 512x512

```
Augmentazione di base
val_aug = A.Compose([
 ...]
```

```
A.Resize(512, 512),
A.Normalize(mean=(0.485, 0.456, 0.406),
 std=(0.229, 0.224, 0.225)),
ToTensorV2()
])
```

- Batch size: 4

- Augmentation: on

```
o # Augmentazione per immagini con classi rare
o rare_aug = A.Compose([
o A.HorizontalFlip(p=0.5),
o A.RandomBrightnessContrast(p=0.4),
o #A.Rotate(limit=20, p=0.3), potenzialmente inutile
o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])
o
o # Augmentazione di base
o base_aug = A.Compose([
o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])
```

```
def _load_and_augment(self, root_dir):
 sample_dirs = []
 for id in self.sample_ids:
 sample_dirs.append(os.path.join(root_dir, id))

 for dir_path in sample_dirs:
 img_path = os.path.join(dir_path, 'rgb.jpg')
 label_path = os.path.join(dir_path, 'labels.png')

 label_img = Image.open(label_path)
 is_rare =
SegmentationDataset.contains_rare_classes(label_img,
RARE CLASSES)
```

```

 if not self.allow_augmentation:
 is_rare = False

 # Sempre aggiungi l'originale
 self.samples.append((img_path, label_path,
 is_rare))

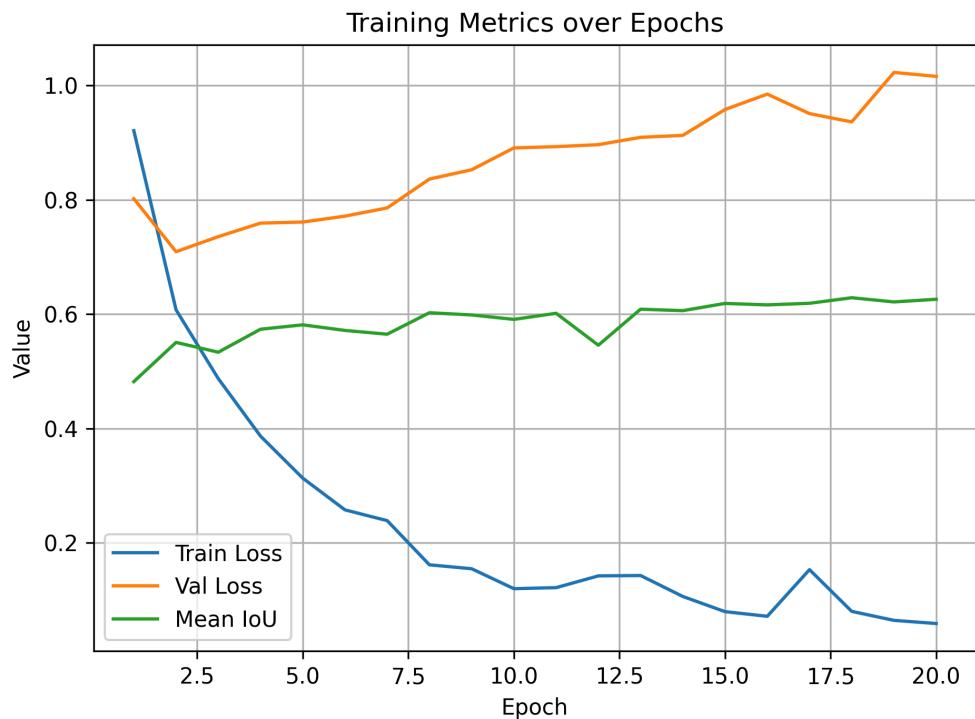
 # Duplicazione per oversampling (solo se contiene
 # classi rare)
 if is_rare and not self.is_val:
 for _ in range(DUPLICATE_FACTOR - 1):
 self.samples.append((img_path,
 label_path, is_rare))

```

- **Scheduler:** off
- **Optimizer:** Adama con lr=1e-4
- **Freezing:** none
- **Auxiliary Head:** off
- **Loss Function / Weighting:** cross-entropy + class weights euristici [1.0, 1.0, 1.0, 1.0, 5.0, 4.0, 3.5, 0.8, 0.8].
- **Main classifier Kernel size:** 1
- **Valutazione risultati:**
  - Migliore Validation Loss ottenuta all'epoca 2:
    - Train Loss: 0.6075 | Val Loss: 0.7091 | Mean IoU: 0.5504
    - ▶ IoU classe 0: 0.5259
    - ▶ IoU classe 1: 0.5440
    - ▶ IoU classe 2: 0.5052
    - ▶ IoU classe 3: 0.5481
    - ▶ IoU classe 4: 0.4803
    - ▶ IoU classe 5: 0.4454
    - ▶ IoU classe 6: 0.1824
    - ▶ IoU classe 7: 0.8097
    - ▶ IoU classe 8: 0.8885
  - Migliore mean IoU ottenuta all'epoca 18:
    - Train Loss: 0.0798 | Val Loss: 0.9362 | Mean IoU: 0.6286
    - ▶ IoU classe 0: 0.5765
    - ▶ IoU classe 1: 0.6088

- ► IoU classe 2: 0.6368
- ► IoU classe 3: 0.6030
- ► IoU classe 4: 0.5410
- ► IoU classe 5: 0.6921
- ► IoU classe 6: 0.2046
- ► IoU classe 7: 0.8388
- ► IoU classe 8: 0.9034

- **Rappresentazione grafica del training:**



- **Massima memoria RAM occupata durante il training:** 8.25GB
- **Log di training:**  
`risultati_training\deeplabv3\resnet101\dlnr1010004\training_log_dlnr1010004.json`
- **Note / Bug Fix / Annotazioni:**
  - con size 512x512 sia sul test set che sul training set, questo modello è stato scartato in quanto troppo esoso in termini di memoria RAM

\* Esperimento: `deeplabv3_resnet101 - dlnr101005`

- **Architettura:** `deeplabv3plus_resnet101`

- **Backbone Weights:** DeepLabV3\_ResNet101\_Weights.DEFAULT

- **Input size:** 256x256

- **Batch size:** 4

- **Augmentation:** on

```

o # Augmentazione per immagini con classi rare
o rare_aug = A.Compose([
o A.HorizontalFlip(p=0.5) ,
o A.RandomBrightnessContrast(p=0.4) ,
o #A.Rotate(limit=20, p=0.3), potenzialmente inutile
o A.Resize(256, 256) ,
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])
o
o
o # Augmentazione di base
o base_aug = A.Compose([
o A.Resize(256, 256) ,
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])

```

```

def _load_and_augment(self, root_dir):
 sample_dirs = []
 for id in self.sample_ids:
 sample_dirs.append(os.path.join(root_dir, id))

 for dir_path in sample_dirs:
 img_path = os.path.join(dir_path, 'rgb.jpg')
 label_path = os.path.join(dir_path, 'labels.png')

 label_img = Image.open(label_path)
 is_rare =
SegmentationDataset.contains_rare_classes(label_img,
RARE_CLASSES)

 if not self.allow_augmentation:
 is_rare = False

 # Sempre aggiungi l'originale

```

```

 self.samples.append((img_path, label_path,
is_rare))

 # Duplicazione per oversampling (solo se contiene
 classi rare)
 if is_rare and not self.is_val:
 for _ in range(DUPLICATE_FACTOR - 1):
 self.samples.append((img_path,
label_path, is_rare))

```

- **Scheduler:** off

- **Splitting:** on

```

o import numpy as np
o from PIL import Image
o import os
o from iterstrat.ml_stratifiers import
 MultilabelStratifiedShuffleSplit
o import json
o
o # sample male etichettati
o excluded_dirs = {
o "0007", "0009", "0090", "0095", "0101", "0104",
"0105",
o "0162", "0163", "0284", "0305", "0306", "0307",
"0308", "0309", "0310", "0311",
o "0351", "0372", "0373", "0376",
o "0498", "0499", "0500", "0501", "0526", "0527",
"0530", "0531", "0542",
o "0564", "0585", "0586", "0587", "0588", "0589",
"0590",
o #seguono nuove immagini da escludere
o "0000", "0001", "0052"
o }
o
o # Elenco di cartelle numeriche valide
o all_dirs = sorted([
o d for d in os.listdir(DATASET_DIR)
o if os.path.isdir(os.path.join(DATASET_DIR, d)) and
d.isdigit() and d not in excluded_dirs
o])
o
o # -----
o # COSTRUZIONE DELLA MATRICE MULTILABEL
o # -----
o def build_presence_matrix(dir_list, n_classes=9):

```

```

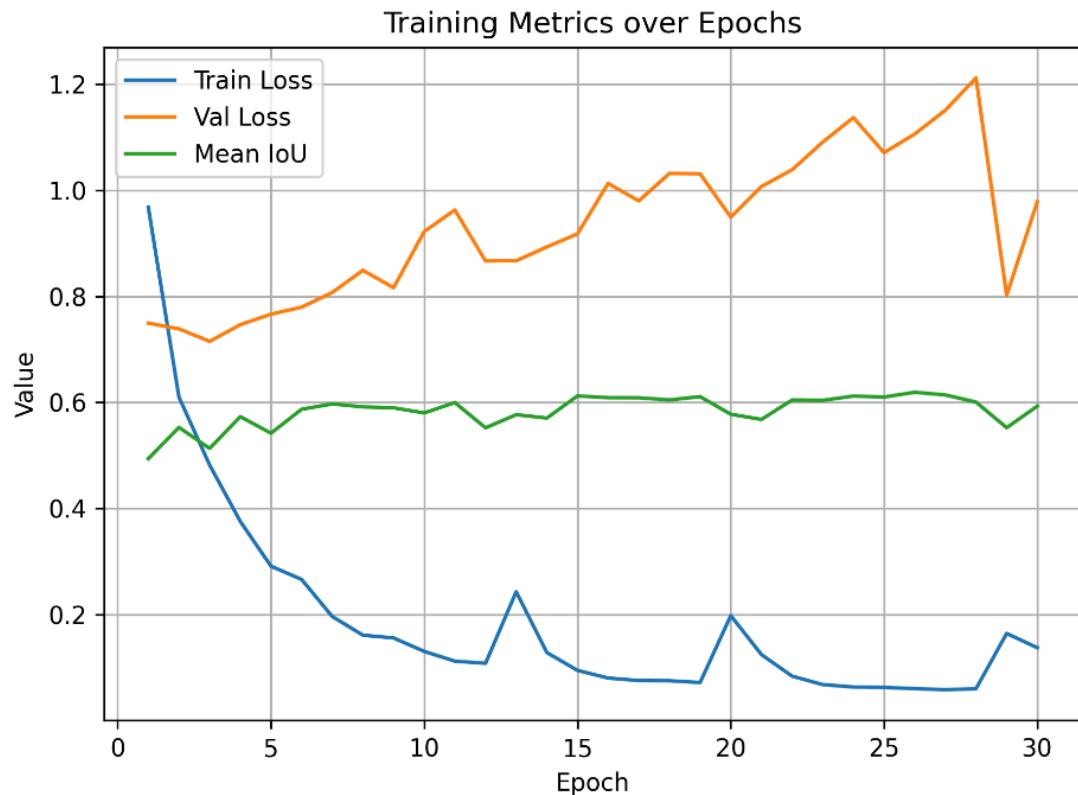
o """Return np.array (n_samples, n_classes) with
o booleans indicating class presence"""
o presence = np.zeros((len(dir_list), n_classes),
o dtype=int)
o
o for idx, d in enumerate(dir_list):
o label_path = os.path.join(DATASET_DIR, d,
o "labels.png")
o lbl = np.array(Image.open(label_path))
o uniq = np.unique(lbl)
o presence[idx, uniq] = 1 # segna le
o classi presenti
o return presence
o
o
o Y = build_presence_matrix(all_dirs, NUM_CLASSES)
o
o
o # -----
o # MULTILABEL STRATIFIED SPLIT
o # -----
o msss = MultilabelStratifiedShuffleSplit(n_splits=1,
o test_size=0.20, random_state=42)
o train_idx, val_idx =
o next(msss.split(np.zeros(len(all_dirs)), Y))
o
o train_ids = [all_dirs[i] for i in train_idx]
o val_ids = [all_dirs[i] for i in val_idx]
o
o print("Campioni:", len(train_ids), "train |",
o len(val_ids), "val")

```

- **Optimizer:** Adama con lr=1e-4
- **Freezing:** none
- **Auxiliary Head:** off
- **Loss Function / Weighting:** cross-entropy + class weights euristici [1.0, 1.0, 1.0, 1.0, 5.0, 4.0, 3.5, 0.8, 0.8].
- **Main classifier Kernel size:** 1
- **Valutazione risultati:**
  - Migliore Validation Loss ottenuta all'epoca 3:
  -  Train Loss: 0.4821 | Val Loss: 0.7150 | Mean IoU: 0.5133
    - ▶ IoU classe 0: 0.4856

- ▶ IoU classe 1: 0.6092
- ▶ IoU classe 2: 0.5384
- ▶ IoU classe 3: 0.6030
- ▶ IoU classe 4: 0.1352
- ▶ IoU classe 5: 0.2619
- ▶ IoU classe 6: 0.3044
- ▶ IoU classe 7: 0.7835
- ▶ IoU classe 8: 0.8706
- 
- Migliore mean IoU ottenuta all'epoca 26:
-  Train Loss: 0.0608 | Val Loss: 1.1057 | Mean IoU: 0.6191
- ▶ IoU classe 0: 0.5361
- ▶ IoU classe 1: 0.7011
- ▶ IoU classe 2: 0.6249
- ▶ IoU classe 3: 0.6641
- ▶ IoU classe 4: 0.1964
- ▶ IoU classe 5: 0.7440
- ▶ IoU classe 6: 0.3067
- ▶ IoU classe 7: 0.8261
- ▶ IoU classe 8: 0.8899

- **Rappresentazione grafica del training:**



- **Massima memoria RAM occupata durante il training:** 2.64GB

- **Log di training:**

```
risultati_training\deeplabv3\resnet101\dlnr1010001\training_log_dlnr1010005.json
```

- **Note / Bug Fix / Annotazioni:**

- picchi anomali ogni 10 epoche circa

## \* Esperimento: deeplabv3\_resnet101 - dlnr101006

- **Architettura:** deeplabv3plus\_resnet101
- **Backbone Weights:** DeepLabV3\_ResNet101\_Weights.DEFAULT
- **Input size:** 256x256
- **Batch size:** 4
- **Augmentation:** on

```

o # Augmentazione per immagini con classi rare
o rare_aug = A.Compose([
o A.HorizontalFlip(p=0.5),
o A.RandomBrightnessContrast(p=0.4),
o #A.Rotate(limit=20, p=0.3), potenzialmente inutile
o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])
o
o # Augmentazione di base
o base_aug = A.Compose([
o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])

```

```

def _load_and_augment(self, root_dir):
 sample_dirs = []
 for id in self.sample_ids:
 sample_dirs.append(os.path.join(root_dir, id))

 for dir_path in sample_dirs:
 img_path = os.path.join(dir_path, 'rgb.jpg')

```

```

 label_path = os.path.join(dir_path, 'labels.png')

 label_img = Image.open(label_path)
 is_rare =
SegmentationDataset.contains_rare_classes(label_img,
RARE_CLASSES)

 if not self.allow_augmentation:
 is_rare = False

 # Sempre aggiungi l'originale
 self.samples.append((img_path, label_path,
is_rare))

 # Duplicazione per oversampling (solo se contiene
classi rare)
 if is_rare and not self.is_val:
 for _ in range(DUPLICATE_FACTOR - 1):
 self.samples.append((img_path,
label_path, is_rare))

```

- **Scheduler:** off

- **Splitting:** on

```

o import numpy as np
o from PIL import Image
o import os
o from iterstrat.ml_stratifiers import
MultilabelStratifiedShuffleSplit
o import json
o
o # sample male etichettati
o excluded_dirs = {
o "0007", "0009", "0090", "0095", "0101", "0104",
"0105",
o "0162", "0163", "0284", "0305", "0306", "0307",
"0308", "0309", "0310", "0311",
o "0351", "0372", "0373", "0376",
o "0498", "0499", "0500", "0501", "0526", "0527",
"0530", "0531", "0542",
o "0564", "0585", "0586", "0587", "0588", "0589",
"0590",
o #seguono nuove immagini da escludere
o "0000", "0001", "0052"
o }
o

```

```

o # Elenco di cartelle numeriche valide
o all_dirs = sorted([
o d for d in os.listdir(DATASET_DIR)
o if os.path.isdir(os.path.join(DATASET_DIR, d)) and
o d.isdigit() and d not in excluded_dirs
o])
o
o #
o # COSTRUZIONE DELLA MATRICE MULTILABEL
o #
o def build_presence_matrix(dir_list, n_classes=9):
o """Return np.array (n_samples, n_classes) with
o booleans indicating class presence"""
o presence = np.zeros((len(dir_list), n_classes),
o dtype=int)
o
o for idx, d in enumerate(dir_list):
o label_path = os.path.join(DATASET_DIR, d,
o "labels.png")
o lbl = np.array(Image.open(label_path))
o uniq = np.unique(lbl)
o presence[idx, uniq] = 1 # segna le
o classi presenti
o
o return presence
o
o
o Y = build_presence_matrix(all_dirs, NUM_CLASSES)
o
o #
o # MULTILABEL STRATIFIED SPLIT
o #
o msss = MultilabelStratifiedShuffleSplit(n_splits=1,
o test_size=0.20, random_state=42)
o train_idx, val_idx =
o next(msss.split(np.zeros(len(all_dirs)), Y))
o
o train_ids = [all_dirs[i] for i in train_idx]
o val_ids = [all_dirs[i] for i in val_idx]
o
o print("Campioni:", len(train_ids), "train |",
o len(val_ids), "val")

```

- **Optimizer:** Adama con lr=1e-4

- **Freezing:** none

- **Auxiliary Head:** off
- **Loss Function / Weighting:** cross-entropy + class weights euristici [1.0, 1.0, 1.0, 1.0, 5.0, 4.0, 3.5, 0.8, 0.8].
- **Main classifier Kernel size:** 5
- **ASPP:** stride [3, 7, 3]

```

o model =
 models.segmentation.deeplabv3_resnet101(weights=weights)
o # versione custom del classifier
o model.classifier = CustomDeepLabHead(2048, NUM_CLASSES)

```

```

class CustomDeepLabHead(nn.Sequential):
 def __init__(self, in_channels, num_classes):
 super().__init__(
 ASPP(in_channels, [3, 7, 3]), # ASPP con
 dilatazioni personalizzate
 nn.Conv2d(256, 256, kernel_size=5,
 padding=2), # kernel 5x5
 nn.ReLU(),
 nn.Dropout(0),
 nn.Conv2d(256, num_classes, kernel_size=1)
 # output classi
)

```

- **Valutazione risultati:**

Migliore Validation Loss ottenuta all'epoca 2:

 Train Loss: 0.5513 | Val Loss: 0.6654 | Mean IoU: 0.5404

- ▶ IoU classe 0: 0.5130
- ▶ IoU classe 1: 0.6549
- ▶ IoU classe 2: 0.4952
- ▶ IoU classe 3: 0.6251
- ▶ IoU classe 4: 0.5213
- ▶ IoU classe 5: 0.1592
- ▶ IoU classe 6: 0.2579

- ▶ IoU classe 7: 0.7432

- ▶ IoU classe 8: 0.8667

Migliore mean IoU ottenuta all'epoca 12:

 Train Loss: 0.1404 | Val Loss: 1.0972 | Mean IoU: 0.6302

- ▶ IoU classe 0: 0.5576

- ▶ IoU classe 1: 0.6834

- ▶ IoU classe 2: 0.5835

- ▶ IoU classe 3: 0.6061

- ▶ IoU classe 4: 0.6778

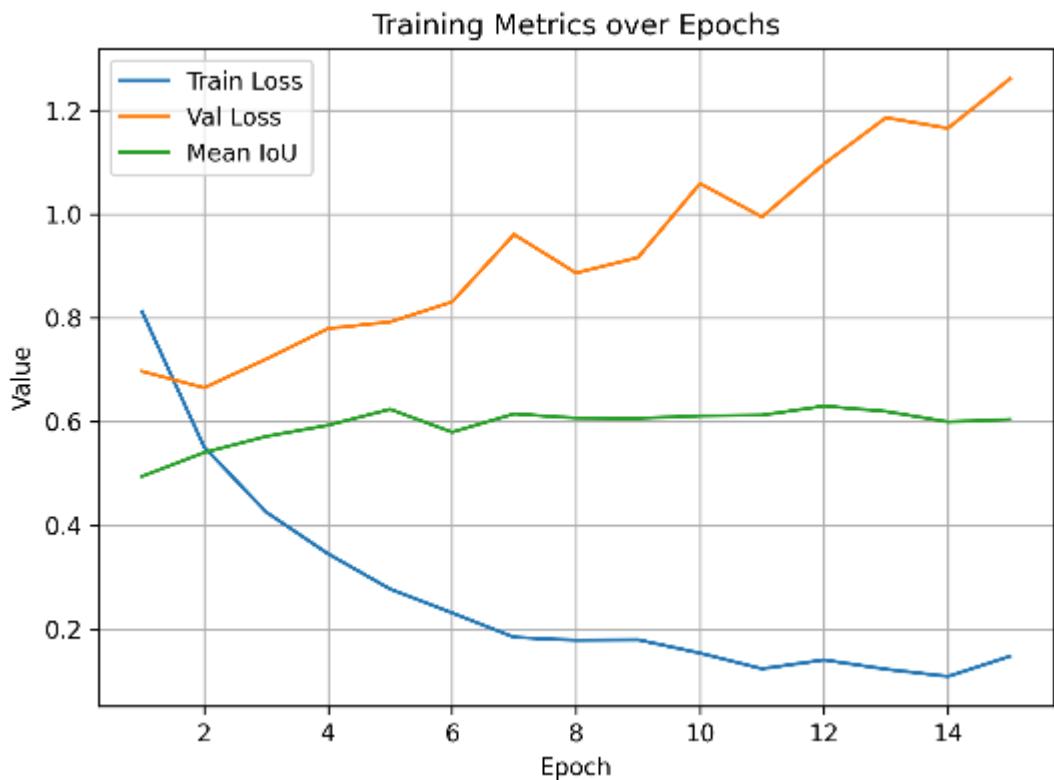
- ▶ IoU classe 5: 0.4874

- ▶ IoU classe 6: 0.2635

- ▶ IoU classe 7: 0.8385

- ▶ IoU classe 8: 0.9014

- **Rappresentazione grafica del training:**



- **Massima memoria RAM occupata durante il training:** 3.57GB
- **Log di training:**  
`risultati_training\deeplabv3\resnet101\dlnr101001\training_log_dlnr101006.json`
- **Note / Bug Fix / Annotazioni:**
  -

### \* Esperimento: deeplabv3\_resnet101 - dlnr101007

uguale a dlnr101006 ma con dilazione ASPP [1, 3, 5] con risultati peggiori perchè Migliore Validation Loss ottenuta all'epoca 2 con Mean IoU: 0.53

### \* Esperimento: deeplabv3\_resnet101 - dlnr101008

- **Architettura:** deeplabv3plus\_resnet101
- **Backbone Weights:** DeepLabV3\_ResNet101\_Weights.DEFAULT
- **Input size:** 256x256
- **Batch size:** 4
- **Augmentation:** on
  - # Augmentazione per immagini con classi rare

```

o rare_aug = A.Compose([
o A.HorizontalFlip(p=0.5),
o A.RandomBrightnessContrast(p=0.4),
o #A.Rotate(limit=20, p=0.3), potenzialmente inutile
o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])
o
o
o # Augmentazione di base
o base_aug = A.Compose([
o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])

```

```

def _load_and_augment(self, root_dir):
 sample_dirs = []
 for id in self.sample_ids:
 sample_dirs.append(os.path.join(root_dir, id))

 for dir_path in sample_dirs:
 img_path = os.path.join(dir_path, 'rgb.jpg')
 label_path = os.path.join(dir_path, 'labels.png')

 label_img = Image.open(label_path)
 is_rare =
SegmentationDataset.contains_rare_classes(label_img,
RARE_CLASSES)

 if not self.allow_augmentation:
 is_rare = False

 # Sempre aggiungi l'originale
 self.samples.append((img_path, label_path,
is_rare))

 # Duplicazione per oversampling (solo se contiene
classi rare)
 if is_rare and not self.is_val:
 for _ in range(DUPLICATE_FACTOR - 1):
 self.samples.append((img_path,
label_path, is_rare))

```

- **Scheduler:** off

- **Splitting:** on

```
o import numpy as np
o from PIL import Image
o import os
o from iterstrat.ml_stratifiers import
 MultilabelStratifiedShuffleSplit
o import json
o
o # sample male etichettati
o excluded_dirs = {
o "0007", "0009", "0090", "0095", "0101", "0104",
o "0105",
o "0162", "0163", "0284", "0305", "0306", "0307",
o "0308", "0309", "0310", "0311",
o "0351", "0372", "0373", "0376",
o "0498", "0499", "0500", "0501", "0526", "0527",
o "0530", "0531", "0542",
o "0564", "0585", "0586", "0587", "0588", "0589",
o "0590",
o #seguono nuovo immagini da escludere
o "0000", "0001", "0052"
o }
o
o # Elenco di cartelle numeriche valide
o all_dirs = sorted([
o d for d in os.listdir(DATASET_DIR)
o if os.path.isdir(os.path.join(DATASET_DIR, d)) and
o d.isdigit() and d not in excluded_dirs
o])
o
o # -----
o # COSTRUZIONE DELLA MATRICE MULTILABEL
o # -----
o def build_presence_matrix(dir_list, n_classes=9):
o """Return np.array (n_samples, n_classes) with
o booleans indicating class presence"""
o presence = np.zeros((len(dir_list), n_classes),
o dtype=int)
o
o for idx, d in enumerate(dir_list):
o label_path = os.path.join(DATASET_DIR, d,
o "labels.png")
o lbl = np.array(Image.open(label_path))
```

```

o uniq = np.unique(lbl)
o presence[idx, uniq] = 1 # segna le
o classi presenti
o return presence
o
o
o Y = build_presence_matrix(all_dirs, NUM_CLASSES)
o
o
o # -----
o # MULTILABEL STRATIFIED SPLIT
o # -----
o msss = MultilabelStratifiedShuffleSplit(n_splits=1,
o test_size=0.20, random_state=42)
o train_idx, val_idx =
o next(msss.split(np.zeros(len(all_dirs)), Y))
o
o train_ids = [all_dirs[i] for i in train_idx]
o val_ids = [all_dirs[i] for i in val_idx]
o
o print("Campioni:", len(train_ids), "train |",
o len(val_ids), "val")

```

- **Optimizer:** Adama con lr=1e-4
- **Freezing:** none
- **Auxiliary Head:** off
- **Loss Function / Weighting:** cross-entropy + class weights euristici [1.0, 1.0, 1.0, 1.0, 5.0, 4.0, 3.5, 0.8, 0.8].
- **Main classifier Kernel size:** 5
- **ASPP e dropout :** stride [3, 7, 3] e 0.1

```

o model =
o models.segmentation.deeplabv3_resnet101(weights=weights)
o # versione custom del classifier
o model.classifier = CustomDeepLabHead(2048, NUM_CLASSES)

```

```

class CustomDeepLabHead(nn.Sequential):
 def __init__(self, in_channels, num_classes):
 super().__init__(
 ASPP(in_channels, [3, 7, 3]), # ASPP con
 dilatazioni personalizzate
 nn.Conv2d(256, 256, kernel_size=5,
 padding=2), # kernel 5x5

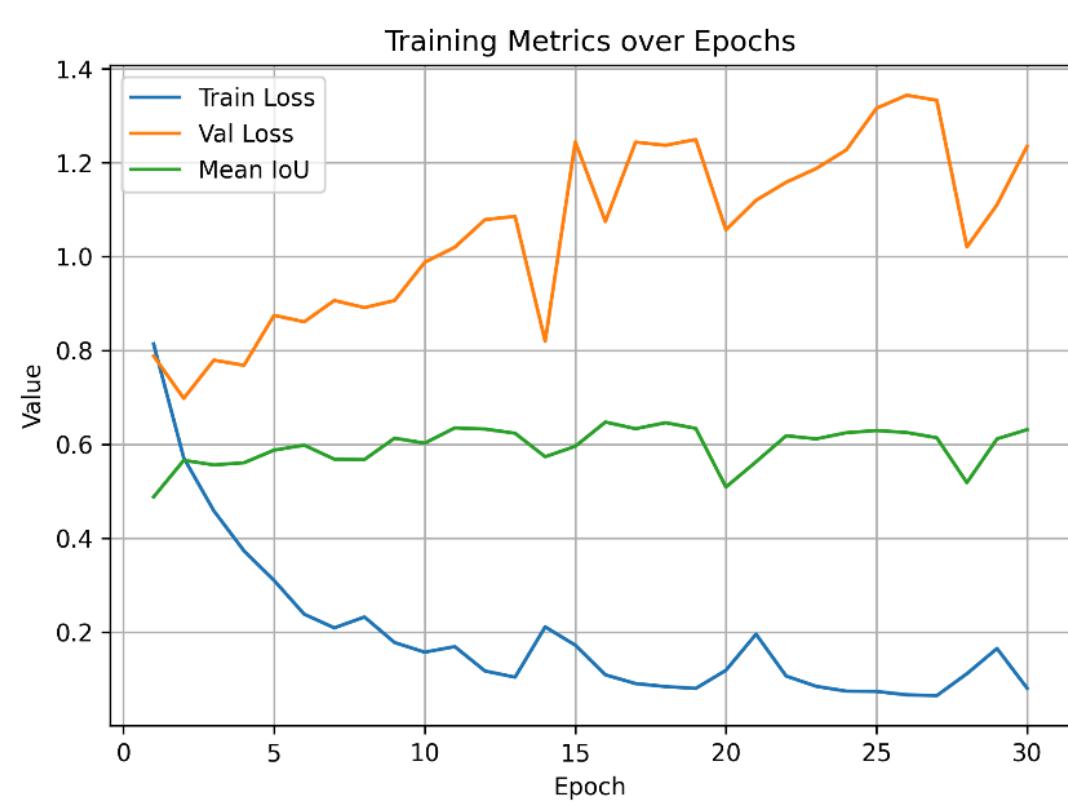
```

```
 nn.ReLU(),
 nn.Dropout(0.1),
 nn.Conv2d(256, num_classes, kernel_size=1)
 # output classi
)
```

- **Valutazione risultati:**

- Migliore Validation Loss ottenuta all'epoca 2:  
     Train Loss: 0.5711 | Val Loss: 0.6978 | Mean IoU: 0.5658
  - ▶ IoU classe 0: 0.4895
  - ▶ IoU classe 1: 0.6365
  - ▶ IoU classe 2: 0.3813
  - ▶ IoU classe 3: 0.5872
  - ▶ IoU classe 4: 0.6044
  - ▶ IoU classe 5: 0.3580
  - ▶ IoU classe 6: 0.2886
  - ▶ IoU classe 7: 0.7859
  - ▶ IoU classe 8: 0.8849
  -
- Migliore mean IoU ottenuta all'epoca 16:  
     Train Loss: 0.1089 | Val Loss: 1.0747 | Mean IoU: 0.6474
  - ▶ IoU classe 0: 0.5505
  - ▶ IoU classe 1: 0.6855
  - ▶ IoU classe 2: 0.5956
  - ▶ IoU classe 3: 0.6407
  - ▶ IoU classe 4: 0.7616
  - ▶ IoU classe 5: 0.4601
  - ▶ IoU classe 6: 0.3005
  - ▶ IoU classe 7: 0.8384
  - ▶ IoU classe 8: 0.8966

- **Rappresentazione grafica del training:**



- **Massima memoria RAM occupata durante il training:** 2.658GB
- **Log di training:**  
`risultati_training\deeplabv3\resnet101\dlnr1010001\training_log_dlnr1010008.json`
- **Note / Bug Fix / Annotazioni:**
  - picchi anomali ogni 10 epoche circa

## \* Esperimento: `deeplabv3_resnet101 - dlnr101009`

- **Architettura:** `deeplabv3plus_resnet101`
- **Backbone Weights:** DeepLabV3\_ResNet101\_Weights.DEFAULT
- **Input size:** 256x256
- **Batch size:** 4
- **Augmentation:** on

```

◦ # Augmentazione per immagini con classi rare
◦ rare_aug = A.Compose([
 A.HorizontalFlip(p=0.5),
 A.RandomBrightnessContrast(p=0.4),
 ...
])

```

```

o #A.Rotate(limit=20, p=0.3), potenzialmente inutile
o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])
o
o # Augmentazione di base
o base_aug = A.Compose([
o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])

```

---

```

def _load_and_augment(self, root_dir):
 sample_dirs = []
 for id in self.sample_ids:
 sample_dirs.append(os.path.join(root_dir, id))

 for dir_path in sample_dirs:
 img_path = os.path.join(dir_path, 'rgb.jpg')
 label_path = os.path.join(dir_path, 'labels.png')

 label_img = Image.open(label_path)
 is_rare =
SegmentationDataset.contains_rare_classes(label_img,
RARE_CLASSES)

 if not self.allow_augmentation:
 is_rare = False

 # Sempre aggiungi l'originale
 self.samples.append((img_path, label_path,
is_rare))

 # Duplicazione per oversampling (solo se contiene
classi rare)
 if is_rare and not self.is_val:
 for _ in range(DUPLICATE_FACTOR - 1):
 self.samples.append((img_path,
label_path, is_rare))

```

- **Scheduler:** off
- **Splitting:** on



```

o return presence
o
o Y = build_presence_matrix(all_dirs, NUM_CLASSES)
o
o # -----
o # MULTILABEL STRATIFIED SPLIT
o # -----
o msss = MultilabelStratifiedShuffleSplit(n_splits=1,
o test_size=0.20, random_state=42)
o train_idx, val_idx =
o next(msss.split(np.zeros(len(all_dirs)), Y))
o
o train_ids = [all_dirs[i] for i in train_idx]
o val_ids = [all_dirs[i] for i in val_idx]
o
o print("Campioni:", len(train_ids), "train |",
o len(val_ids), "val")

```

- **Optimizer:** Adama con lr=1e-4
- **Freezing:** none
- **Auxiliary Head:** off
- **Loss Function / Weighting:** cross-entropy + class weights euristici [1.0, 1.0, 1.0, 1.0, 5.0, 4.0, 3.5, 0.8, 0.8].
- **Main classifier Kernel size:** 5
- **ASPP e dropout :** stride [3, 7, 3] e 0.2

```

o model =
o models.segmentation.deeplabv3_resnet101(weights=weights)
o # versione custom del classifier
o model.classifier = CustomDeepLabHead(2048, NUM_CLASSES)

```

```

class CustomDeepLabHead(nn.Sequential):
 def __init__(self, in_channels, num_classes):
 super().__init__(
 ASPP(in_channels, [3, 7, 3]), # ASPP con
 dilatazioni personalizzate
 nn.Conv2d(256, 256, kernel_size=5,
padding=2), # kernel 5x5
 nn.ReLU(),
 nn.Dropout(0.2),

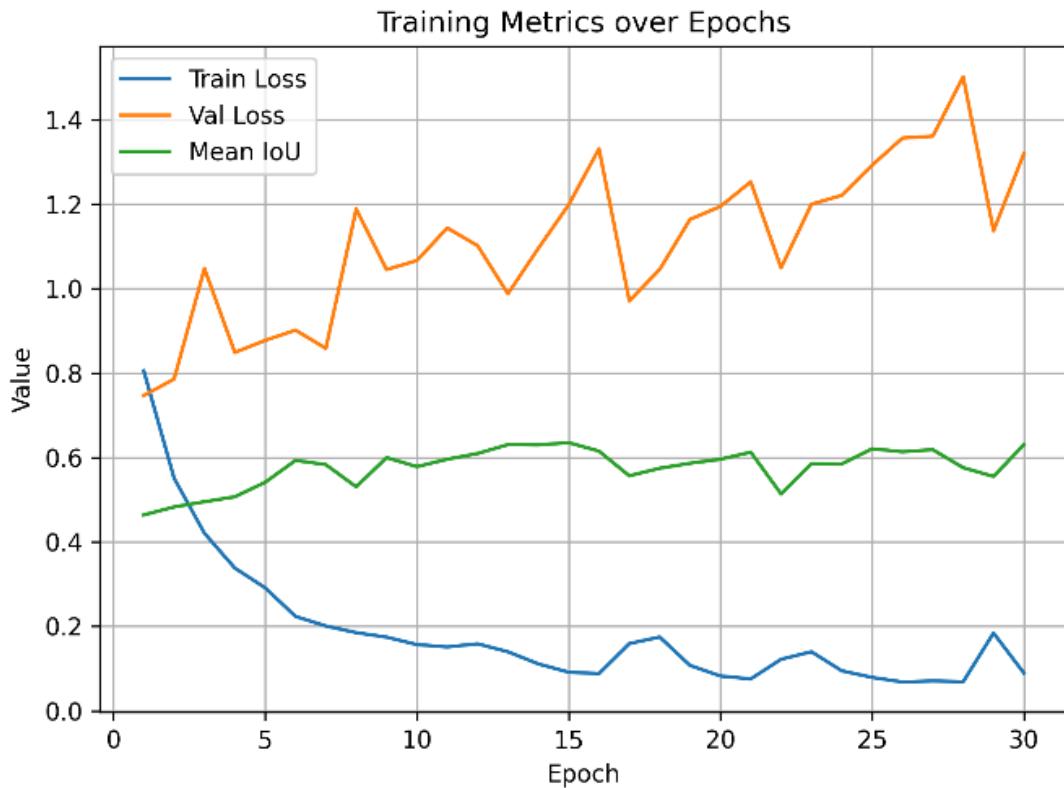
```

```
 nn.Conv2d(256, num_classes, kernel_size=1)
 # output classi
)
```

- **Valutazione risultati:**

- Migliore Validation Loss ottenuta all'epoca 1:  
     Train Loss: 0.8053 | Val Loss: 0.7467 | Mean IoU: 0.4639
  - ▶ IoU classe 0: 0.4377
  - ▶ IoU classe 1: 0.5977
  - ▶ IoU classe 2: 0.4377
  - ▶ IoU classe 3: 0.5946
  - ▶ IoU classe 4: 0.0000
  - ▶ IoU classe 5: 0.2050
  - ▶ IoU classe 6: 0.2629
  - ▶ IoU classe 7: 0.7605
  - ▶ IoU classe 8: 0.8529
  -
- Migliore mean IoU ottenuta all'epoca 15:  
     Train Loss: 0.0914 | Val Loss: 1.1984 | Mean IoU: 0.6351
  - ▶ IoU classe 0: 0.5441
  - ▶ IoU classe 1: 0.6598
  - ▶ IoU classe 2: 0.5785
  - ▶ IoU classe 3: 0.6332
  - ▶ IoU classe 4: 0.7150
  - ▶ IoU classe 5: 0.4863
  - ▶ IoU classe 6: 0.2818
  - ▶ IoU classe 7: 0.8304
  - ▶ IoU classe 8: 0.8961

- **Rappresentazione grafica del training:**



- **Massima memoria RAM occupata durante il training: 3.57GB**
- **Log di training:**  
`risultati_training\deeplabv3\resnet101\dlnr1010001\training_log_dlnr1010009.json`
- **Note / Bug Fix / Annotazioni:**

\* Esperimento: `deeplabv3_resnet101 - dlnr101010`

- **Architettura:** `deeplabv3plus_resnet101`
- **Backbone Weights:** DeepLabV3\_ResNet101\_Weights.DEFAULT
- **Input size:** 256x256
- **Batch size:** 4
- **Augmentation:** on

```

o # Augmentazione per immagini con classi rare
o rare_aug = A.Compose([
o A.HorizontalFlip(p=0.5),
o A.RandomBrightnessContrast(p=0.4),
o #A.Rotate(limit=20, p=0.3), potenzialmente inutile

```

```

o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])
o
o # Augmentazione di base
o base_aug = A.Compose([
o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])

```

```

def _load_and_augment(self, root_dir):
 sample_dirs = []
 for id in self.sample_ids:
 sample_dirs.append(os.path.join(root_dir, id))

 for dir_path in sample_dirs:
 img_path = os.path.join(dir_path, 'rgb.jpg')
 label_path = os.path.join(dir_path, 'labels.png')

 label_img = Image.open(label_path)
 is_rare =
 SegmentationDataset.contains_rare_classes(label_img,
RARE_CLASSES)

 if not self.allow_augmentation:
 is_rare = False

 # Sempre aggiungi l'originale
 self.samples.append((img_path, label_path,
is_rare))

 # Duplicazione per oversampling (solo se contiene
 classi rare)
 if is_rare and not self.is_val:
 for _ in range(DUPLICATE_FACTOR - 1):
 self.samples.append((img_path,
label_path, is_rare))

```

- **Scheduler:** off

- **Splitting:** on

```

o import numpy as np

```

```

o from PIL import Image
o import os
o from iterstrat.ml_stratifiers import
 MultilabelStratifiedShuffleSplit
o import json
o
o
o # sample male etichettati
o excluded_dirs = {
o "0007", "0009", "0090", "0095", "0101", "0104",
o "0105",
o "0162", "0163", "0284", "0305", "0306", "0307",
o "0308", "0309", "0310", "0311",
o "0351", "0372", "0373", "0376",
o "0498", "0499", "0500", "0501", "0526", "0527",
o "0530", "0531", "0542",
o "0564", "0585", "0586", "0587", "0588", "0589",
o "0590",
o #seguono nuove immagini da escludere
o "0000", "0001", "0052"
o }
o
o
o # Elenco di cartelle numeriche valide
o all_dirs = sorted([
o d for d in os.listdir(DATASET_DIR)
o if os.path.isdir(os.path.join(DATASET_DIR, d)) and
o d.isdigit() and d not in excluded_dirs
o])
o
o
o # -----
o # COSTRUZIONE DELLA MATRICE MULTILABEL
o # -----
o def build_presence_matrix(dir_list, n_classes=9):
o """Return np.array (n_samples, n_classes) with
o booleans indicating class presence"""
o presence = np.zeros((len(dir_list), n_classes),
o dtype=int)
o
o for idx, d in enumerate(dir_list):
o label_path = os.path.join(DATASET_DIR, d,
o "labels.png")
o lbl = np.array(Image.open(label_path))
o uniq = np.unique(lbl)
o presence[idx, uniq] = 1 # segna le
o classi presenti
o
o return presence

```

```

o
o Y = build_presence_matrix(all_dirs, NUM_CLASSES)
o
o # -----
o # MULTILABEL STRATIFIED SPLIT
o # -----
o msss = MultilabelStratifiedShuffleSplit(n_splits=1,
 test_size=0.20, random_state=42)
o train_idx, val_idx =
 next(msss.split(np.zeros(len(all_dirs)), Y))
o
o train_ids = [all_dirs[i] for i in train_idx]
o val_ids = [all_dirs[i] for i in val_idx]
o
o print("Campioni:", len(train_ids), "train |",
 len(val_ids), "val")

```

- **Optimizer:** Adama con lr=1e-4
- **Freezing:** none
- **Auxiliary Head:** off
- **Loss Function / Weighting:** cross-entropy + class weights euristici [1.0, 1.0, 1.0, 1.0, 5.0, 4.0, 3.5, 0.8, 0.8].
- **Main classifier Kernel size:** 5 e 3

```

o # ⚡ Primo layer convoluzionale: kernel 3x3
o nn.Conv2d(256, 256, kernel_size=3, padding=1), #
 mantiene dimensione
o nn.ReLU(),
o # ⚡ Secondo layer convoluzionale: kernel 5x5
o nn.Conv2d(256, 256, kernel_size=5, padding=2), # campo
 visivo più ampio
o nn.ReLU(),

```

- **ASPP:** stride [3, 7, 3]

```

o model =
 models.segmentation.deeplabv3_resnet101(weights=weights)
o # versione custom del classifier
o model.classifier = CustomDeepLabHead(2048, NUM_CLASSES)

```

```

-
class CustomDeepLabHead(nn.Sequential):
 def __init__(self, in_channels, num_classes):
 super().__init__(

```

```

 ASPP(in_channels, [3, 7, 3]), # ASPP con
 dilatazioni personalizzate
 nn.Conv2d(256, 256, kernel_size=5,
padding=2), # kernel 5x5
 nn.ReLU(),
 nn.Dropout(0),
 nn.Conv2d(256, num_classes, kernel_size=1)
output classi
)

```

- **Valutazione risultati:**

Migliore Validation Loss ottenuta all'epoca 3:

 Train Loss: 0.4911 | Val Loss: 0.7059 | Mean IoU: 0.5146

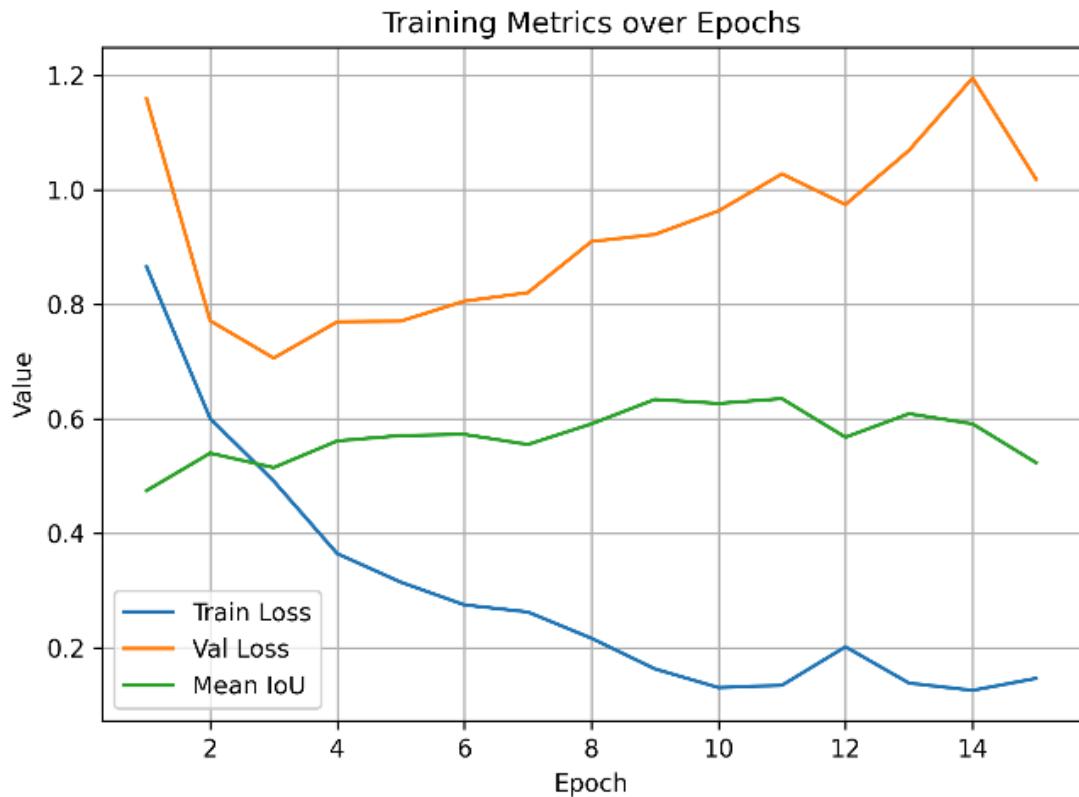
- ▶ IoU classe 0: 0.5091
- ▶ IoU classe 1: 0.6581
- ▶ IoU classe 2: 0.5004
- ▶ IoU classe 3: 0.6067
- ▶ IoU classe 4: 0.1908
- ▶ IoU classe 5: 0.1882
- ▶ IoU classe 6: 0.2894
- ▶ IoU classe 7: 0.8116
- ▶ IoU classe 8: 0.8712

Migliore mean IoU ottenuta all'epoca 11:

 Train Loss: 0.1342 | Val Loss: 1.0277 | Mean IoU: 0.6350

- ▶ IoU classe 0: 0.5409
- ▶ IoU classe 1: 0.6768
- ▶ IoU classe 2: 0.6031
- ▶ IoU classe 3: 0.6327
- ▶ IoU classe 4: 0.6595
- ▶ IoU classe 5: 0.4610
- ▶ IoU classe 6: 0.3177
- ▶ IoU classe 7: 0.8350
- ▶ IoU classe 8: 0.8941

- **Rappresentazione grafica del training:**



- **Massima memoria RAM occupata durante il training:** 4.03GB
- **Log di training:**  
`risultati_training\deeplabv3\resnet101\dlnr1010001\training_log_dlnr1010010.json`
- **Note / Bug Fix / Annotazioni:**

\* Esperimento: `deeplabv3_resnet101 - dlnr101011`

- **Architettura:** `deeplabv3plus_resnet101`
- **Backbone Weights:** DeepLabV3\_ResNet101\_Weights.DEFAULT
- **Input size:** 256x256
- **Batch size:** 4
- **Augmentation:** on

```

o # Augmentazione per immagini con classi rare
o rare_aug = A.Compose([
o A.HorizontalFlip(p=0.5),
o A.RandomBrightnessContrast(p=0.4),
o #A.Rotate(limit=20, p=0.3), potenzialmente inutile

```

```

o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])
o
o # Augmentazione di base
o base_aug = A.Compose([
o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])

```

```

def _load_and_augment(self, root_dir):
 sample_dirs = []
 for id in self.sample_ids:
 sample_dirs.append(os.path.join(root_dir, id))

 for dir_path in sample_dirs:
 img_path = os.path.join(dir_path, 'rgb.jpg')
 label_path = os.path.join(dir_path, 'labels.png')

 label_img = Image.open(label_path)
 is_rare =
 SegmentationDataset.contains_rare_classes(label_img,
RARE_CLASSES)

 if not self.allow_augmentation:
 is_rare = False

 # Sempre aggiungi l'originale
 self.samples.append((img_path, label_path,
is_rare))

 # Duplicazione per oversampling (solo se contiene
 classi rare)
 if is_rare and not self.is_val:
 for _ in range(DUPLICATE_FACTOR - 1):
 self.samples.append((img_path,
label_path, is_rare))

```

- **Scheduler:** off

- **Splitting:** on

```

o import numpy as np

```

```

o from PIL import Image
o import os
o from iterstrat.ml_stratifiers import
 MultilabelStratifiedShuffleSplit
o import json
o
o
o # sample male etichettati
o excluded_dirs = {
o "0007", "0009", "0090", "0095", "0101", "0104",
o "0105",
o "0162", "0163", "0284", "0305", "0306", "0307",
o "0308", "0309", "0310", "0311",
o "0351", "0372", "0373", "0376",
o "0498", "0499", "0500", "0501", "0526", "0527",
o "0530", "0531", "0542",
o "0564", "0585", "0586", "0587", "0588", "0589",
o "0590",
o #seguono nuovo immagini da escludere
o "0000", "0001", "0052"
o }
o
o
o # Elenco di cartelle numeriche valide
o all_dirs = sorted([
o d for d in os.listdir(DATASET_DIR)
o if os.path.isdir(os.path.join(DATASET_DIR, d)) and
d.isdigit() and d not in excluded_dirs
o])
o
o
o # -----
o # COSTRUZIONE DELLA MATRICE MULTILABEL
o # -----
o def build_presence_matrix(dir_list, n_classes=9):
o """Return np.array (n_samples, n_classes) with
o booleans indicating class presence"""
o presence = np.zeros((len(dir_list), n_classes),
o dtype=int)
o
o for idx, d in enumerate(dir_list):
o label_path = os.path.join(DATASET_DIR, d,
o "labels.png")
o lbl = np.array(Image.open(label_path))
o uniq = np.unique(lbl)
o presence[idx, uniq] = 1 # segna le
o classi presenti
o
o return presence

```

```

o
o Y = build_presence_matrix(all_dirs, NUM_CLASSES)
o
o # -----
o # MULTILABEL STRATIFIED SPLIT
o # -----
o msss = MultilabelStratifiedShuffleSplit(n_splits=1,
o test_size=0.20, random_state=42)
o train_idx, val_idx =
o next(msss.split(np.zeros(len(all_dirs)), Y))
o
o train_ids = [all_dirs[i] for i in train_idx]
o val_ids = [all_dirs[i] for i in val_idx]
o
o print("Campioni:", len(train_ids), "train |",
o len(val_ids), "val")

```

- **Optimizer:** Adama con lr=1e-4
- **Freezing:** none
- **Auxiliary Head:** off
- **Loss Function / Weighting:** cross-entropy + class weights euristici [1.0, 1.0, 1.0, 1.0, 5.0, 4.0, 3.5, 0.8, 0.8].
- **Main classifier Kernel size:** due di dimensione 5

```

⚡ Primo layer convoluzionale: kernel 3x3
- nn.Conv2d(256, 256, kernel_size=5, padding=2), #
 mantiene dimensione
- nn.ReLU(),
⚡ Secondo layer convoluzionale: kernel 5x5
- nn.Conv2d(256, 256, kernel_size=5, padding=2), # campo
 visivo più ampio
- nn.ReLU(),

```

- **ASPP:** stride [3, 7, 3]

```

o model =
o models.segmentation.deeplabv3_resnet101(weights=weights)
o # versione custom del classifier
o model.classifier = CustomDeepLabHead(2048, NUM_CLASSES)

```

```

-
- class CustomDeepLabHead(nn.Sequential):
- def __init__(self, in_channels, num_classes):
- super().__init__(
-
```

```
 ASPP(in_channels, [3, 7, 3]), # ASPP con
 dilatazioni personalizzate
 nn.Conv2d(256, 256, kernel_size=5,
padding=2), # kernel 5x5
 nn.ReLU(),
 nn.Dropout(0),
 nn.Conv2d(256, num_classes, kernel_size=1)
output classi
)
```

- **Valutazione risultati:**

Migliore Validation Loss ottenuta all'epoca 2:

 Train Loss: 0.6026 | Val Loss: 0.6464 | Mean IoU: 0.5708

- ▶ IoU classe 0: 0.4810
- ▶ IoU classe 1: 0.6194
- ▶ IoU classe 2: 0.5079
- ▶ IoU classe 3: 0.6208
- ▶ IoU classe 4: 0.6019
- ▶ IoU classe 5: 0.2128
- ▶ IoU classe 6: 0.3210
- ▶ IoU classe 7: 0.8015
- ▶ IoU classe 8: 0.8805

Migliore mean IoU ottenuta all'epoca 13:

 Train Loss: 0.1472 | Val Loss: 1.0552 | Mean IoU: 0.6213

- ▶ IoU classe 0: 0.5387
- ▶ IoU classe 1: 0.6678
- ▶ IoU classe 2: 0.5658
- ▶ IoU classe 3: 0.5696
- ▶ IoU classe 4: 0.7548

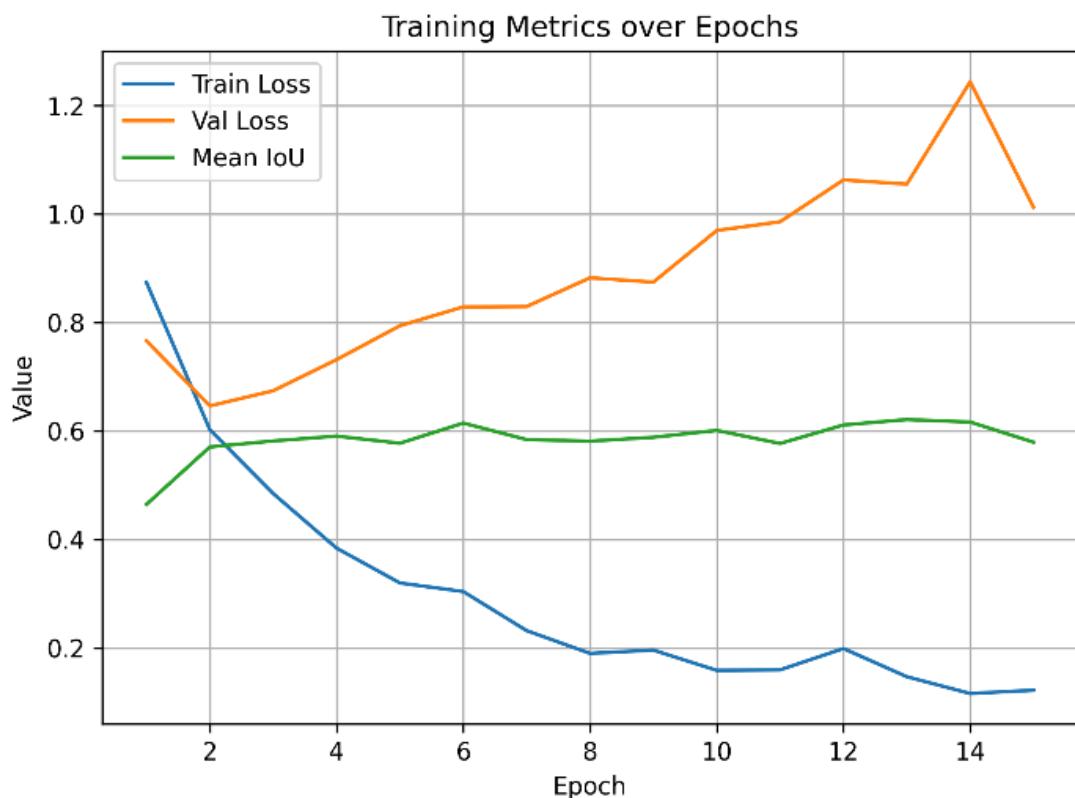
- ▶ IoU classe 5: 0.3527

- ▶ IoU classe 6: 0.3220

- ▶ IoU classe 7: 0.8402

- ▶ IoU classe 8: 0.8972

- **Rappresentazione grafica del training:**



- **Massima memoria RAM occupata durante il training:** 2.7GB

- **Log di training:**

`risultati_training\deeplabv3\resnet101\dlnr1010001\training_log_dlnr101011.json`

- **Note / Bug Fix / Annotazioni:**

\* Esperimento: `deeplabv3_resnet101 - dlnr101012`

- **Architettura:** `deeplabv3plus_resnet101`

- **Backbone Weights:** DeepLabV3\_ResNet101\_Weights.DEFAULT

- **Input size:** 256x256

- **Batch size:** 4

- **Augmentation:** on

```

o # Augmentazione per immagini con classi rare
o rare_aug = A.Compose([
o A.HorizontalFlip(p=0.5),
o A.RandomBrightnessContrast(p=0.4),
o #A.Rotate(limit=20, p=0.3), potenzialmente inutile
o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])
o
o # Augmentazione di base
o base_aug = A.Compose([
o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])

```

```

def _load_and_augment(self, root_dir):
 sample_dirs = []
 for id in self.sample_ids:
 sample_dirs.append(os.path.join(root_dir, id))

 for dir_path in sample_dirs:
 img_path = os.path.join(dir_path, 'rgb.jpg')
 label_path = os.path.join(dir_path, 'labels.png')

 label_img = Image.open(label_path)
 is_rare =
SegmentationDataset.contains_rare_classes(label_img,
RARE_CLASSES)

 if not self.allow_augmentation:
 is_rare = False

 # Sempre aggiungi l'originale
 self.samples.append((img_path, label_path,
is_rare))

 # Duplicazione per oversampling (solo se contiene
 classi rare)

```

```

 if is_rare and not self.is_val:
 for _ in range(DUPLICATE_FACTOR - 1):
 self.samples.append((img_path,
label_path, is_rare))

```

- **Scheduler:** off
- **Splitting:** on

```

o import numpy as np
o from PIL import Image
o import os
o from iterstrat.ml_stratifiers import
 MultilabelStratifiedShuffleSplit
o import json
o
o # sample male etichettati
o excluded_dirs = {
o "0007", "0009", "0090", "0095", "0101", "0104",
"0105",
o "0162", "0163", "0284", "0305", "0306", "0307",
"0308", "0309", "0310", "0311",
o "0351", "0372", "0373", "0376",
o "0498", "0499", "0500", "0501", "0526", "0527",
"0530", "0531", "0542",
o "0564", "0585", "0586", "0587", "0588", "0589",
"0590",
o #seguono nuovo immagini da escludere
o "0000", "0001", "0052"
o }
o
o # Elenco di cartelle numeriche valide
o all_dirs = sorted([
o d for d in os.listdir(DATASET_DIR)
o if os.path.isdir(os.path.join(DATASET_DIR, d)) and
d.isdigit() and d not in excluded_dirs
o])
o
o # -----
o # COSTRUZIONE DELLA MATRICE MULTILABEL
o # -----
o def build_presence_matrix(dir_list, n_classes=9):
o """Return np.array (n_samples, n_classes) with
booleans indicating class presence"""
o presence = np.zeros((len(dir_list), n_classes),
dtype=int)
o

```

```

o for idx, d in enumerate(dir_list):
o label_path = os.path.join(DATASET_DIR, d,
o "labels.png")
o lbl = np.array(Image.open(label_path))
o uniq = np.unique(lbl)
o presence[idx, uniq] = 1 # segna le
o classi presenti
o
o return presence
o
o
o Y = build_presence_matrix(all_dirs, NUM_CLASSES)
o
o
o # -----
o # MULTILABEL STRATIFIED SPLIT
o # -----
o msss = MultilabelStratifiedShuffleSplit(n_splits=1,
o test_size=0.20, random_state=42)
o train_idx, val_idx =
o next(msss.split(np.zeros(len(all_dirs)), Y))
o
o train_ids = [all_dirs[i] for i in train_idx]
o val_ids = [all_dirs[i] for i in val_idx]
o
o print("Campioni:", len(train_ids), "train |",
o len(val_ids), "val")

```

- **Optimizer:** Adama con lr=1e-4
- **Freezing:** none
- **Auxiliary Head:** off
- **Loss Function / Weighting:** cross-entropy + class weights euristici [1.0, 1.0, 1.0, 1.0, 5.0, 4.0, 3.5, 0.8, 0.8].
- **Main classifier Kernel size:** due di dimensione 5
  - # Primo layer convoluzionale: kernel 3x3  
nn.Conv2d(256, 256, kernel\_size=5, padding=2), # mantiene dimensione  
nn.ReLU(),
  - # Secondo layer convoluzionale: kernel 5x5  
nn.Conv2d(256, 256, kernel\_size=5, padding=2), # campo visivo più ampio  
nn.ReLU(),
- **ASPP:** stride [3, 7, 3]

```
o model =
 models.segmentation.deeplabv3_resnet101(weights=weights)
o # versione custom del classifier
o model.classifier = CustomDeepLabHead(2048, NUM_CLASSES)
```

```
class CustomDeepLabHead(nn.Sequential):
 def __init__(self, in_channels, num_classes):
 super().__init__(
 ASPP(in_channels, [3, 7, 3]), # ASPP con
 dilatazioni personalizzate
 nn.Conv2d(256, 256, kernel_size=5,
 padding=2), # kernel 5x5
 nn.ReLU(),
 nn.Dropout(0.05),
 nn.Conv2d(256, num_classes, kernel_size=1)
 # output classi
)
```

- **Valutazione risultati:**

Migliore Validation Loss ottenuta all'epoca 5:

 Train Loss: 0.3375 | Val Loss: 0.7033 | Mean IoU: 0.5997

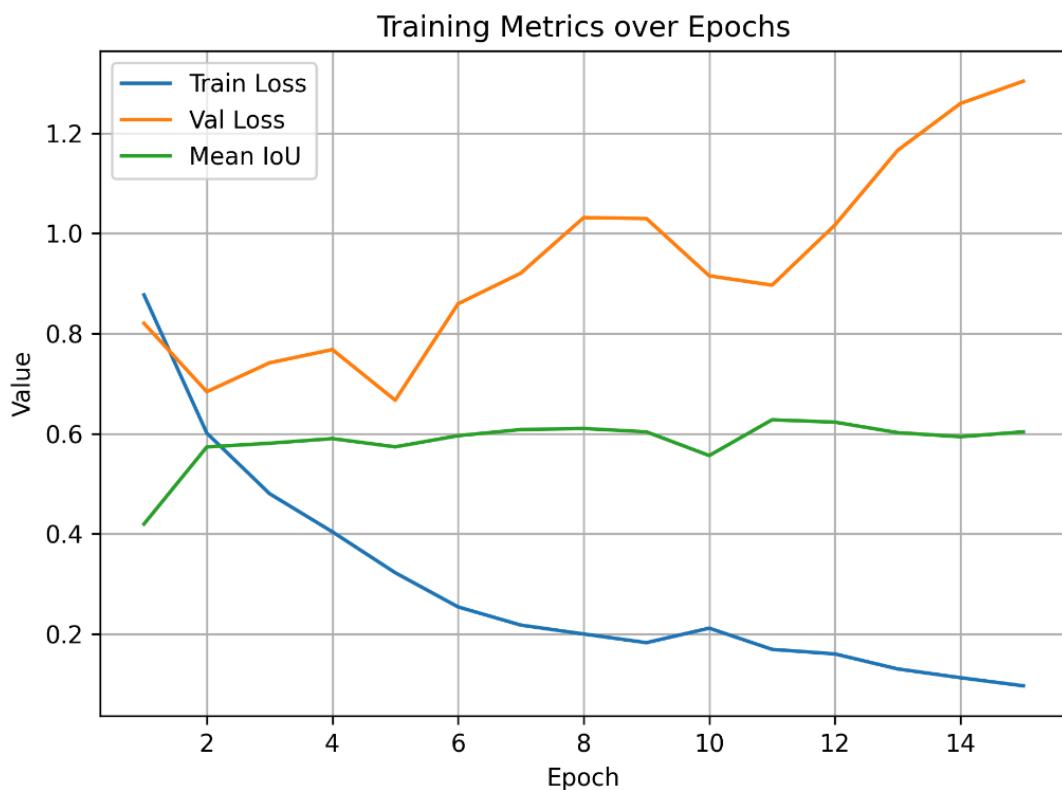
- ▶ IoU classe 0: 0.5237
- ▶ IoU classe 1: 0.6338
- ▶ IoU classe 2: 0.5476
- ▶ IoU classe 3: 0.5924
- ▶ IoU classe 4: 0.6271
- ▶ IoU classe 5: 0.4224
- ▶ IoU classe 6: 0.3016
- ▶ IoU classe 7: 0.7931
- ▶ IoU classe 8: 0.8793

Migliore mean IoU ottenuta all'epoca 9:

 Train Loss: 0.1911 | Val Loss: 0.9139 | Mean IoU: 0.6395

- ▶ IoU classe 0: 0.5093
- ▶ IoU classe 1: 0.6897
- ▶ IoU classe 2: 0.5568
- ▶ IoU classe 3: 0.6291
- ▶ IoU classe 4: 0.7707
- ▶ IoU classe 5: 0.4354
- ▶ IoU classe 6: 0.2997
- ▶ IoU classe 7: 0.8368
- ▶ IoU classe 8: 0.8975

- **Rappresentazione grafica del training:**



- **Massima memoria RAM occupata durante il training:** 2.7GB

- **Log di training:**  
risultati\_training\deeplabv3\resnet101\dlnr1010001\training\_log\_dlnr101012.json
- **Note / Bug Fix / Annotazioni:**

\* Esperimento: deeplabv3\_resnet101 - dlnr101013

- **Architettura:** deeplabv3plus\_resnet101
- **Backbone Weights:** DeepLabV3\_ResNet101\_Weights.DEFAULT
- **Input size:** 256x256
- **Batch size:** 4
- **Augmentation:** on
 

```

o # Augmentazione per immagini con classi rare
o rare_aug = A.Compose([
o A.HorizontalFlip(p=0.5),
o A.RandomBrightnessContrast(p=0.4),
o #A.Rotate(limit=20, p=0.3), potenzialmente inutile
o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])
o
o # Augmentazione di base
o base_aug = A.Compose([
o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])

```

```

def _load_and_augment(self, root_dir):
 sample_dirs = []
 for id in self.sample_ids:
 sample_dirs.append(os.path.join(root_dir, id))

 for dir_path in sample_dirs:
 img_path = os.path.join(dir_path, 'rgb.jpg')
 label_path = os.path.join(dir_path, 'labels.png')

```

```

 label_img = Image.open(label_path)
 is_rare =
SegmentationDataset.contains_rare_classes(label_img,
RARE_CLASSES)

 if not self.allow_augmentation:
 is_rare = False

 # Sempre aggiungi l'originale
 self.samples.append((img_path, label_path,
is_rare))

 # Duplicazione per oversampling (solo se contiene
 classi rare)
 if is_rare and not self.is_val:
 for _ in range(DUPLICATE_FACTOR - 1):
 self.samples.append((img_path,
label_path, is_rare))

```

- **Scheduler:** off

- **Splitting:** on

```

o import numpy as np
o from PIL import Image
o import os
o from iterstrat.ml_stratifiers import
 MultilabelStratifiedShuffleSplit
o import json
o
o
o # sample male etichettati
o excluded_dirs = {
o "0007", "0009", "0090", "0095", "0101", "0104",
"0105",
o "0162", "0163", "0284", "0305", "0306", "0307",
"0308", "0309", "0310", "0311",
o "0351", "0372", "0373", "0376",
o "0498", "0499", "0500", "0501", "0526", "0527",
"0530", "0531", "0542",
o "0564", "0585", "0586", "0587", "0588", "0589",
"0590",
o #seguono nuove immagini da escludere
o "0000", "0001", "0052"
o }
o
o # Elenco di cartelle numeriche valide

```

```

o all_dirs = sorted([
o d for d in os.listdir(DATASET_DIR)
o if os.path.isdir(os.path.join(DATASET_DIR, d)) and
d.isdigit() and d not in excluded_dirs
o])
o
o # -----
o # COSTRUZIONE DELLA MATRICE MULTILABEL
o # -----
o def build_presence_matrix(dir_list, n_classes=9):
o """Return np.array (n_samples, n_classes) with
o booleans indicating class presence"""
o presence = np.zeros((len(dir_list), n_classes),
o dtype=int)
o
o for idx, d in enumerate(dir_list):
o label_path = os.path.join(DATASET_DIR, d,
o "labels.png")
o lbl = np.array(Image.open(label_path))
o uniq = np.unique(lbl)
o presence[idx, uniq] = 1 # segna le
o classi presenti
o
o return presence
o
o
o Y = build_presence_matrix(all_dirs, NUM_CLASSES)
o
o # -----
o # MULTILABEL STRATIFIED SPLIT
o # -----
o msss = MultilabelStratifiedShuffleSplit(n_splits=1,
o test_size=0.20, random_state=42)
o train_idx, val_idx =
o next(msss.split(np.zeros(len(all_dirs)), Y))
o
o train_ids = [all_dirs[i] for i in train_idx]
o val_ids = [all_dirs[i] for i in val_idx]
o
o print("Campioni:", len(train_ids), "train |",
o len(val_ids), "val")

```

- **Optimizer:** Adama con lr=1e-4
- **Freezing:** none

- **Auxiliary Head:** off
- **Loss Function / Weighting:** cross-entropy + class weights euristici [1.0, 1.0, 1.0, 1.0, 5.0, 4.0, 3.5, 0.8, 0.8].
- **Main classifier Kernel size:** due di dimensione 5

```

🔍 Primo layer convoluzionale: kernel 3x3
nn.Conv2d(256, 256, kernel_size=5, padding=2), #
mantiene dimensione
nn.ReLU(),
🔍 Secondo layer convoluzionale: kernel 5x5
nn.Conv2d(256, 256, kernel_size=5, padding=2), # campo
visivo più ampio
nn.ReLU(),

```

- **ASPP:** stride [3, 7, 3]

```

o model =
 models.segmentation.deeplabv3_resnet101(weights=weights)
o # versione custom del classifier
o model.classifier = CustomDeepLabHead(2048, NUM_CLASSES)

```

```

class CustomDeepLabHead(nn.Sequential):
 def __init__(self, in_channels, num_classes):
 super().__init__(
 ASPP(in_channels, [3, 7, 3]), # ASPP con
 dilatazioni personalizzate
 nn.Conv2d(256, 256, kernel_size=5,
padding=2), # kernel 5x5
 nn.ReLU(),
 nn.Dropout(0.1),
 nn.Conv2d(256, num_classes, kernel_size=1)
 # output classi
)

```

- **Valutazione risultati:**

Migliore Validation Loss ottenuta all'epoca 2:

 Train Loss: 0.5890 | Val Loss: 0.6494 | Mean IoU: 0.5656

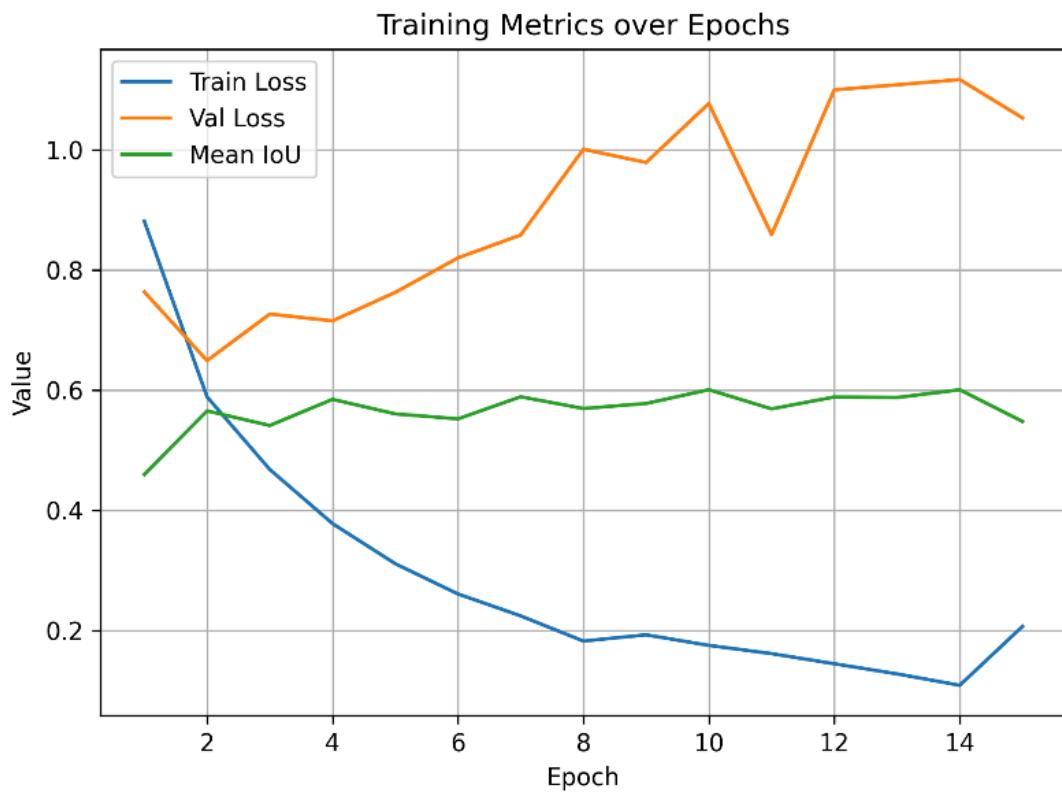
- ▶ IoU classe 0: 0.4966
- ▶ IoU classe 1: 0.6386
- ▶ IoU classe 2: 0.5564

- ▶ IoU classe 3: 0.6029
- ▶ IoU classe 4: 0.4459
- ▶ IoU classe 5: 0.2749
- ▶ IoU classe 6: 0.3078
- ▶ IoU classe 7: 0.8223
- ▶ IoU classe 8: 0.8765

Migliore mean IoU ottenuta all'epoca 10:

 Train Loss: 0.1752 | Val Loss: 1.0776 | Mean IoU: 0.6009

- ▶ IoU classe 0: 0.5302
  - ▶ IoU classe 1: 0.6489
  - ▶ IoU classe 2: 0.5761
  - ▶ IoU classe 3: 0.5994
  - ▶ IoU classe 4: 0.5974
  - ▶ IoU classe 5: 0.3799
  - ▶ IoU classe 6: 0.2832
  - ▶ IoU classe 7: 0.8323
  - ▶ IoU classe 8: 0.8902
- **Rappresentazione grafica del training:**



- **Massima memoria RAM occupata durante il training:** 2.6GB
- **Log di training:**  
`risultati_training\deeplabv3\resnet101\dlnr1010001\training_log_dlnr101013.json`
- **Note / Bug Fix / Annotazioni:**

\* Esperimento: `deeplabv3_resnet101 - dlnr101014`

- **Architettura:** `deeplabv3plus_resnet101`
- **Backbone Weights:** DeepLabV3\_ResNet101\_Weights.DEFAULT
- **Input size:** 256x256
- **Batch size:** 4
- **Augmentation:** on

```

o # Augmentazione per immagini con classi rare
o rare_aug = A.Compose([
o A.HorizontalFlip(p=0.5),
o A.RandomBrightnessContrast(p=0.4),
o #A.Rotate(limit=20, p=0.3), potenzialmente inutile

```

```

o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])
o
o # Augmentazione di base
o base_aug = A.Compose([
o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])

```

```

def _load_and_augment(self, root_dir):
 sample_dirs = []
 for id in self.sample_ids:
 sample_dirs.append(os.path.join(root_dir, id))

 for dir_path in sample_dirs:
 img_path = os.path.join(dir_path, 'rgb.jpg')
 label_path = os.path.join(dir_path, 'labels.png')

 label_img = Image.open(label_path)
 is_rare =
 SegmentationDataset.contains_rare_classes(label_img,
RARE_CLASSES)

 if not self.allow_augmentation:
 is_rare = False

 # Sempre aggiungi l'originale
 self.samples.append((img_path, label_path,
is_rare))

 # Duplicazione per oversampling (solo se contiene
 classi rare)
 if is_rare and not self.is_val:
 for _ in range(DUPLICATE_FACTOR - 1):
 self.samples.append((img_path,
label_path, is_rare))

```

- **Scheduler:** off

- **Splitting:** on

```

o import numpy as np

```

```

o from PIL import Image
o import os
o from iterstrat.ml_stratifiers import
 MultilabelStratifiedShuffleSplit
o import json
o
o # sample male etichettati
o excluded_dirs = {
o "0007", "0009", "0090", "0095", "0101", "0104",
o "0105",
o "0162", "0163", "0284", "0305", "0306", "0307",
o "0308", "0309", "0310", "0311",
o "0351", "0372", "0373", "0376",
o "0498", "0499", "0500", "0501", "0526", "0527",
o "0530", "0531", "0542",
o "0564", "0585", "0586", "0587", "0588", "0589",
o "0590",
o #seguono nuove immagini da escludere
o "0000", "0001", "0052"
o }
o
o # Elenco di cartelle numeriche valide
o all_dirs = sorted([
o d for d in os.listdir(DATASET_DIR)
o if os.path.isdir(os.path.join(DATASET_DIR, d)) and
d.isdigit() and d not in excluded_dirs
o])
o
o # -----
o # COSTRUZIONE DELLA MATRICE MULTILABEL
o # -----
o def build_presence_matrix(dir_list, n_classes=9):
o """Return np.array (n_samples, n_classes) with
o booleans indicating class presence"""
o presence = np.zeros((len(dir_list), n_classes),
o dtype=int)
o
o for idx, d in enumerate(dir_list):
o label_path = os.path.join(DATASET_DIR, d,
o "labels.png")
o lbl = np.array(Image.open(label_path))
o uniq = np.unique(lbl)
o presence[idx, uniq] = 1 # segna le
o classi presenti
o
o return presence

```

```

o
o Y = build_presence_matrix(all_dirs, NUM_CLASSES)
o
o # -----
o # MULTILABEL STRATIFIED SPLIT
o # -----
o msss = MultilabelStratifiedShuffleSplit(n_splits=1,
o test_size=0.20, random_state=42)
o train_idx, val_idx =
o next(msss.split(np.zeros(len(all_dirs)), Y))
o
o train_ids = [all_dirs[i] for i in train_idx]
o val_ids = [all_dirs[i] for i in val_idx]
o
o print("Campioni:", len(train_ids), "train |",
o len(val_ids), "val")

```

- **Optimizer:** Adama con lr=1e-4
- **Freezing:** none
- **Auxiliary Head:** off
- **Loss Function / Weighting:** cross-entropy + class weights euristici [1.0, 1.0, 1.0, 1.0, 5.0, 4.0, 3.5, 0.8, 0.8].
- **Main classifier Kernel size:** due di dimensione 5

```

⚡ Primo layer convoluzionale: kernel 3x3
- nn.Conv2d(256, 256, kernel_size=5, padding=2), #
- mantiene dimensione
- nn.ReLU(),
⚡ Secondo layer convoluzionale: kernel 5x5
- nn.Conv2d(256, 256, kernel_size=5, padding=2), # campo
- visivo più ampio
- nn.ReLU(),

```

- **ASPP:** stride [3, 7, 3]

```

o model =
o models.segmentation.deeplabv3_resnet101(weights=weights)
o # versione custom del classifier
o model.classifier = CustomDeepLabHead(2048, NUM_CLASSES)

```

```

-
- class CustomDeepLabHead(nn.Sequential):
- def __init__(self, in_channels, num_classes):
- super().__init__(
-
```

```

 ASPP(in_channels, [3, 7, 3]), # ASPP con
 dilatazioni personalizzate
 nn.Conv2d(256, 256, kernel_size=5,
padding=2), # kernel 5x5
 nn.ReLU(),
 nn.Dropout(0.07),
 nn.Conv2d(256, num_classes, kernel_size=1)
output classi
)

```

- **Valutazione risultati:**

Migliore Validation Loss ottenuta all'epoca 3:

 Train Loss: 0.4850 | Val Loss: 0.7054 | Mean IoU: 0.5732

- ▶ IoU classe 0: 0.5110
- ▶ IoU classe 1: 0.6382
- ▶ IoU classe 2: 0.5217
- ▶ IoU classe 3: 0.5848
- ▶ IoU classe 4: 0.5679
- ▶ IoU classe 5: 0.3303
- ▶ IoU classe 6: 0.2764
- ▶ IoU classe 7: 0.7965
- ▶ IoU classe 8: 0.8696

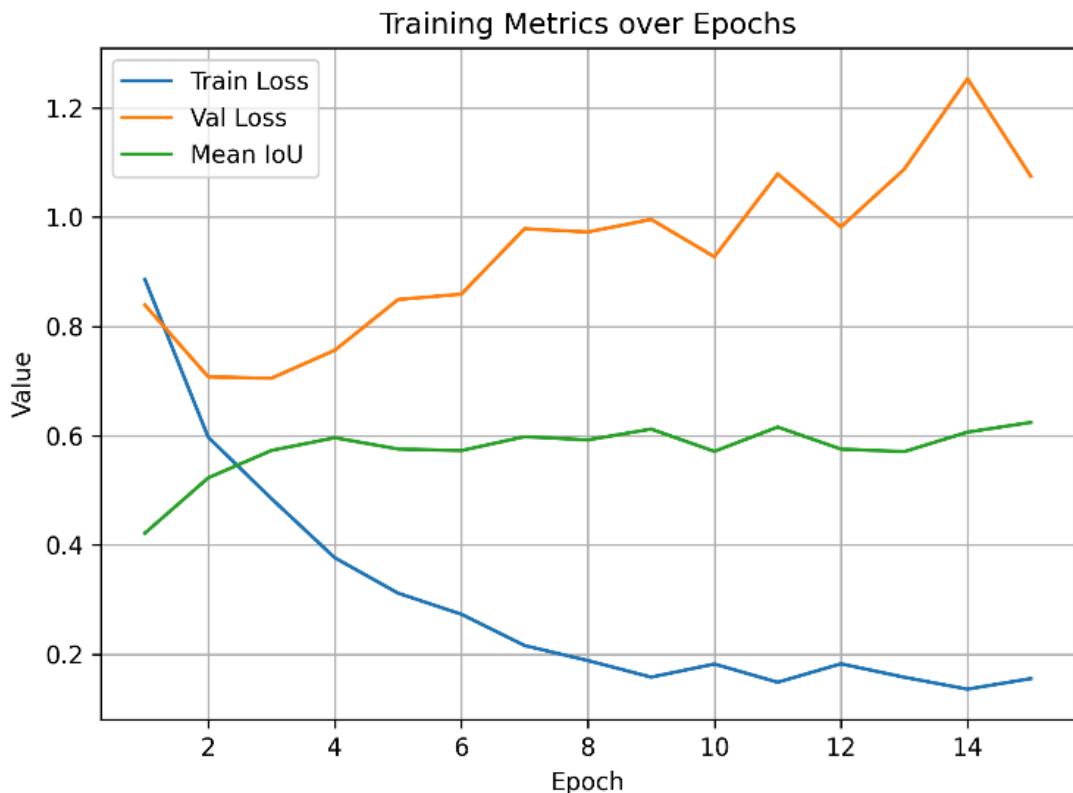
Migliore mean IoU ottenuta all'epoca 15:

 Train Loss: 0.1557 | Val Loss: 1.0749 | Mean IoU: 0.6245

- ▶ IoU classe 0: 0.5427
- ▶ IoU classe 1: 0.6423
- ▶ IoU classe 2: 0.5548
- ▶ IoU classe 3: 0.6181
- ▶ IoU classe 4: 0.6214

- ▶ IoU classe 5: 0.4975
- ▶ IoU classe 6: 0.3262
- ▶ IoU classe 7: 0.8368
- ▶ IoU classe 8: 0.8985

- **Rappresentazione grafica del training:**



- **Massima memoria RAM occupata durante il training:** 2.7GB
- **Log di training:**  
`risultati_training\deeplabv3\resnet101\dlnr1010001\training_log_dlnr101014.json`
- **Note / Bug Fix / Annotazioni:**

\* Esperimento: `deeplabv3_resnet101 - dlnr101015`

- **Architettura:** `deeplabv3plus_resnet101`
- **Backbone Weights:** DeepLabV3\_ResNet101\_Weights.DEFAULT
- **Input size:** 256x256

- **Batch size:** 4

- **Augmentation:** on

```

o # Augmentazione per immagini con classi rare
o rare_aug = A.Compose([
o A.HorizontalFlip(p=0.5),
o A.RandomBrightnessContrast(p=0.4),
o #A.Rotate(limit=20, p=0.3), potenzialmente inutile
o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])
o
o # Augmentazione di base
o base_aug = A.Compose([
o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])

```

```

def _load_and_augment(self, root_dir):
 sample_dirs = []
 for id in self.sample_ids:
 sample_dirs.append(os.path.join(root_dir, id))

 for dir_path in sample_dirs:
 img_path = os.path.join(dir_path, 'rgb.jpg')
 label_path = os.path.join(dir_path, 'labels.png')

 label_img = Image.open(label_path)
 is_rare =
SegmentationDataset.contains_rare_classes(label_img,
RARE_CLASSES)

 if not self.allow_augmentation:
 is_rare = False

 # Sempre aggiungi l'originale
 self.samples.append((img_path, label_path,
is_rare))

 # Duplicazione per oversampling (solo se contiene
 classi rare)

```

```

 if is_rare and not self.is_val:
 for _ in range(DUPLICATE_FACTOR - 1):
 self.samples.append((img_path,
label_path, is_rare))

```

- **Scheduler:** off
- **Splitting:** on

```

o import numpy as np
o from PIL import Image
o import os
o from iterstrat.ml_stratifiers import
 MultilabelStratifiedShuffleSplit
o import json
o
o # sample male etichettati
o excluded_dirs = {
o "0007", "0009", "0090", "0095", "0101", "0104",
"0105",
o "0162", "0163", "0284", "0305", "0306", "0307",
"0308", "0309", "0310", "0311",
o "0351", "0372", "0373", "0376",
o "0498", "0499", "0500", "0501", "0526", "0527",
"0530", "0531", "0542",
o "0564", "0585", "0586", "0587", "0588", "0589",
"0590",
o #seguono nuovo immagini da escludere
o "0000", "0001", "0052"
o }
o
o # Elenco di cartelle numeriche valide
o all_dirs = sorted([
o d for d in os.listdir(DATASET_DIR)
o if os.path.isdir(os.path.join(DATASET_DIR, d)) and
d.isdigit() and d not in excluded_dirs
o])
o
o # -----
o # COSTRUZIONE DELLA MATRICE MULTILABEL
o # -----
o def build_presence_matrix(dir_list, n_classes=9):
o """Return np.array (n_samples, n_classes) with
booleans indicating class presence"""
o presence = np.zeros((len(dir_list), n_classes),
dtype=int)
o

```

```

o for idx, d in enumerate(dir_list):
o label_path = os.path.join(DATASET_DIR, d,
o "labels.png")
o lbl = np.array(Image.open(label_path))
o uniq = np.unique(lbl)
o presence[idx, uniq] = 1 # segna le
o classi presenti
o
o return presence
o
o
o Y = build_presence_matrix(all_dirs, NUM_CLASSES)
o
o
o # -----
o # MULTILABEL STRATIFIED SPLIT
o # -----
o msss = MultilabelStratifiedShuffleSplit(n_splits=1,
o test_size=0.20, random_state=42)
o train_idx, val_idx =
o next(msss.split(np.zeros(len(all_dirs)), Y))
o
o train_ids = [all_dirs[i] for i in train_idx]
o val_ids = [all_dirs[i] for i in val_idx]
o
o print("Campioni:", len(train_ids), "train |",
o len(val_ids), "val")

```

- **Optimizer:** Adama con lr=1e-4
- **Freezing:** none
- **Auxiliary Head:** off
- **Loss Function / Weighting:** cross-entropy + class weights euristici [1.0, 1.0, 1.0, 1.0, 5.0, 4.0, 3.5, 0.8, 0.8].
- **Main classifier Kernel size:** 5, 7
  - # Primo layer convoluzionale: kernel 3x3  
nn.Conv2d(256, 256, kernel\_size=5, padding=2), # mantiene dimensione  
nn.ReLU(),
  - # Secondo layer convoluzionale: kernel 7x7  
nn.Conv2d(256, 256, kernel\_size=7, padding=3), # campo visivo più ampio  
nn.ReLU(),
- **ASPP:** stride [3, 7, 3]

```

o model =
 models.segmentation.deeplabv3_resnet101(weights=weights)
o # versione custom del classifier
o model.classifier = CustomDeepLabHead(2048, NUM_CLASSES)

```

```

class CustomDeepLabHead(nn.Sequential):
 def __init__(self, in_channels, num_classes):
 super().__init__(
 ASPP(in_channels, [3, 7, 3]), # ASPP con
 dilatazioni personalizzate
 nn.Conv2d(256, 256, kernel_size=5,
 padding=2), # kernel 5x5
 nn.ReLU(),
 nn.Dropout(0),
 nn.Conv2d(256, num_classes, kernel_size=1)
 # output classi
)

```

- **Valutazione risultati:**

Migliore Validation Loss ottenuta all'epoca 2:

 Train Loss: 0.6266 | Val Loss: 0.6906 | Mean IoU: 0.5099

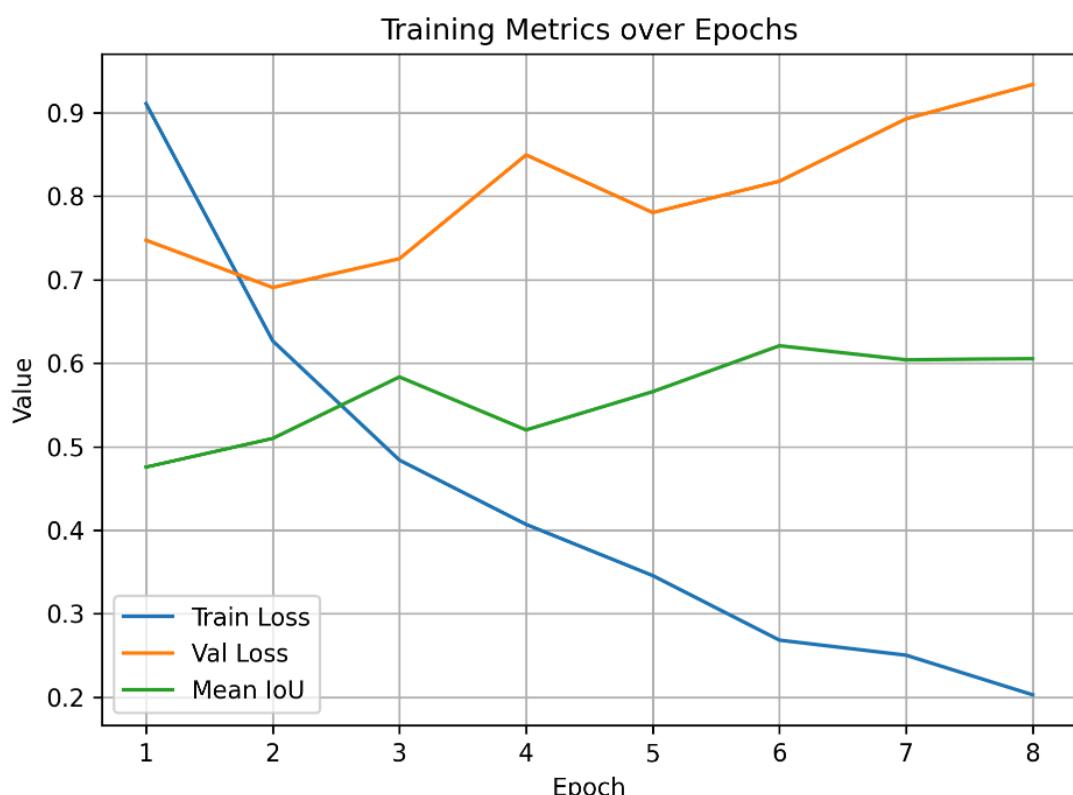
- ▶ IoU classe 0: 0.4924
- ▶ IoU classe 1: 0.5962
- ▶ IoU classe 2: 0.4811
- ▶ IoU classe 3: 0.5892
- ▶ IoU classe 4: 0.2433
- ▶ IoU classe 5: 0.1892
- ▶ IoU classe 6: 0.3024
- ▶ IoU classe 7: 0.8110
- ▶ IoU classe 8: 0.8667

Migliore mean IoU ottenuta all'epoca 6:

 Train Loss: 0.2683 | Val Loss: 0.8179 | Mean IoU: 0.6209

- ▶ IoU classe 0: 0.5330
- ▶ IoU classe 1: 0.6726
- ▶ IoU classe 2: 0.5424
- ▶ IoU classe 3: 0.6475
- ▶ IoU classe 4: 0.7008
- ▶ IoU classe 5: 0.3720
- ▶ IoU classe 6: 0.3107
- ▶ IoU classe 7: 0.8321
- ▶ IoU classe 8: 0.8888

- **Rappresentazione grafica del training:**



- **Massima memoria RAM occupata durante il training:** 3.15GB
- **Log di training:**  
`risultati_training\deeplabv3\resnet101\dlnr1010001\training_log_dlnr101015.json`

- **Note / Bug Fix / Annotazioni:**

\* Esperimento: **deeplabv3\_resnet101 - dlrn101016**

- **Architettura:** **deeplabv3plus\_resnet101**
- **Backbone Weights:** DeepLabV3\_ResNet101\_Weights.DEFAULT
- **Input size:** 256x256
- **Batch size:** 4
- **Augmentation:** on

```

o # Augmentazione per immagini con classi rare
o rare_aug = A.Compose([
o A.HorizontalFlip(p=0.5),
o A.RandomBrightnessContrast(p=0.4),
o #A.Rotate(limit=20, p=0.3), potenzialmente inutile
o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])
o
o
o # Augmentazione di base
o base_aug = A.Compose([
o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])

```

```

def _load_and_augment(self, root_dir):
 sample_dirs = []
 for id in self.sample_ids:
 sample_dirs.append(os.path.join(root_dir, id))

 for dir_path in sample_dirs:
 img_path = os.path.join(dir_path, 'rgb.jpg')
 label_path = os.path.join(dir_path, 'labels.png')

 label_img = Image.open(label_path)

```

```

 is_rare =
SegmentationDataset.contains_rare_classes(label_img,
RARE_CLASSES)

 if not self.allow_augmentation:
 is_rare = False

 # Sempre aggiungi l'originale
 self.samples.append((img_path, label_path,
is_rare))

 # Duplicazione per oversampling (solo se contiene
 classi rare)
 if is_rare and not self.is_val:
 for _ in range(DUPLICATE_FACTOR - 1):
 self.samples.append((img_path,
label_path, is_rare))

```

- **Scheduler:** off

- **Splitting:** on

```

o import numpy as np
o from PIL import Image
o import os
o from iterstrat.ml_stratifiers import
 MultilabelStratifiedShuffleSplit
o import json
o
o # sample male etichettati
o excluded_dirs = {
o "0007", "0009", "0090", "0095", "0101", "0104",
"0105",
o "0162", "0163", "0284", "0305", "0306", "0307",
"0308", "0309", "0310", "0311",
o "0351", "0372", "0373", "0376",
o "0498", "0499", "0500", "0501", "0526", "0527",
"0530", "0531", "0542",
o "0564", "0585", "0586", "0587", "0588", "0589",
"0590",
o #seguono nuove immagini da escludere
o "0000", "0001", "0052"
o }
o
o # Elenco di cartelle numeriche valide
o all_dirs = sorted([
o d for d in os.listdir(DATASET_DIR)

```

```

o if os.path.isdir(os.path.join(DATASET_DIR, d)) and
o d.isdigit() and d not in excluded_dirs
o])
o
o # -----
o # COSTRUZIONE DELLA MATRICE MULTILABEL
o # -----
o def build_presence_matrix(dir_list, n_classes=9):
o """Return np.array (n_samples, n_classes) with
o booleans indicating class presence"""
o presence = np.zeros((len(dir_list), n_classes),
o dtype=int)
o
o for idx, d in enumerate(dir_list):
o label_path = os.path.join(DATASET_DIR, d,
o "labels.png")
o lbl = np.array(Image.open(label_path))
o uniq = np.unique(lbl)
o presence[idx, uniq] = 1 # segna le
o classi presenti
o
o return presence
o
o
o Y = build_presence_matrix(all_dirs, NUM_CLASSES)
o
o # -----
o # MULTILABEL STRATIFIED SPLIT
o # -----
o msss = MultilabelStratifiedShuffleSplit(n_splits=1,
o test_size=0.20, random_state=42)
o train_idx, val_idx =
o next(msss.split(np.zeros(len(all_dirs)), Y))
o
o train_ids = [all_dirs[i] for i in train_idx]
o val_ids = [all_dirs[i] for i in val_idx]
o
o print("Campioni:", len(train_ids), "train |",
o len(val_ids), "val")

```

- **Optimizer:** Adama con lr=1e-4
- **Freezing:** none
- **Auxiliary Head:** off

- **Loss Function / Weighting:** cross-entropy + class weights euristici [1.0, 1.0, 1.0, 1.0, 5.0, 4.0, 3.5, 0.8, 0.8].

- **Main classifier Kernel size: 5, 7**

```
⚡ Primo layer convoluzionale: kernel 5x5
nn.Conv2d(256, 256, kernel_size=5, padding=2), # mantiene dimensione
nn.ReLU(),
⚡ Secondo layer convoluzionale: kernel 5x5
nn.Conv2d(256, 256, kernel_size=5, padding=2), # campo visivo più ampio
nn.ReLU(),
```

- **ASPP: stride [3, 7, 3]**

- model =  
models.segmentation.deeplabv3\_resnet101(weights=weights)
- # versione custom del classifier
- model.classifier = CustomDeepLabHead(2048, NUM\_CLASSES)

```
class CustomDeepLabHead(nn.Sequential):
 def __init__(self, in_channels, num_classes):
 super().__init__(
 ASPP(in_channels, [3, 7, 3]), # ASPP con dilatazioni personalizzate
 nn.Conv2d(256, 256, kernel_size=5,
padding=2), # kernel 5x5
 nn.ReLU(),
 nn.Dropout(0.02),
 nn.Conv2d(256, num_classes, kernel_size=1)
 # output class
)
```

- **Valutazione risultati:**

Migliore Validation Loss ottenuta all'epoca 2:

 Train Loss: 0.5938 | Val Loss: 0.6856 | Mean IoU: 0.5735

▶ IoU classe 0: 0.4865

▶ IoU classe 1: 0.6379

▶ IoU classe 2: 0.5090

▶ IoU classe 3: 0.5968

▶ IoU classe 4: 0.4395

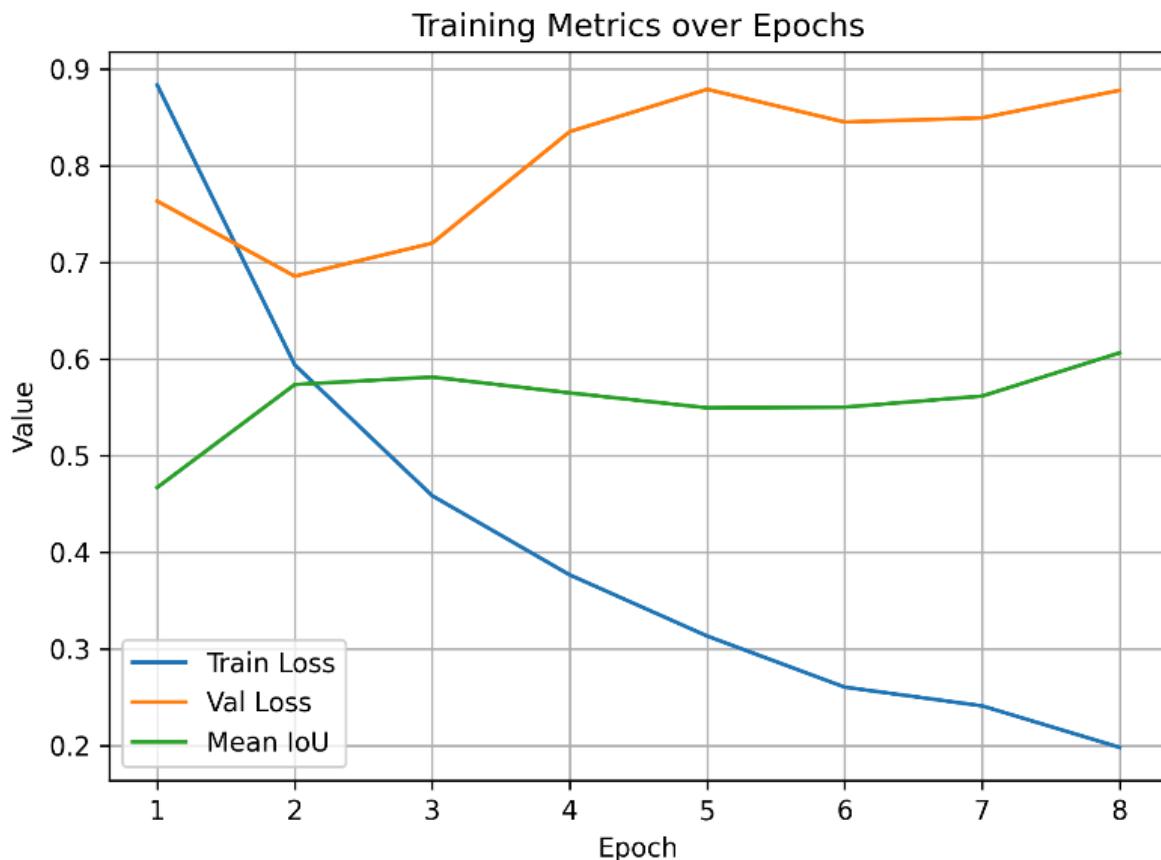
- ▶ IoU classe 5: 0.4210
- ▶ IoU classe 6: 0.3014
- ▶ IoU classe 7: 0.8162
- ▶ IoU classe 8: 0.8660

Migliore mean IoU ottenuta all'epoca 8:

 Train Loss: 0.1980 | Val Loss: 0.8780 | Mean IoU: 0.6063

- ▶ IoU classe 0: 0.5290
- ▶ IoU classe 1: 0.6790
- ▶ IoU classe 2: 0.5932
- ▶ IoU classe 3: 0.6270
- ▶ IoU classe 4: 0.5069
- ▶ IoU classe 5: 0.4022
- ▶ IoU classe 6: 0.3390
- ▶ IoU classe 7: 0.8210
- ▶ IoU classe 8: 0.8819

- **Rappresentazione grafica del training:**



- **Massima memoria RAM occupata durante il training:** 2.7GB
- **Log di training:**  
`risultati_training\deeplabv3\resnet101\dln101001\training_log_dln101016.json`
- **Note / Bug Fix / Annotazioni:**

\* Esperimento: `deeplabv3_resnet101 - dln101017`

- **Architettura:** `deeplabv3plus_resnet101`
- **Backbone Weights:** DeepLabV3\_ResNet101\_Weights.DEFAULT
- **Input size:** 256x256
- **Batch size:** 4
- **Augmentation:** on

```

o # Augmentazione per immagini con classi rare
o rare_aug = A.Compose([
o A.HorizontalFlip(p=0.5),
o A.RandomBrightnessContrast(p=0.4),

```

```

o #A.Rotate(limit=20, p=0.3), potenzialmente inutile
o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])
o
o # Augmentazione di base
o base_aug = A.Compose([
o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])

```

---

```

def _load_and_augment(self, root_dir):
 sample_dirs = []
 for id in self.sample_ids:
 sample_dirs.append(os.path.join(root_dir, id))

 for dir_path in sample_dirs:
 img_path = os.path.join(dir_path, 'rgb.jpg')
 label_path = os.path.join(dir_path, 'labels.png')

 label_img = Image.open(label_path)
 is_rare =
SegmentationDataset.contains_rare_classes(label_img,
RARE_CLASSES)

 if not self.allow_augmentation:
 is_rare = False

 # Sempre aggiungi l'originale
 self.samples.append((img_path, label_path,
is_rare))

 # Duplicazione per oversampling (solo se contiene
classi rare)
 if is_rare and not self.is_val:
 for _ in range(DUPLICATE_FACTOR - 1):
 self.samples.append((img_path,
label_path, is_rare))

```

- **Scheduler:** off
- **Splitting:** on



```

o return presence
o
o Y = build_presence_matrix(all_dirs, NUM_CLASSES)
o
o # -----
o # MULTILABEL STRATIFIED SPLIT
o # -----
o msss = MultilabelStratifiedShuffleSplit(n_splits=1,
o test_size=0.20, random_state=42)
o train_idx, val_idx =
o next(msss.split(np.zeros(len(all_dirs)), Y))
o
o train_ids = [all_dirs[i] for i in train_idx]
o val_ids = [all_dirs[i] for i in val_idx]
o
o print("Campioni:", len(train_ids), "train |",
o len(val_ids), "val")

```

- **Optimizer:** Adama con lr=1e-4
- **Freezing:** none
- **Auxiliary Head:** off
- **Loss Function / Weighting:** cross-entropy + class weights euristici [1.0, 1.0, 1.0, 1.0, 5.0, 4.0, 3.5, 0.8, 0.8].
- **Cross entropy loss:**

```

#FOCUSSED CROSS ENTROPY LOSS

o
o
o class FocusedCrossEntropyLoss(nn.Module):
o def __init__(self, weight=None, lambda_confusion=2.0,
o class_a=1, class_b=3):
o super(FocusedCrossEntropyLoss, self).__init__()
o self.ce_loss = nn.CrossEntropyLoss(weight=weight,
o ignore_index=255)
o self.lambda_confusion = lambda_confusion
o self.class_a = class_a
o self.class_b = class_b
o
o
o def forward(self, pred, target):
o # pred: (B, C, H, W), logits
o # target: (B, H, W), long

```

```

 ce = self.ce_loss(pred, target)

 probs = F.softmax(pred, dim=1)
 pred_labels = torch.argmax(probs, dim=1)

 confusion_mask = (
 ((target == self.class_a) & (pred_labels ==
self.class_b)) |
 ((target == self.class_b) & (pred_labels ==
self.class_a))
)

 if confusion_mask.sum() > 0:
 a_probs = probs[:, self.class_a, :, :]
 b_probs = probs[:, self.class_b, :, :]
 confusion_penalty = (a_probs + b_probs).mean()
 else:
 confusion_penalty = torch.tensor(0.0,
device=pred.device)

 return ce + self.lambda_confusion *
confusion_penalty

```

- **Main classifier Kernel size:** due di dimensione 5

```

🔍 Primo layer convoluzionale: kernel 3x3
nn.Conv2d(256, 256, kernel_size=5, padding=2), #
mantiene dimensione
nn.ReLU(),
🔍 Secondo layer convoluzionale: kernel 5x5
nn.Conv2d(256, 256, kernel_size=5, padding=2), # campo
visivo più ampio
nn.ReLU(),

```

- **ASPP:** stride [3, 7, 3]

```

o model =
 models.segmentation.deeplabv3_resnet101(weights=weights)
o # versione custom del classifier

```

```
o model.classifier = CustomDeepLabHead(2048, NUM_CLASSES)
```

```
class CustomDeepLabHead(nn.Sequential):
 def __init__(self, in_channels, num_classes):
 super().__init__(
 ASPP(in_channels, [3, 7, 3]), # ASPP con
 dilatazioni personalizzate
 nn.Conv2d(256, 256, kernel_size=5,
 padding=2), # kernel 5x5
 nn.ReLU(),
 nn.Dropout(0.05),
 nn.Conv2d(256, num_classes, kernel_size=1)
 # output classi
)
```

- **Valutazione risultati:**

Migliore Validation Loss ottenuta all'epoca 3:

📊 Train Loss: 1.1621 | Val Loss: 1.2142 | Mean IoU: 0.2926

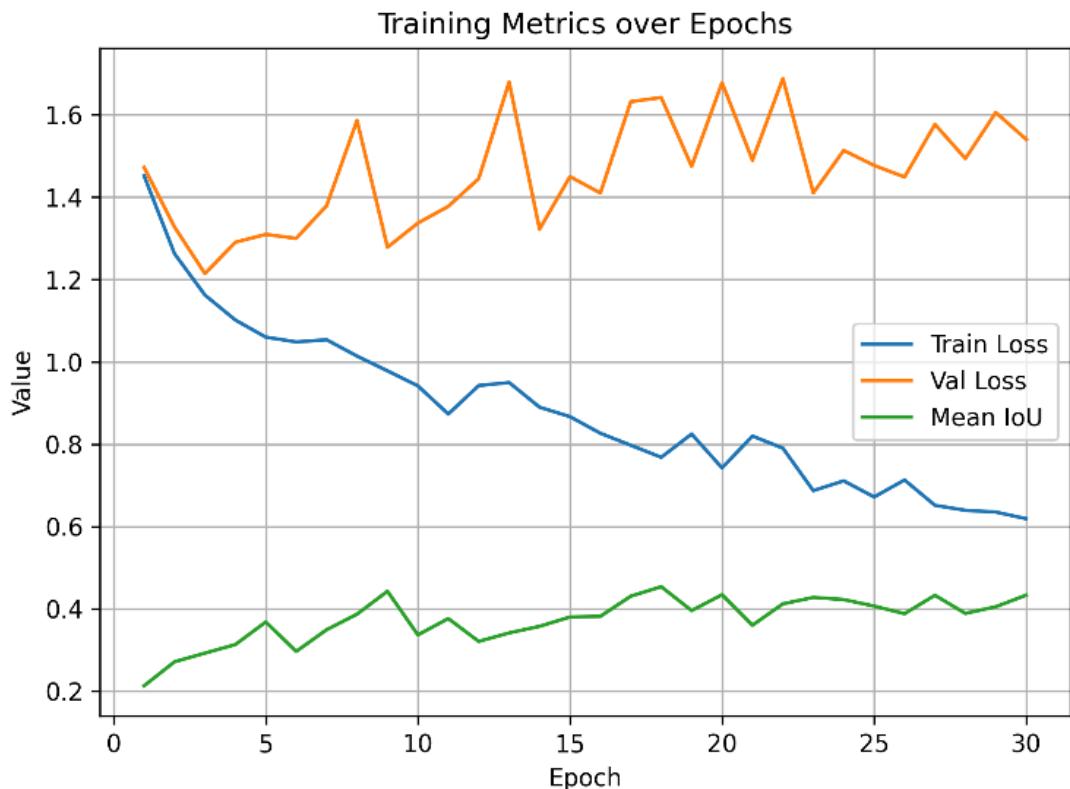
- ▶ IoU classe 0: 0.3136
- ▶ IoU classe 1: 0.0633
- ▶ IoU classe 2: 0.2910
- ▶ IoU classe 3: 0.0000
- ▶ IoU classe 4: 0.0290
- ▶ IoU classe 5: 0.0218
- ▶ IoU classe 6: 0.2568
- ▶ IoU classe 7: 0.8052
- ▶ IoU classe 8: 0.8737

Migliore mean IoU ottenuta all'epoca 18:

📊 Train Loss: 0.7681 | Val Loss: 1.6418 | Mean IoU: 0.4540

- ▶ IoU classe 0: 0.3728
- ▶ IoU classe 1: 0.4314
- ▶ IoU classe 2: 0.3458
- ▶ IoU classe 3: 0.4196
- ▶ IoU classe 4: 0.0413
- ▶ IoU classe 5: 0.4209
- ▶ IoU classe 6: 0.2676
- ▶ IoU classe 7: 0.8149
- ▶ IoU classe 8: 0.8900

- **Rappresentazione grafica del training:**



- **Massima memoria RAM occupata durante il training:** 3.6GB
- **Log di training:**  
`risultati_training\deeplabv3\resnet101\dlnr101001\training_log_dlnr101017.json`
- **Note / Bug Fix / Annotazioni:**

\* Esperimento: `deeplabv3_resnet101 - dlnr101018`

- **Architettura:** `deeplabv3plus_resnet101`
- **Backbone Weights:** DeepLabV3\_ResNet101\_Weights.DEFAULT
- **Input size:** 256x256
- **Batch size:** 4
- **Augmentation:** on

- # Augmentazione per immagini con classi rare
- rare\_aug = A.Compose([
 ○ A.HorizontalFlip(`p=0.5`) ,
 ○ A.RandomBrightnessContrast(`p=0.4`) ,
 ○ #A.Rotate(limit=20, p=0.3), potenzialmente inutile
 ])

```

o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])
o
o # Augmentazione di base
o base_aug = A.Compose([
o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])

```

```

def _load_and_augment(self, root_dir):
 sample_dirs = []
 for id in self.sample_ids:
 sample_dirs.append(os.path.join(root_dir, id))

 for dir_path in sample_dirs:
 img_path = os.path.join(dir_path, 'rgb.jpg')
 label_path = os.path.join(dir_path, 'labels.png')

 label_img = Image.open(label_path)
 is_rare =
 SegmentationDataset.contains_rare_classes(label_img,
RARE_CLASSES)

 if not self.allow_augmentation:
 is_rare = False

 # Sempre aggiungi l'originale
 self.samples.append((img_path, label_path,
is_rare))

 # Duplicazione per oversampling (solo se contiene
 classi rare)
 if is_rare and not self.is_val:
 for _ in range(DUPLICATE_FACTOR - 1):
 self.samples.append((img_path,
label_path, is_rare))

```

- **Scheduler:** off

- **Splitting:** on

```

o import numpy as np

```

```

o from PIL import Image
o import os
o from iterstrat.ml_stratifiers import
 MultilabelStratifiedShuffleSplit
o import json
o
o # sample male etichettati
o excluded_dirs = {
o "0007", "0009", "0090", "0095", "0101", "0104",
o "0105",
o "0162", "0163", "0284", "0305", "0306", "0307",
o "0308", "0309", "0310", "0311",
o "0351", "0372", "0373", "0376",
o "0498", "0499", "0500", "0501", "0526", "0527",
o "0530", "0531", "0542",
o "0564", "0585", "0586", "0587", "0588", "0589",
o "0590",
o #seguono nuove immagini da escludere
o "0000", "0001", "0052"
o }
o
o # Elenco di cartelle numeriche valide
o all_dirs = sorted([
o d for d in os.listdir(DATASET_DIR)
o if os.path.isdir(os.path.join(DATASET_DIR, d)) and
d.isdigit() and d not in excluded_dirs
o])
o
o # -----
o # COSTRUZIONE DELLA MATRICE MULTILABEL
o # -----
o def build_presence_matrix(dir_list, n_classes=9):
o """Return np.array (n_samples, n_classes) with
o booleans indicating class presence"""
o presence = np.zeros((len(dir_list), n_classes),
o dtype=int)
o
o for idx, d in enumerate(dir_list):
o label_path = os.path.join(DATASET_DIR, d,
o "labels.png")
o lbl = np.array(Image.open(label_path))
o uniq = np.unique(lbl)
o presence[idx, uniq] = 1 # segna le
o classi presenti
o
o return presence

```

```

o
o Y = build_presence_matrix(all_dirs, NUM_CLASSES)
o
o # -----
o # MULTILABEL STRATIFIED SPLIT
o # -----
o msss = MultilabelStratifiedShuffleSplit(n_splits=1,
o test_size=0.20, random_state=42)
o train_idx, val_idx =
o next(msss.split(np.zeros(len(all_dirs)), Y))
o
o train_ids = [all_dirs[i] for i in train_idx]
o val_ids = [all_dirs[i] for i in val_idx]
o
o print("Campioni:", len(train_ids), "train |",
o len(val_ids), "val")

```

- **Optimizer:** Adama con lr=1e-4
- **Freezing:** none
- **Auxiliary Head:** off
- **Loss Function / Weighting:** cross-entropy + class weights euristici [1.0, 1.0, 1.0, 1.0, 5.0, 4.0, 3.5, 0.8, 0.8].
- **Cross penalty loss:**

```

o #FOCUSSED CROSS ENTROPY LOSS
o
o
o class FocusedCrossEntropyLoss(nn.Module):
o def __init__(self, weight=None, lambda_confusion=0.5,
o class_a=1, class_b=3):
o super(FocusedCrossEntropyLoss, self).__init__()
o self.ce_loss = nn.CrossEntropyLoss(weight=weight,
o ignore_index=255)
o self.lambda_confusion = lambda_confusion
o self.class_a = class_a
o self.class_b = class_b
o
o
o def forward(self, pred, target):
o # pred: (B, C, H, W), logits
o # target: (B, H, W), long
o ce = self.ce_loss(pred, target)

```

```

 probs = F.softmax(pred, dim=1)
 pred_labels = torch.argmax(probs, dim=1)

 confusion_mask = (
 ((target == self.class_a) & (pred_labels ==
self.class_b)) |
 ((target == self.class_b) & (pred_labels ==
self.class_a))
)

 if confusion_mask.sum() > 0:
 a_probs = probs[:, self.class_a, :, :]
 b_probs = probs[:, self.class_b, :, :]
 confusion_penalty = (a_probs + b_probs).mean()
 else:
 confusion_penalty = torch.tensor(0.0,
device=pred.device)

 return ce + self.lambda_confusion *
confusion_penalty

```

- **Main classifier Kernel size:** due di dimensione 5

```

Primo layer convoluzionale: kernel 3x3
nn.Conv2d(256, 256, kernel_size=5, padding=2), #
mantiene dimensione
nn.ReLU(),
Secondo layer convoluzionale: kernel 5x5
nn.Conv2d(256, 256, kernel_size=5, padding=2), # campo
visivo più ampio
nn.ReLU(),

```

- **ASPP:** stride [3, 7, 3]

```

model =
models.segmentation.deeplabv3_resnet101(weights=weights)
versione custom del classifier
model.classifier = CustomDeepLabHead(2048, NUM_CLASSES)

```

```

class CustomDeepLabHead(nn.Sequential):
 def __init__(self, in_channels, num_classes):
 super().__init__(
 ASPP(in_channels, [3, 7, 3]), # ASPP con
 dilatazioni personalizzate
 nn.Conv2d(256, 256, kernel_size=5,
 padding=2), # kernel 5x5
 nn.ReLU(),
 nn.Dropout(0.05),
 nn.Conv2d(256, num_classes, kernel_size=1)
 # output classi
)

```

- **Valutazione risultati:**

 Train Loss: 1.1693 | Val Loss: 1.1890 | Mean IoU: 0.3508

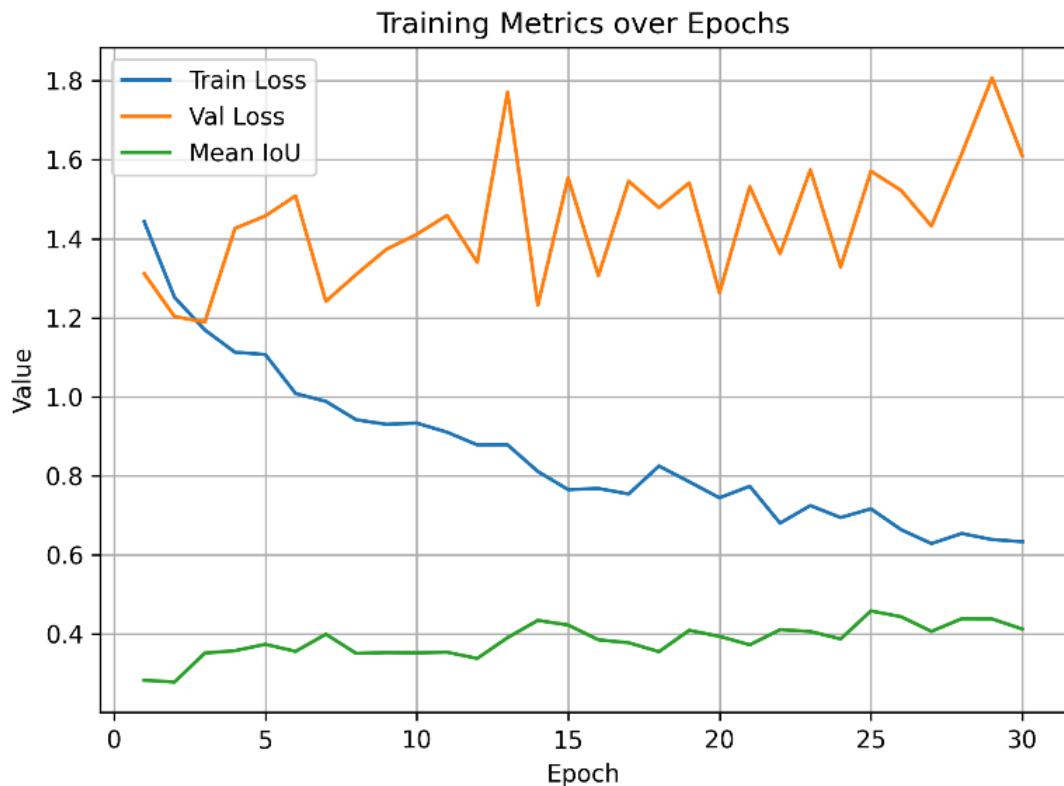
- ▶ IoU classe 0: 0.3534
- ▶ IoU classe 1: 0.0094
- ▶ IoU classe 2: 0.3062
- ▶ IoU classe 3: 0.3140
- ▶ IoU classe 4: 0.0145
- ▶ IoU classe 5: 0.2450
- ▶ IoU classe 6: 0.2446
- ▶ IoU classe 7: 0.8077
- ▶ IoU classe 8: 0.8651

Migliore mean IoU ottenuta all'epoca 25:

 Train Loss: 0.7161 | Val Loss: 1.5710 | Mean IoU: 0.4576

- ▶ IoU classe 0: 0.3657
- ▶ IoU classe 1: 0.4367
- ▶ IoU classe 2: 0.4148
- ▶ IoU classe 3: 0.5200
- ▶ IoU classe 4: 0.0254
- ▶ IoU classe 5: 0.3010
- ▶ IoU classe 6: 0.2411
- ▶ IoU classe 7: 0.8300
- ▶ IoU classe 8: 0.8915

- **Rappresentazione grafica del training:**



- **Massima memoria RAM occupata durante il training:** 3.6GB
- **Log di training:**  
`risultati_training\deeplabv3\resnet101\dlnr1010001\training_log_dlnr101018.json`
- **Note / Bug Fix / Annotazioni:**

\* Esperimento: `deeplabv3_resnet101 - dlnr101019`

- **Architettura:** `deeplabv3plus_resnet101`
- **Backbone Weights:** DeepLabV3\_ResNet101\_Weights.DEFAULT
- **Input size:** 256x256
- **Batch size:** 4
- **Augmentation:** on
  - # Augmentazione per immagini con classi rare
  - rare\_aug = A.Compose([
 A.HorizontalFlip(**p=0.5**) ,
 A.RandomBrightnessContrast(**p=0.4**) ,
 #A.Rotate(limit=20, p=0.3), potenzialmente inutile

```

o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])
o
o # Augmentazione di base
o base_aug = A.Compose([
o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])

```

```

def _load_and_augment(self, root_dir):
 sample_dirs = []
 for id in self.sample_ids:
 sample_dirs.append(os.path.join(root_dir, id))

 for dir_path in sample_dirs:
 img_path = os.path.join(dir_path, 'rgb.jpg')
 label_path = os.path.join(dir_path, 'labels.png')

 label_img = Image.open(label_path)
 is_rare =
 SegmentationDataset.contains_rare_classes(label_img,
RARE_CLASSES)

 if not self.allow_augmentation:
 is_rare = False

 # Sempre aggiungi l'originale
 self.samples.append((img_path, label_path,
is_rare))

 # Duplicazione per oversampling (solo se contiene
 classi rare)
 if is_rare and not self.is_val:
 for _ in range(DUPLICATE_FACTOR - 1):
 self.samples.append((img_path,
label_path, is_rare))

```

- **Scheduler:** off

- **Splitting:** on

```

o import numpy as np

```

```

o from PIL import Image
o import os
o from iterstrat.ml_stratifiers import
 MultilabelStratifiedShuffleSplit
o import json
o
o # sample male etichettati
o excluded_dirs = {
o "0007", "0009", "0090", "0095", "0101", "0104",
o "0105",
o "0162", "0163", "0284", "0305", "0306", "0307",
o "0308", "0309", "0310", "0311",
o "0351", "0372", "0373", "0376",
o "0498", "0499", "0500", "0501", "0526", "0527",
o "0530", "0531", "0542",
o "0564", "0585", "0586", "0587", "0588", "0589",
o "0590",
o #seguono nuove immagini da escludere
o "0000", "0001", "0052"
o }
o
o # Elenco di cartelle numeriche valide
o all_dirs = sorted([
o d for d in os.listdir(DATASET_DIR)
o if os.path.isdir(os.path.join(DATASET_DIR, d)) and
d.isdigit() and d not in excluded_dirs
o])
o
o # -----
o # COSTRUZIONE DELLA MATRICE MULTILABEL
o # -----
o def build_presence_matrix(dir_list, n_classes=9):
o """Return np.array (n_samples, n_classes) with
o booleans indicating class presence"""
o presence = np.zeros((len(dir_list), n_classes),
o dtype=int)
o
o for idx, d in enumerate(dir_list):
o label_path = os.path.join(DATASET_DIR, d,
o "labels.png")
o lbl = np.array(Image.open(label_path))
o uniq = np.unique(lbl)
o presence[idx, uniq] = 1 # segna le
o classi presenti
o
o return presence

```

```

o
o Y = build_presence_matrix(all_dirs, NUM_CLASSES)
o
o # -----
o # MULTILABEL STRATIFIED SPLIT
o # -----
o msss = MultilabelStratifiedShuffleSplit(n_splits=1,
o test_size=0.20, random_state=42)
o train_idx, val_idx =
o next(msss.split(np.zeros(len(all_dirs)), Y))
o
o train_ids = [all_dirs[i] for i in train_idx]
o val_ids = [all_dirs[i] for i in val_idx]
o
o print("Campioni:", len(train_ids), "train |",
o len(val_ids), "val")

```

- **Optimizer:** Adama con lr=1e-4
- **Freezing:** none
- **Auxiliary Head:** off
- **Loss Function / Weighting:** cross-entropy + class weights euristici [1.0, 1.0, 1.0, 1.0, 5.0, 4.0, 3.5, 0.8, 0.8].
- **Cross penalty loss:**

```

o #FOCUSSED CROSS ENTROPY LOSS
o
o
o class FocusedCrossEntropyLoss(nn.Module):
o def __init__(self, weight=None, lambda_confusion=0.1,
o class_a=1, class_b=3):
o super(FocusedCrossEntropyLoss, self).__init__()
o self.ce_loss = nn.CrossEntropyLoss(weight=weight,
o ignore_index=255)
o self.lambda_confusion = lambda_confusion
o self.class_a = class_a
o self.class_b = class_b
o
o
o def forward(self, pred, target):
o # pred: (B, C, H, W), logits
o # target: (B, H, W), long
o ce = self.ce_loss(pred, target)

```

```

 probs = F.softmax(pred, dim=1)
 pred_labels = torch.argmax(probs, dim=1)

 confusion_mask = (
 ((target == self.class_a) & (pred_labels ==
self.class_b)) |
 ((target == self.class_b) & (pred_labels ==
self.class_a))
)

 if confusion_mask.sum() > 0:
 a_probs = probs[:, self.class_a, :, :]
 b_probs = probs[:, self.class_b, :, :]
 confusion_penalty = (a_probs + b_probs).mean()
 else:
 confusion_penalty = torch.tensor(0.0,
device=pred.device)

 return ce + self.lambda_confusion *
confusion_penalty

```

- **Main classifier Kernel size:** due di dimensione 5

```

Primo layer convoluzionale: kernel 3x3
nn.Conv2d(256, 256, kernel_size=5, padding=2), #
mantiene dimensione
nn.ReLU(),
Secondo layer convoluzionale: kernel 5x5
nn.Conv2d(256, 256, kernel_size=5, padding=2), # campo
visivo più ampio
nn.ReLU(),

```

- **ASPP:** stride [3, 7, 3]

```

model =
models.segmentation.deeplabv3_resnet101(weights=weights)
versione custom del classifier
model.classifier = CustomDeepLabHead(2048, NUM_CLASSES)

```

```

class CustomDeepLabHead(nn.Sequential):
 def __init__(self, in_channels, num_classes):
 super().__init__(
 ASPP(in_channels, [3, 7, 3]), # ASPP con
 dilatazioni personalizzate
 nn.Conv2d(256, 256, kernel_size=5,
 padding=2), # kernel 5x5
 nn.ReLU(),
 nn.Dropout(0.05),
 nn.Conv2d(256, num_classes, kernel_size=1)
 # output classi
)

```

- **Valutazione risultati:**

 Train Loss: 1.0955 | Val Loss: 1.2171 | Mean IoU: 0.3295

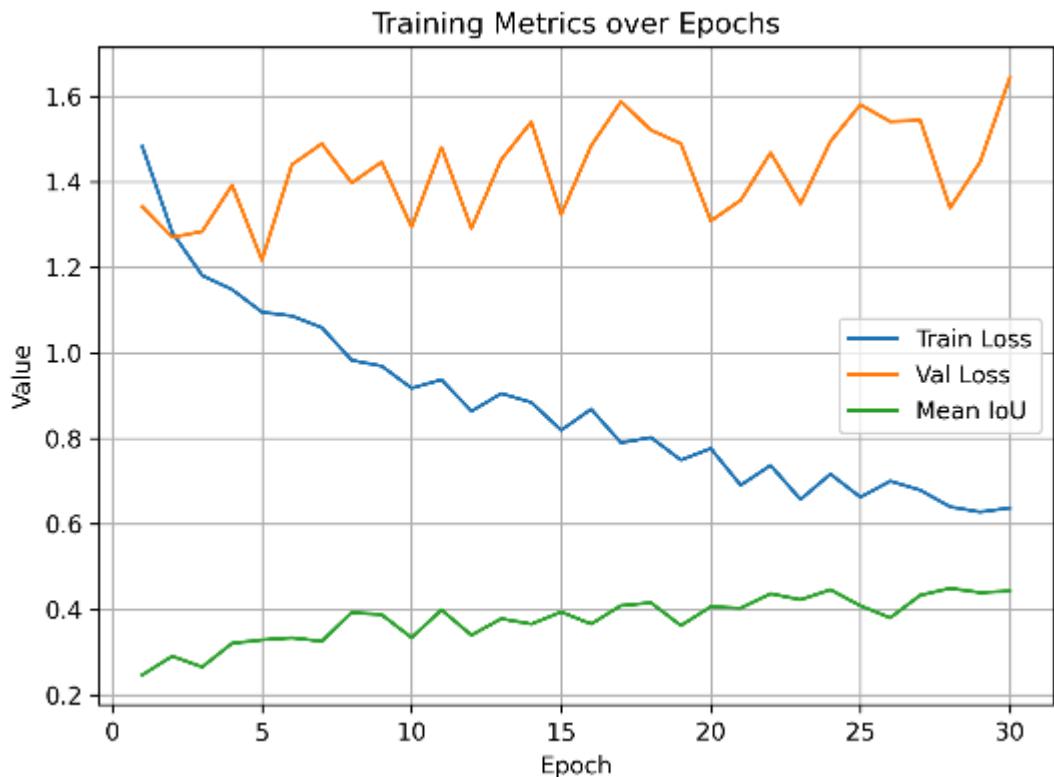
- ▶ IoU classe 0: 0.3609
- ▶ IoU classe 1: 0.1535
- ▶ IoU classe 2: 0.2541
- ▶ IoU classe 3: 0.0047
- ▶ IoU classe 4: 0.0114
- ▶ IoU classe 5: 0.3014
- ▶ IoU classe 6: 0.2420
- ▶ IoU classe 7: 0.7911
- ▶ IoU classe 8: 0.8778

Migliore mean IoU ottenuta all'epoca 28:

 Train Loss: 0.6404 | Val Loss: 1.3389 | Mean IoU: 0.4501

- ▶ IoU classe 0: 0.2255
- ▶ IoU classe 1: 0.4809
- ▶ IoU classe 2: 0.4056
- ▶ IoU classe 3: 0.4222
- ▶ IoU classe 4: 0.0484
- ▶ IoU classe 5: 0.2713
- ▶ IoU classe 6: 0.2620
- ▶ IoU classe 7: 0.8187
- ▶ IoU classe 8: 0.8914

- **Rappresentazione grafica del training:**



- **Massima memoria RAM occupata durante il training:** 4.6GB
- **Log di training:**  
`risultati_training\deeplabv3\resnet101\dlnr1010001\training_log_dlnr101019.json`
- **Note / Bug Fix / Annotazioni:**

### \* Esperimento: deeplabv3\_resnet101 - dlnr101020

- **Architettura:** `deeplabv3plus_resnet101`
- **Backbone Weights:** DeepLabV3\_ResNet101\_Weights.DEFAULT
- **Input size:** 256x256
- **Batch size:** 4
- **Augmentation:** on

```

o # Augmentazione per immagini con classi rare
o rare_aug = A.Compose([
o A.HorizontalFlip(p=0.5),
o A.RandomBrightnessContrast(p=0.4),
o #A.Rotate(limit=20, p=0.3), potenzialmente inutile

```

```

o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])
o
o
o # Augmentazione mirata per immagini che contengono sia
o # classi 1 che 3
o confused_aug = A.Compose([
o A.OneOf([
o A.MotionBlur(p=0.5),
o A.GaussianBlur(p=0.5),
o A.MedianBlur(blur_limit=5, p=0.5),
o A.GaussNoise(var_limit=(10.0, 50.0), p=0.5),
o A.Sharpen(p=0.5),
o], p=1.0),
o A.RandomBrightnessContrast(p=0.3),
o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])
o
o
o # Augmentazione di base
o base_aug = A.Compose([
o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])

```

```

def _load_and_augment(self, root_dir):
 sample_dirs = []
 for id in self.sample_ids:
 sample_dirs.append(os.path.join(root_dir, id))

 for dir_path in sample_dirs:
 img_path = os.path.join(dir_path, 'rgb.jpg')
 label_path = os.path.join(dir_path, 'labels.png')

 label_img = Image.open(label_path)
 is_rare =
SegmentationDataset.contains_rare_classes(label_img,
RARE_CLASSES)

```

```

 if not self.allow_augmentation:
 is_rare = False

 # Sempre aggiungi l'originale
 self.samples.append((img_path, label_path,
 is_rare))

 # Duplicazione per oversampling (solo se contiene
 # classi rare)
 if is_rare and not self.is_val:
 for _ in range(DUPLICATE_FACTOR - 1):
 self.samples.append((img_path,
 label_path, is_rare))

 if self.is_val:
 aug = SegmentationDataset.base_aug
 else:
 # Carica la maschera come immagine PIL per
 # controllare le classi
 label_img = Image.open(label_path)
 is_confused =
 SegmentationDataset.contains_confused_classes(label_img)

 if is_confused:
 aug = SegmentationDataset.confused_aug
 elif is_rare:
 aug = SegmentationDataset.rare_aug
 else:
 aug = SegmentationDataset.base_aug

```

- **Scheduler:** off

- **Splitting:** on

```

o import numpy as np
o from PIL import Image
o import os
o from iterstrat.ml_stratifiers import
 MultilabelStratifiedShuffleSplit
o import json
o
o # sample male etichettati
o excluded_dirs = {
o "0007", "0009", "0090", "0095", "0101", "0104",
o "0105",
o "0162", "0163", "0284", "0305", "0306", "0307",
o "0308", "0309", "0310", "0311",

```

```

o "0351", "0372", "0373", "0376",
o "0498", "0499", "0500", "0501", "0526", "0527",
o "0530", "0531", "0542",
o "0564", "0585", "0586", "0587", "0588", "0589",
o "0590",
o #seguono nuove immagini da escludere
o "0000", "0001", "0052"
o }
o
o # Elenco di cartelle numeriche valide
o all_dirs = sorted([
o d for d in os.listdir(DATASET_DIR)
o if os.path.isdir(os.path.join(DATASET_DIR, d)) and
o d.isdigit() and d not in excluded_dirs
o])
o
o # -----
o # COSTRUZIONE DELLA MATRICE MULTILABEL
o # -----
o def build_presence_matrix(dir_list, n_classes=9):
o """Return np.array (n_samples, n_classes) with
o booleans indicating class presence"""
o presence = np.zeros((len(dir_list), n_classes),
o dtype=int)
o
o for idx, d in enumerate(dir_list):
o label_path = os.path.join(DATASET_DIR, d,
o "labels.png")
o lbl = np.array(Image.open(label_path))
o uniq = np.unique(lbl)
o presence[idx, uniq] = 1 # segna le
o classi presenti
o return presence
o
o Y = build_presence_matrix(all_dirs, NUM_CLASSES)
o
o # -----
o # MULTILABEL STRATIFIED SPLIT
o # -----
o msss = MultilabelStratifiedShuffleSplit(n_splits=1,
o test_size=0.20, random_state=42)
o train_idx, val_idx =
o next(msss.split(np.zeros(len(all_dirs)), Y))
o
o train_ids = [all_dirs[i] for i in train_idx]

```

```

o val_ids = [all_dirs[i] for i in val_idx]
o
o print("Campioni:", len(train_ids), "train |",
| len(val_ids), "val")

```

- **Optimizer:** Adama con lr=1e-4
- **Freezing:** none
- **Auxiliary Head:** off
- **Loss Function / Weighting:** cross-entropy + class weights euristici [1.0, 1.0, 1.0, 1.0, 5.0, 4.0, 3.5, 0.8, 0.8].
- **Main classifier Kernel size:** due di dimensione 5

```

🎨 Primo layer convoluzionale: kernel 3x3
- nn.Conv2d(256, 256, kernel_size=5, padding=2), # mantiene dimensione
- nn.ReLU(),
🎨 Secondo layer convoluzionale: kernel 5x5
- nn.Conv2d(256, 256, kernel_size=5, padding=2), # campo visivo più ampio
- nn.ReLU(),

```

- **ASPP:** stride [3, 7, 3]

```

o model =
 models.segmentation.deeplabv3_resnet101(weights=weights)
o # versione custom del classifier
o model.classifier = CustomDeepLabHead(2048, NUM_CLASSES)

```

```

class CustomDeepLabHead(nn.Sequential):
 def __init__(self, in_channels, num_classes):
 super().__init__(
 ASPP(in_channels, [3, 7, 3]), # ASPP con dilatazioni personalizzate
 nn.Conv2d(256, 256, kernel_size=5,
padding=2), # kernel 5x5
 nn.ReLU(),
 nn.Dropout(0.05),
 nn.Conv2d(256, num_classes, kernel_size=1)
 # output class
)

```

- **Valutazione risultati:**

Migliore Validation Loss ottenuta all'epoca 3:

Train Loss: 0.4752 | Val Loss: 0.6834 | Mean IoU: 0.6026

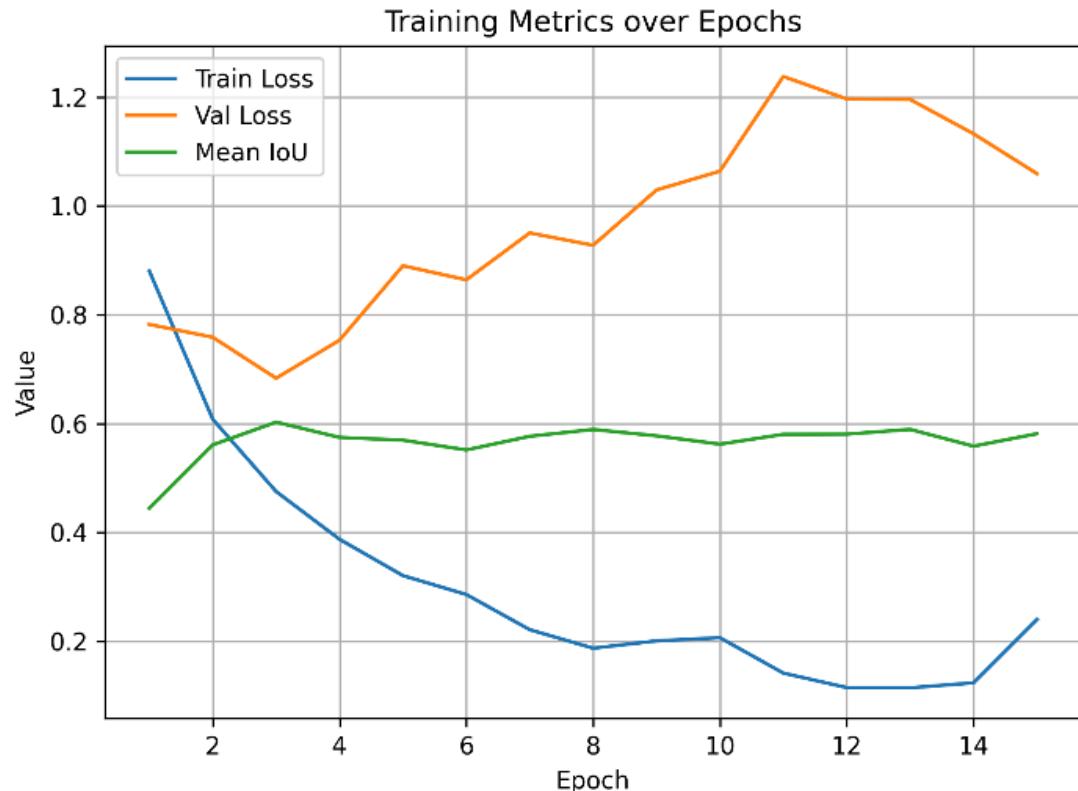
- ▶ IoU classe 0: 0.5033
- ▶ IoU classe 1: 0.6373
- ▶ IoU classe 2: 0.4822
- ▶ IoU classe 3: 0.5902
- ▶ IoU classe 4: 0.6699
- ▶ IoU classe 5: 0.4330
- ▶ IoU classe 6: 0.3208
- ▶ IoU classe 7: 0.8155
- ▶ IoU classe 8: 0.8722

Migliore mean IoU ottenuta all'epoca 3:

Train Loss: 0.4752 | Val Loss: 0.6834 | Mean IoU: 0.6026

- ▶ IoU classe 0: 0.5033
- ▶ IoU classe 1: 0.6373
- ▶ IoU classe 2: 0.4822
- ▶ IoU classe 3: 0.5902
- ▶ IoU classe 4: 0.6699
- ▶ IoU classe 5: 0.4330
- ▶ IoU classe 6: 0.3208
- ▶ IoU classe 7: 0.8155
- ▶ IoU classe 8: 0.8722

- Rappresentazione grafica del training:



- **Massima memoria RAM occupata durante il training:** 4GB
- **Log di training:**  
`risultati_training\deeplabv3\resnet101\dlnr1010001\training_log_dlnr101020.json`
- **Note / Bug Fix / Annotazioni:**

## \* Esperimento: deeplabv3\_resnet101 - dlnr101021

- **Architettura:** deeplabv3plus\_resnet101
- **Backbone Weights:** DeepLabV3\_ResNet101\_Weights.DEFAULT
- **Input size:** 256x256
- **Batch size:** 4
- **Augmentation:** on

```

o # Augmentazione per immagini con classi rare
o rare_aug = A.Compose([
o A.HorizontalFlip(p=0.5),
o A.RandomBrightnessContrast(p=0.4),
o #A.Rotate(limit=20, p=0.3), potenzialmente inutile
o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])
o
o # Augmentazione mirata per immagini che contengono sia
o # classi 1 che 3
o confused_aug = A.Compose([
o A.OneOf([
o A.MotionBlur(p=0.5),
o A.GaussianBlur(p=0.5),
o A.MedianBlur(blur_limit=5, p=0.5),
o A.GaussNoise(var_limit=(10.0, 50.0), p=0.5),
o A.Sharpen(p=0.5),
o], p=1.0),
o A.RandomBrightnessContrast(p=0.3),
o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
```

```
o ToTensorV2()
o])
o
o # Augmentazione di base
o base_aug = A.Compose([
o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])
o

- def _load_and_augment(self, root_dir):
- sample_dirs = []
- for id in self.sample_ids:
- sample_dirs.append(os.path.join(root_dir, id))

- for dir_path in sample_dirs:
- img_path = os.path.join(dir_path, 'rgb.jpg')
- label_path = os.path.join(dir_path, 'labels.png')

- label_img = Image.open(label_path)
- is_rare =
- SegmentationDataset.contains_rare_classes(label_img,
- RARE_CLASSES)

- if not self.allow_augmentation:
- is_rare = False

- # Sempre aggiungi l'originale
- self.samples.append((img_path, label_path,
is_rare))

- # Duplicazione per oversampling (solo se contiene
classi rare)
- if is_rare and not self.is_val:
- for _ in range(DUPLICATE_FACTOR - 1):
- self.samples.append((img_path,
label_path, is_rare))

- if self.is_val:
- aug = SegmentationDataset.base_aug
- else:
- # Carica la maschera come immagine PIL per
controllare le classi
- label_img = Image.open(label_path)
```

```

-- is_confused =
SegmentationDataset.contains_confused_classes(label_img)

-- if is_confused:
-- aug = SegmentationDataset.confused_aug
-- elif is_rare:
-- aug = SegmentationDataset.rare_aug
-- else:
-- aug = SegmentationDataset.base_aug
--
```

- **Scheduler:** off

- **Splitting:** on

```

o import numpy as np
o from PIL import Image
o import os
o from iterstrat.ml_stratifiers import
 MultilabelStratifiedShuffleSplit
o import json
o
o # sample male etichettati
o excluded_dirs = {
o "0007", "0009", "0090", "0095", "0101", "0104",
o "0105",
o "0162", "0163", "0284", "0305", "0306", "0307",
o "0308", "0309", "0310", "0311",
o "0351", "0372", "0373", "0376",
o "0498", "0499", "0500", "0501", "0526", "0527",
o "0530", "0531", "0542",
o "0564", "0585", "0586", "0587", "0588", "0589",
o "0590",
o #seguono nuovo immagini da escludere
o "0000", "0001", "0052"
o }
o
o # Elenco di cartelle numeriche valide
o all_dirs = sorted([
o d for d in os.listdir(DATASET_DIR)
o if os.path.isdir(os.path.join(DATASET_DIR, d)) and
d.isdigit() and d not in excluded_dirs
o])
o
o # -----
o # COSTRUZIONE DELLA MATRICE MULTILABEL
o # -----
o def build_presence_matrix(dir_list, n_classes=9):

```

```

o """Return np.array (n_samples, n_classes) with
o booleans indicating class presence"""
o
o presence = np.zeros((len(dir_list), n_classes),
o dtype=int)
o
o for idx, d in enumerate(dir_list):
o label_path = os.path.join(DATASET_DIR, d,
o "labels.png")
o lbl = np.array(Image.open(label_path))
o uniq = np.unique(lbl)
o presence[idx, uniq] = 1 # segna le
o classi presenti
o
o return presence
o
o
o Y = build_presence_matrix(all_dirs, NUM_CLASSES)
o
o
o # -----
o # MULTILABEL STRATIFIED SPLIT
o # -----
o msss = MultilabelStratifiedShuffleSplit(n_splits=1,
o test_size=0.20, random_state=42)
o train_idx, val_idx =
o next(msss.split(np.zeros(len(all_dirs)), Y))
o
o train_ids = [all_dirs[i] for i in train_idx]
o val_ids = [all_dirs[i] for i in val_idx]
o
o print("Campioni:", len(train_ids), "train |",
o len(val_ids), "val")

```

- **Optimizer:** Adama con lr=1e-4
- **Freezing:** none
- **Auxiliary Head:** on
  - if AUX\_HEAD:
    - model.aux\_classifier[-1] = nn.Conv2d(256, NUM\_CLASSES,
 kernel\_size=1)
    - else:
      - model.aux\_classifier = None
      - 
      - if self.use\_aux and "aux" in outputs:
        - aux\_output = outputs.get("aux")
        - aux\_loss = self.criterion(aux\_output, masks)
        - loss += 0.4 \* aux\_loss

- AUX\_HEAD = True
- 
- trainer = Trainer(
 ○ model=model,
 ○ train\_loader=train\_loader,
 ○ val\_loader=val\_loader,
 ○ device=DEVICE,
 ○ criterion=criterion,
 ○ optimizer=optimizer,
 ○ json\_log\_path=".//training\_log.json",
 ○ scheduler=None,
 ○ use\_aux=AUX\_HEAD
 ○ )
 ○

- **Loss Function / Weighting:** cross-entropy + class weights euristici [1.0, 1.0, 1.0, 1.0, 5.0, 4.0, 3.5, 0.8, 0.8].

- **Main classifier Kernel size:** due di dimensione 5

```

🔍 Primo layer convoluzionale: kernel 3x3
- nn.Conv2d(256, 256, kernel_size=5, padding=2), # mantiene dimensione
- nn.ReLU(),
🔍 Secondo layer convoluzionale: kernel 5x5
- nn.Conv2d(256, 256, kernel_size=5, padding=2), # campo visivo più ampio
- nn.ReLU(),

```

- **ASPP:** stride [3, 7, 3]

- model =
 ○ models.segmentation.deeplabv3\_resnet101(weights=weights)
- # versione custom del classifier
- model.classifier = CustomDeepLabHead(2048, NUM\_CLASSES)

```

class CustomDeepLabHead(nn.Sequential):
 def __init__(self, in_channels, num_classes):
 super().__init__(
 ASPP(in_channels, [3, 7, 3]), # ASPP con dilatazioni personalizzate
 nn.Conv2d(256, 256, kernel_size=5,
padding=2), # kernel 5x5
 nn.ReLU(),
 nn.Dropout(0.05),
 nn.Conv2d(256, num_classes, kernel_size=1)
 # output class
)

```

- **Valutazione risultati:**

Migliore Validation Loss ottenuta all'epoca 3:

 Train Loss: 0.6645 | Val Loss: 0.6926 | Mean IoU: 0.5837

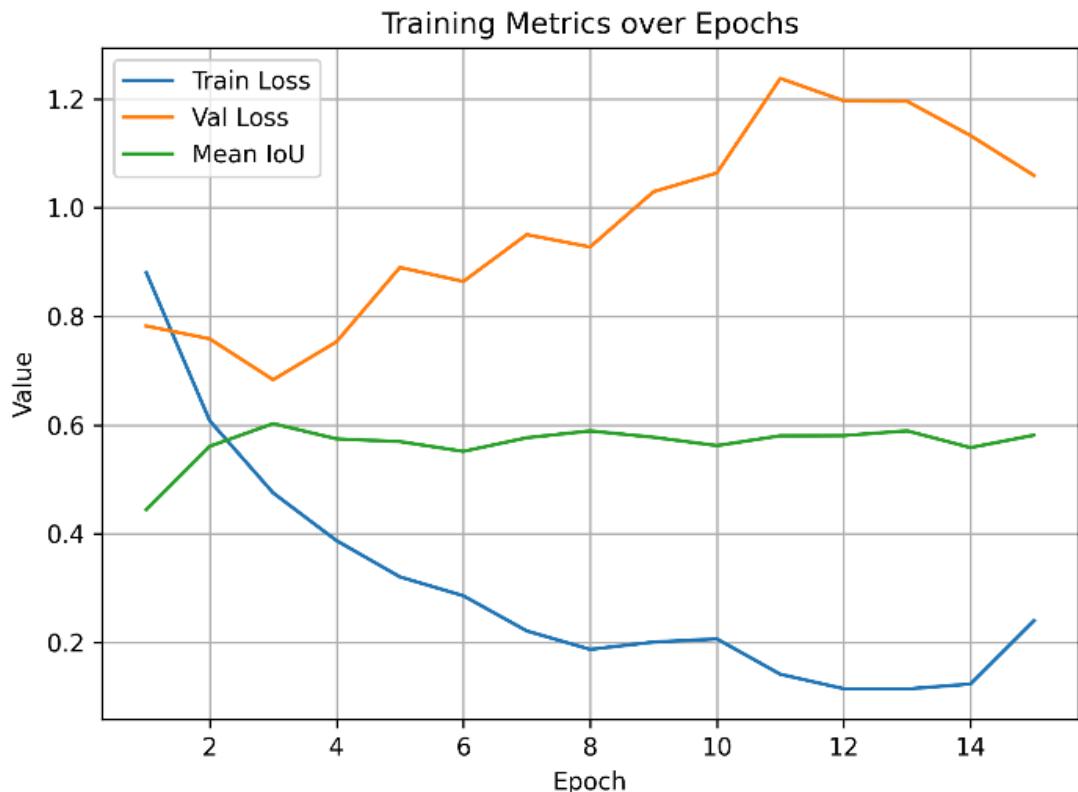
- ▶ IoU classe 0: 0.5200
- ▶ IoU classe 1: 0.5868
- ▶ IoU classe 2: 0.4574
- ▶ IoU classe 3: 0.5728
- ▶ IoU classe 4: 0.7273
- ▶ IoU classe 5: 0.3785
- ▶ IoU classe 6: 0.2812
- ▶ IoU classe 7: 0.7813
- ▶ IoU classe 8: 0.8846

Migliore mean IoU ottenuta all'epoca 9:

 Train Loss: 0.2170 | Val Loss: 0.9923 | Mean IoU: 0.6363

- ▶ IoU classe 0: 0.5503
- ▶ IoU classe 1: 0.6650
- ▶ IoU classe 2: 0.5687
- ▶ IoU classe 3: 0.5805
- ▶ IoU classe 4: 0.7819
- ▶ IoU classe 5: 0.4737
- ▶ IoU classe 6: 0.2927
- ▶ IoU classe 7: 0.8277
- ▶ IoU classe 8: 0.9003

- **Rappresentazione grafica del training:**



- **Massima memoria RAM occupata durante il training:** 2.7GB
- **Log di training:**  
`risultati_training\deeplabv3\resnet101\dlnr1010001\training_log_dlnr101021.json`
- **Note / Bug Fix / Annotazioni:**

\* Esperimento: `deeplabv3_resnet101 - dlnr101022`

- **Architettura:** `deeplabv3plus_resnet101`
- **Backbone Weights:** DeepLabV3\_ResNet101\_Weights.DEFAULT
- **Input size:** 256x256
- **Batch size:** 4
- **Augmentation:** on
  - # Augmentazione per immagini con classi rare
  - rare\_aug = A.Compose([
 A.HorizontalFlip(`p=0.5`) ,
 A.RandomBrightnessContrast(`p=0.4`) ,

```

o #A.Rotate(limit=20, p=0.3), potenzialmente inutile
o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])
o
o
o # Augmentazione mirata per immagini che contengono sia
o # classi 1 che 3
o confused_aug = A.Compose([
o A.OneOf([
o A.MotionBlur(p=0.5),
o A.GaussianBlur(p=0.5),
o A.MedianBlur(blur_limit=5, p=0.5),
o A.GaussNoise(var_limit=(10.0, 50.0), p=0.5),
o A.Sharpen(p=0.5),
o], p=1.0),
o A.RandomBrightnessContrast(p=0.3),
o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])
o
o
o # Augmentazione di base
o base_aug = A.Compose([
o A.Resize(256, 256),
o A.Normalize(mean=(0.485, 0.456, 0.406),
o std=(0.229, 0.224, 0.225)),
o ToTensorV2()
o])

```

```

- def _load_and_augment(self, root_dir):
- sample_dirs = []
- for id in self.sample_ids:
- sample_dirs.append(os.path.join(root_dir, id))
-
- for dir_path in sample_dirs:
- img_path = os.path.join(dir_path, 'rgb.jpg')
- label_path = os.path.join(dir_path, 'labels.png')
-
- label_img = Image.open(label_path)
-
```

```

 is_rare =
SegmentationDataset.contains_rare_classes(label_img,
RARE_CLASSES)

 if not self.allow_augmentation:
 is_rare = False

 # Sempre aggiungi l'originale
 self.samples.append((img_path, label_path,
is_rare))

 # Duplicazione per oversampling (solo se contiene
classi rare)
 if is_rare and not self.is_val:
 for _ in range(DUPLICATE_FACTOR - 1):
 self.samples.append((img_path,
label_path, is_rare))

if self.is_val:
 aug = SegmentationDataset.base_aug
else:
 # Carica la maschera come immagine PIL per
controllare le classi
 label_img = Image.open(label_path)
 is_confused =
SegmentationDataset.contains_confused_classes(label_img)

 if is_confused:
 aug = SegmentationDataset.confused_aug
 elif is_rare:
 aug = SegmentationDataset.rare_aug
 else:
 aug = SegmentationDataset.base_aug

```

- **Scheduler:** off
- **Splitting:** on

```

o import numpy as np
o from PIL import Image
o import os
o from iterstrat.ml_stratifiers import
 MultilabelStratifiedShuffleSplit
o import json
o
o # sample male etichettati
o excluded_dirs = {

```

```

o "0007", "0009", "0090", "0095", "0101", "0104",
o "0105",
o "0162", "0163", "0284", "0305", "0306", "0307",
o "0308", "0309", "0310", "0311",
o "0351", "0372", "0373", "0376",
o "0498", "0499", "0500", "0501", "0526", "0527",
"0530", "0531", "0542",
"0564", "0585", "0586", "0587", "0588", "0589",
"0590",
o # seguono nuove immagini da escludere
o "0000", "0001", "0052"
o }

o

o # Elenco di cartelle numeriche valide
o all_dirs = sorted([
o d for d in os.listdir(DATASET_DIR)
o if os.path.isdir(os.path.join(DATASET_DIR, d)) and
d.isdigit() and d not in excluded_dirs
o])
o

o # -----
o # COSTRUZIONE DELLA MATRICE MULTILABEL
o # -----
o def build_presence_matrix(dir_list, n_classes=9):
o """Return np.array (n_samples, n_classes) with
booleans indicating class presence"""
o presence = np.zeros((len(dir_list), n_classes),
o dtype=int)

o

o for idx, d in enumerate(dir_list):
o label_path = os.path.join(DATASET_DIR, d,
"labels.png")
o lbl = np.array(Image.open(label_path))
o uniq = np.unique(lbl)
o presence[idx, uniq] = 1 # segna le
classi presenti
o return presence

o

o Y = build_presence_matrix(all_dirs, NUM_CLASSES)

o

o # -----
o # MULTILABEL STRATIFIED SPLIT
o # -----
o msss = MultilabelStratifiedShuffleSplit(n_splits=1,
test_size=0.20, random_

```

```

o train_idx, val_idx =
 next(msss.split(np.zeros(len(all_dirs)), Y))

o

o train_ids = [all_dirs[i] for i in train_idx]
o val_ids = [all_dirs[i] for i in val_idx]
o

o print("Campioni:", len(train_ids), "train |",
 len(val_ids), "val")

```

- **Optimizer:** Adama con lr=1e-4

- **Freezing:** none

- **Auxiliary Head:** on

```

o if AUX_HEAD:
o model.aux_classifier[-1] = nn.Conv2d(256, NUM_CLASSES,
 kernel_size=1)
o else:
o model.aux_classifier = None
o
o if self.use_aux and "aux" in outputs:
o aux_output = outputs.get("aux")
o aux_loss = self.criterion(aux_output, masks)
o loss += 0.3 * aux_loss
o AUX_HEAD = True
o
o trainer = Trainer(
o model=model,
o train_loader=train_loader,
o val_loader=val_loader,
o device=DEVICE,
o criterion=criterion,
o optimizer=optimizer,
o json_log_path=".//training_log.json",
o scheduler=None,
o use_aux=AUX_HEAD
o)
o

```

- **Loss Function / Weighting:** cross-entropy + class weights euristici [1.0, 1.0, 1.0, 1.0, 5.0, 4.0, 3.5, 0.8, 0.8].

- **Main classifier Kernel size:** due di dimensione 5

```

- # 🔍 Primo layer convoluzionale: kernel 3x3
- nn.Conv2d(256, 256, kernel_size=5, padding=2), #
| mantiene dimensione

```

```

 nn.ReLU(),
 # Secondo layer convoluzionale: kernel 5x5
 nn.Conv2d(256, 256, kernel_size=5, padding=2), # campo
 visivo più ampio
 nn.ReLU(),

```

- **ASPP: stride [3, 7, 3]**

```

○ model =
 models.segmentation.deeplabv3_resnet101(weights=weights)
○ # versione custom del classifier
○ model.classifier = CustomDeepLabHead(2048, NUM_CLASSES)

```

```

class CustomDeepLabHead(nn.Sequential):
 def __init__(self, in_channels, num_classes):
 super().__init__(
 ASPP(in_channels, [3, 7, 3]), # ASPP con
 dilatazioni personalizzate
 nn.Conv2d(256, 256, kernel_size=5,
 padding=2), # kernel 5x5
 nn.ReLU(),
 nn.Dropout(0.05),
 nn.Conv2d(256, num_classes, kernel_size=1)
 # output classi
)

```

- **Valutazione risultati:**

Migliore Validation Loss ottenuta all'epoca 2:

 Train Loss: 0.7801 | Val Loss: 0.7652 | Mean IoU: 0.5487

- ▶ IoU classe 0: 0.5069
- ▶ IoU classe 1: 0.6413
- ▶ IoU classe 2: 0.3969
- ▶ IoU classe 3: 0.5744
- ▶ IoU classe 4: 0.5398
- ▶ IoU classe 5: 0.3108
- ▶ IoU classe 6: 0.2558
- ▶ IoU classe 7: 0.7924
- ▶ IoU classe 8: 0.8784

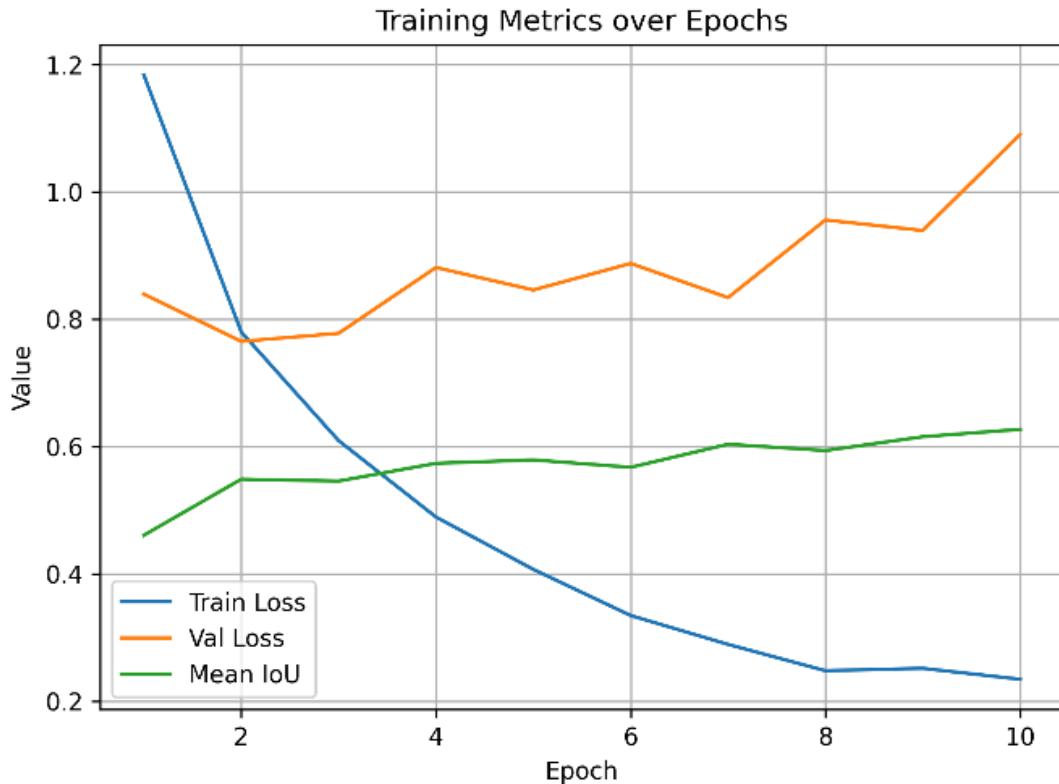
Migliore mean IoU ottenuta all'epoca 10:

 Train Loss: 0.2344 | Val Loss: 1.0903 | Mean IoU: 0.6270

- ▶ IoU classe 0: 0.5482
- ▶ IoU classe 1: 0.6693
- ▶ IoU classe 2: 0.5924
- ▶ IoU classe 3: 0.6273
- ▶ IoU classe 4: 0.6607

- ▶ IoU classe 5: 0.4515
- ▶ IoU classe 6: 0.2718
- ▶ IoU classe 7: 0.8442
- ▶ IoU classe 8: 0.8990

- **Rappresentazione grafica del training:**



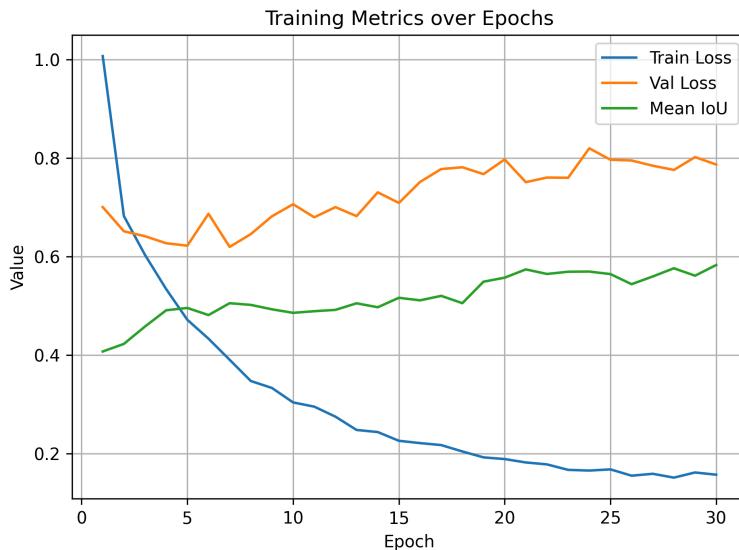
- **Massima memoria RAM occupata durante il training:** 2.7GB
- **Log di training:**  
`risultati_training\deeplabv3\resnet101\dlnr1010001\training_log_dlnr101022.json`
- **Note / Bug Fix / Annotazioni:**

## deeplabv3\_mobilenet\_v3\_large

\* Esperimento: `deeplabv3_mobilenet_v3_large - DLMN000`

- **Architettura:** `deeplabv3plus_mobilenet_v3_large`
- **Backbone Weights:** DeepLabV3\_MobileNet\_V3\_Large\_Weights.DEFAULT
- **Input size:** 256x256

- **Batch size:** 4
- **Augmentation:** off
- **Scheduler:** off
- **Optimizer:** Adama con lr=1e-4
- **Freezing:** none
- **Auxiliary Head:** off
- **Loss Function / Weighting:** cross-entropy
- **Main classifier Kernel size:** 1
- **Valutazione risultati:**
  - Migliore Validation Loss ottenuta all'epoca 7:
    - Train Loss: 0.3904 | Val Loss: 0.6196 | Mean IoU: 0.5052
    - ▶ IoU classe 0: 0.3568
    - ▶ IoU classe 1: 0.5695
    - ▶ IoU classe 2: 0.5950
    - ▶ IoU classe 3: 0.5364
    - ▶ IoU classe 4: 0.0000
    - ▶ IoU classe 5: 0.5020
    - ▶ IoU classe 6: 0.1968
    - ▶ IoU classe 7: 0.8017
    - ▶ IoU classe 8: 0.8404
  - Migliore **mean IoU** ottenuta all'epoca 30:
    - Train Loss: 0.1570 | Val Loss: 0.7868 | Mean IoU: 0.5828
    - ▶ IoU classe 0: 0.4531
    - ▶ IoU classe 1: 0.5564
    - ▶ IoU classe 2: 0.6296
    - ▶ IoU classe 3: 0.5506
    - ▶ IoU classe 4: 0.5626
    - ▶ IoU classe 5: 0.5171
    - ▶ IoU classe 6: 0.1771
    - ▶ IoU classe 7: 0.8165
    - ▶ IoU classe 8: 0.8523
- **Rappresentazione grafica del training:**

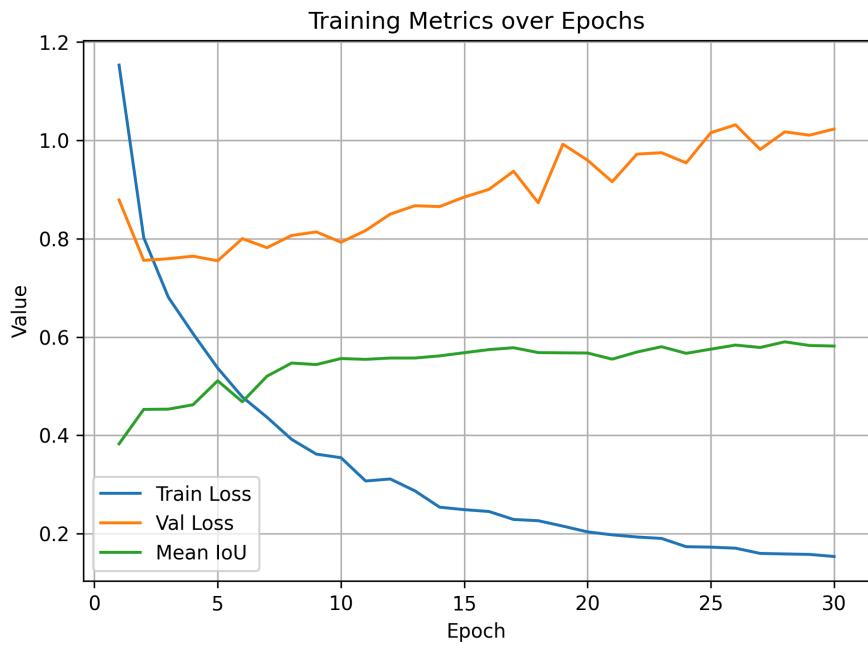


- **Massima memoria RAM occupata durante il training:** 0.56GB
- **Log di training:**  
`risultati_training\deeplabv3\mobilenetv3_large\dlmnv3l0001\training_log_dlmnv3l0000.json`
- **Note / Bug Fix / Annotazioni:**
  - —

\* Esperimento: **deeplabv3\_mobilenet\_v3\_large - DLMN001**

- **Architettura:** `deeplabv3plus_mobilenet_v3_large`
- **Backbone Weights:** DeepLabV3\_MobileNet\_V3\_Large\_Weights.DEFAULT
- **Input size:** 256x256
- **Batch size:** 4
- **Augmentation:** off
- **Scheduler:** off
- **Optimizer:** Adama con lr=1e-4
- **Freezing:** none
- **Auxiliary Head:** off

- **Loss Function / Weighting:** cross-entropy + class weights euristici [1.0, 1.0, 1.0, 1.0, 5.0, 4.0, 3.5, 0.8, 0.8].
- **Main classifier Kernel size: 1**
- **Valutazione risultati:**
  - Migliore Validation Loss ottenuta all'epoca 5:
    - Train Loss: 0.5369 | Val Loss: 0.7551 | Mean IoU: 0.5111
    - ▶ IoU classe 0: 0.3274
    - ▶ IoU classe 1: 0.4677
    - ▶ IoU classe 2: 0.5834
    - ▶ IoU classe 3: 0.5334
    - ▶ IoU classe 4: 0.2439
    - ▶ IoU classe 5: 0.4288
    - ▶ IoU classe 6: 0.2185
    - ▶ IoU classe 7: 0.7925
    - ▶ IoU classe 8: 0.8208
  - Migliore **mean IoU** ottenuta all'epoca 26:
    - Train Loss: 0.1706 | Val Loss: 1.0316 | Mean IoU: 0.5837
    - ▶ IoU classe 0: 0.4317
    - ▶ IoU classe 1: 0.5694
    - ▶ IoU classe 2: 0.5931
    - ▶ IoU classe 3: 0.5216
    - ▶ IoU classe 4: 0.6083
    - ▶ IoU classe 5: 0.5410
    - ▶ IoU classe 6: 0.1721
    - ▶ IoU classe 7: 0.8126
    - ▶ IoU classe 8: 0.8511
- **Rappresentazione grafica del training:**



- **Massima memoria RAM occupata durante il training:** 0.56GB
- **Log di training:**  
`risultati_training\deeplabv3\mobilenetv3_large\dlnnv3l0001\training_log_dlnnv3l0001.json`
- **Note / Bug Fix / Annotazioni:**

○ —